



LightSecAgg: a Lightweight and Versatile Design for Secure Aggregation in Federated Learning

Jinhyun So, Ramy E. Ali, Chaoyang He, Salman Avestimehr

University of Southern California

Joint work with

Chien-Sheng Yang (USC), Songze Li (HKUST), Qian Yu (USC), and Basak Guler (UCR)

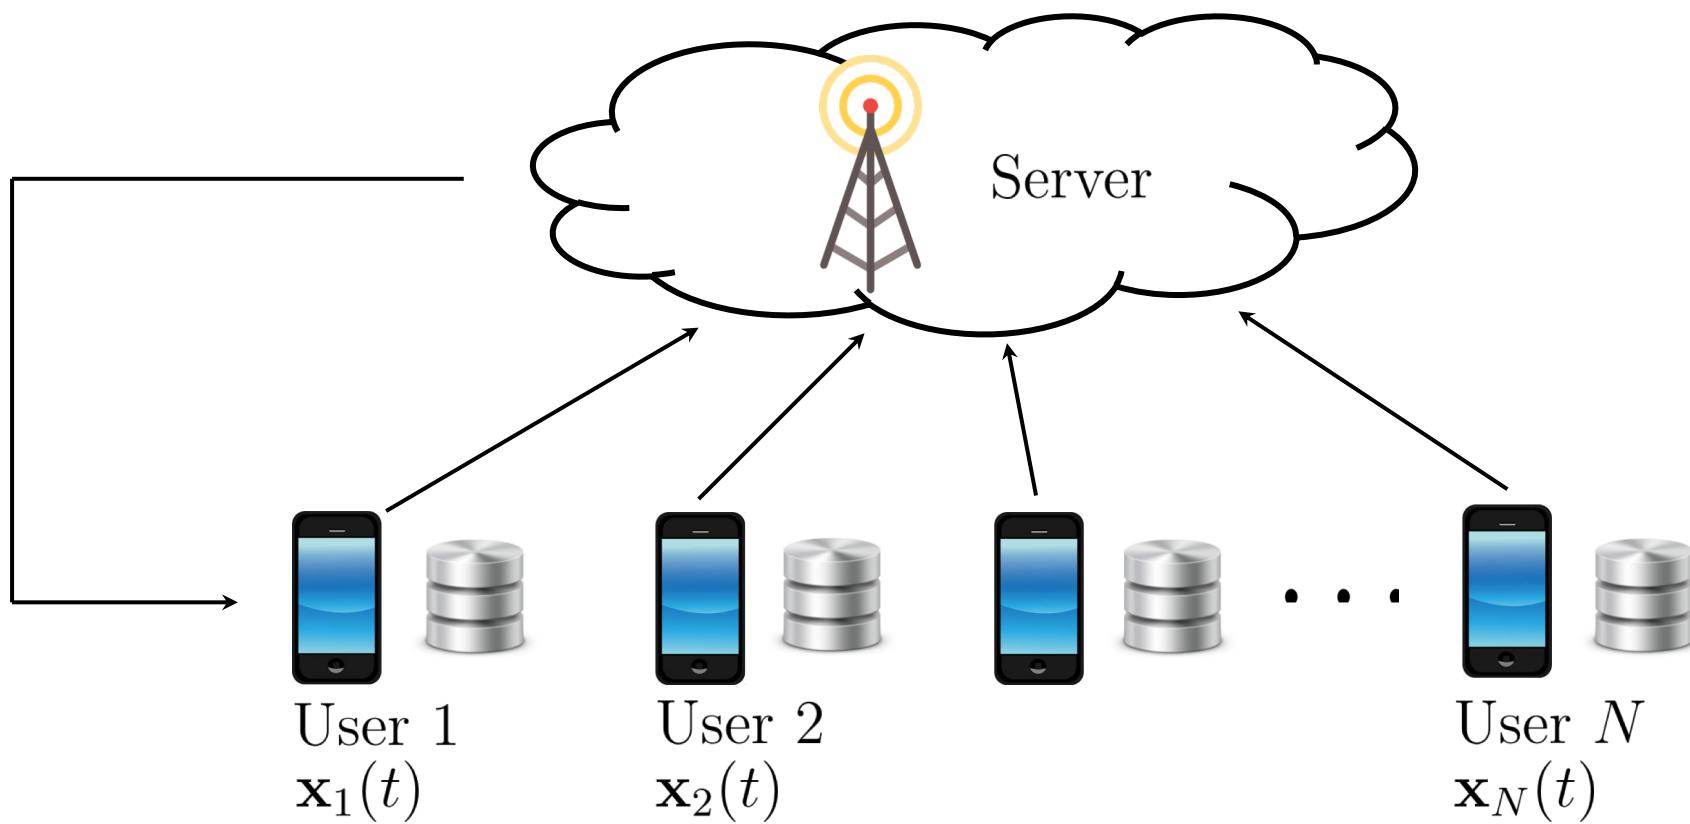
Executive Summary

- Part 1. General description of LightSecAgg
- Part 2. LightSecAgg for Asynchronous Federated Learning
- Part 3. System Design and Experiments

Part 1. General description of LightSecAgg

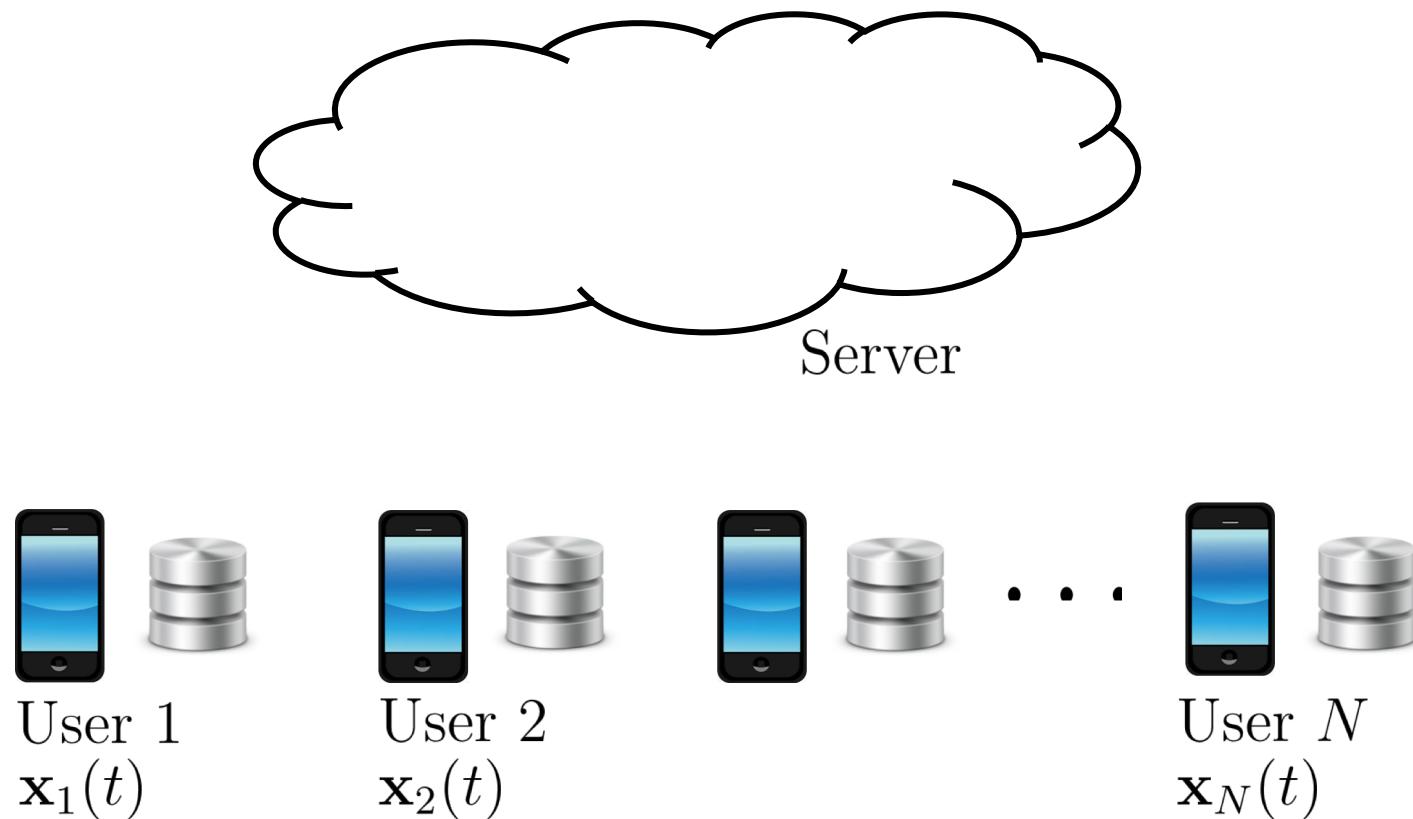
Federated Learning

Machine learning on massive amount of data collected on many users/mobile devices



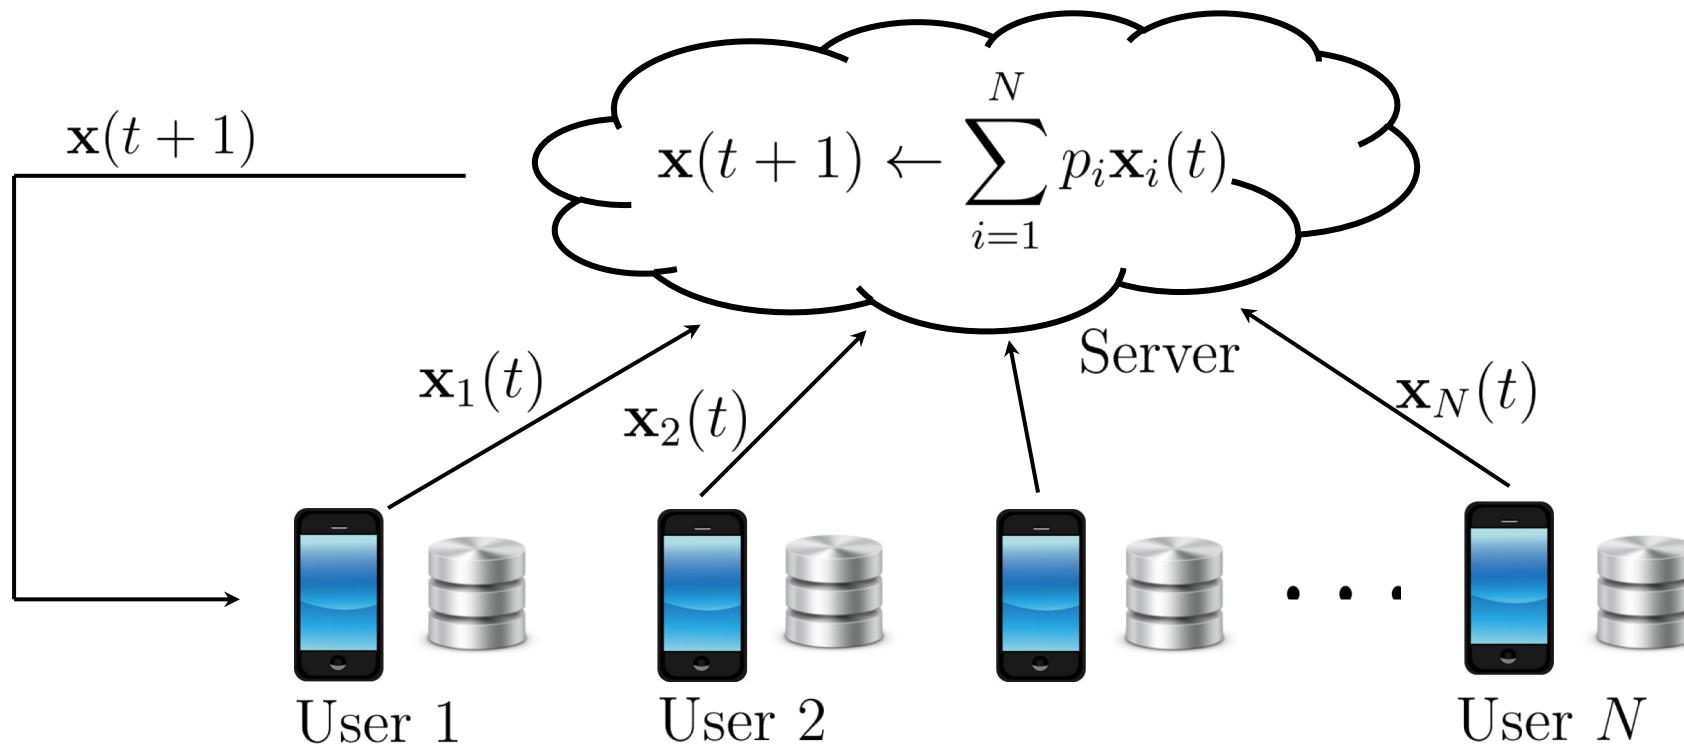
Federated Learning

Step 1: Train **locally** at each user on the available data



Federated Learning

Step 2: Aggregate the local models at the server, update the mode, and send the new global model back to the users

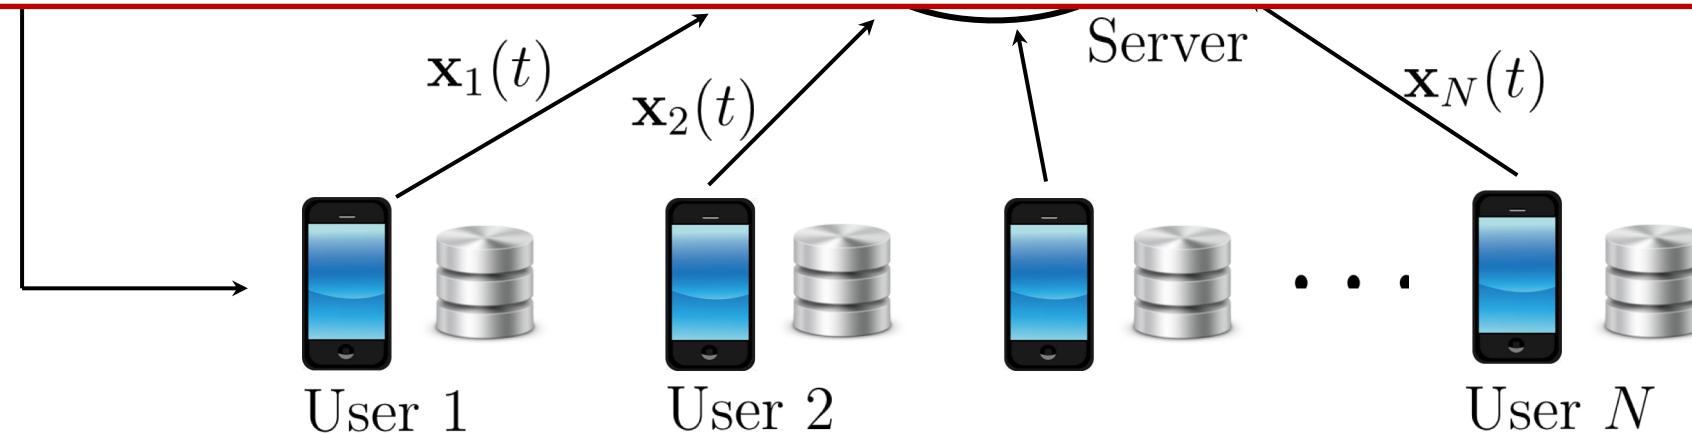


Federated Learning

Step 2: Aggregate the local models at the server, update the mode, and send the new global model back to the users

- **Key Design Principles**

1. Privacy
2. Communication Efficiency



Model Inversion Attack

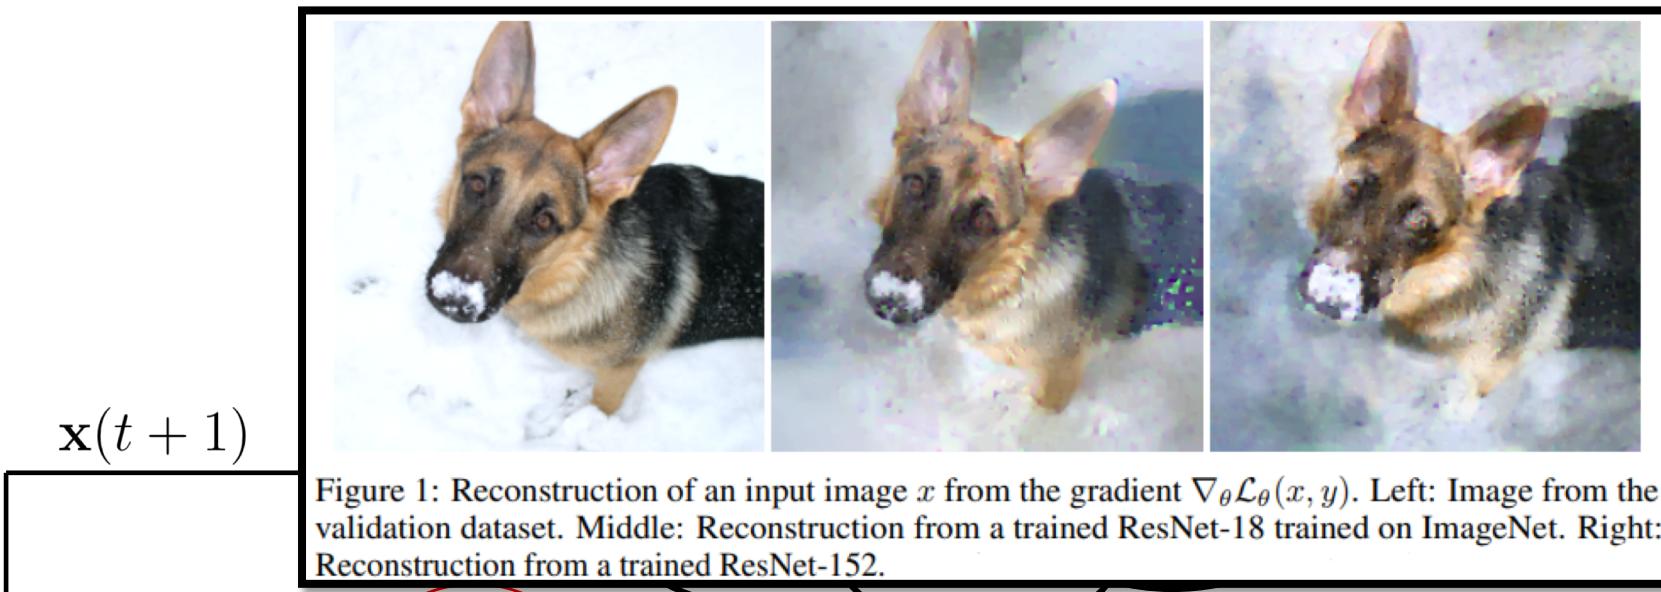
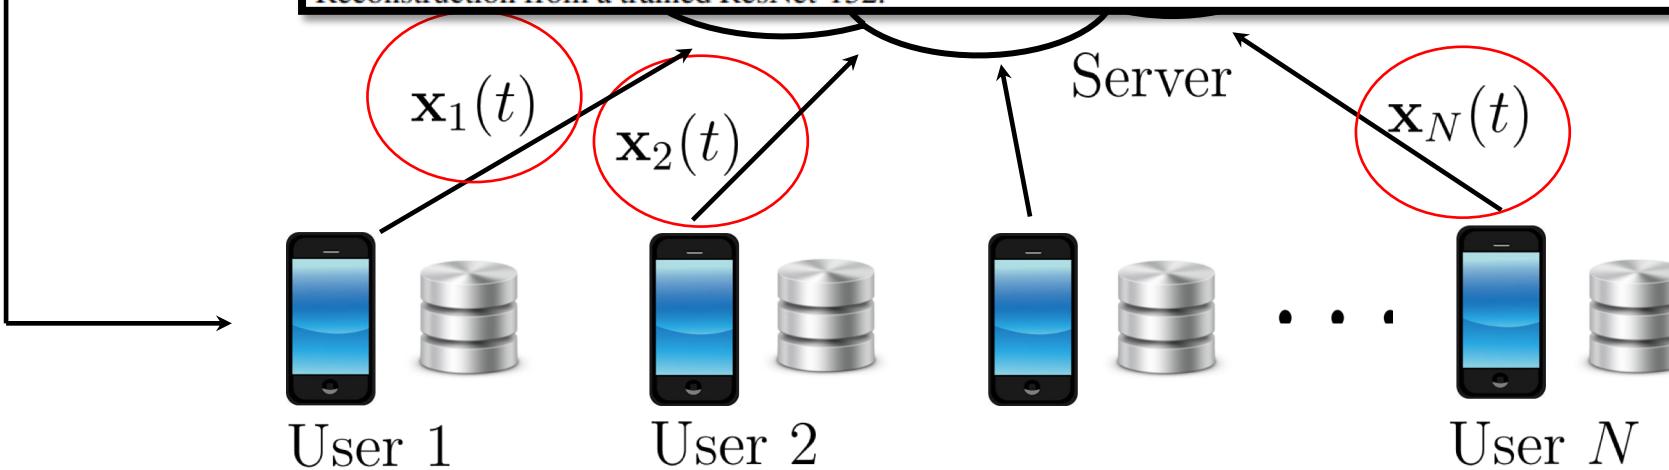


Figure 1: Reconstruction of an input image x from the gradient $\nabla_{\theta} \mathcal{L}_{\theta}(x, y)$. Left: Image from the validation dataset. Middle: Reconstruction from a trained ResNet-18 trained on ImageNet. Right: Reconstruction from a trained ResNet-152.

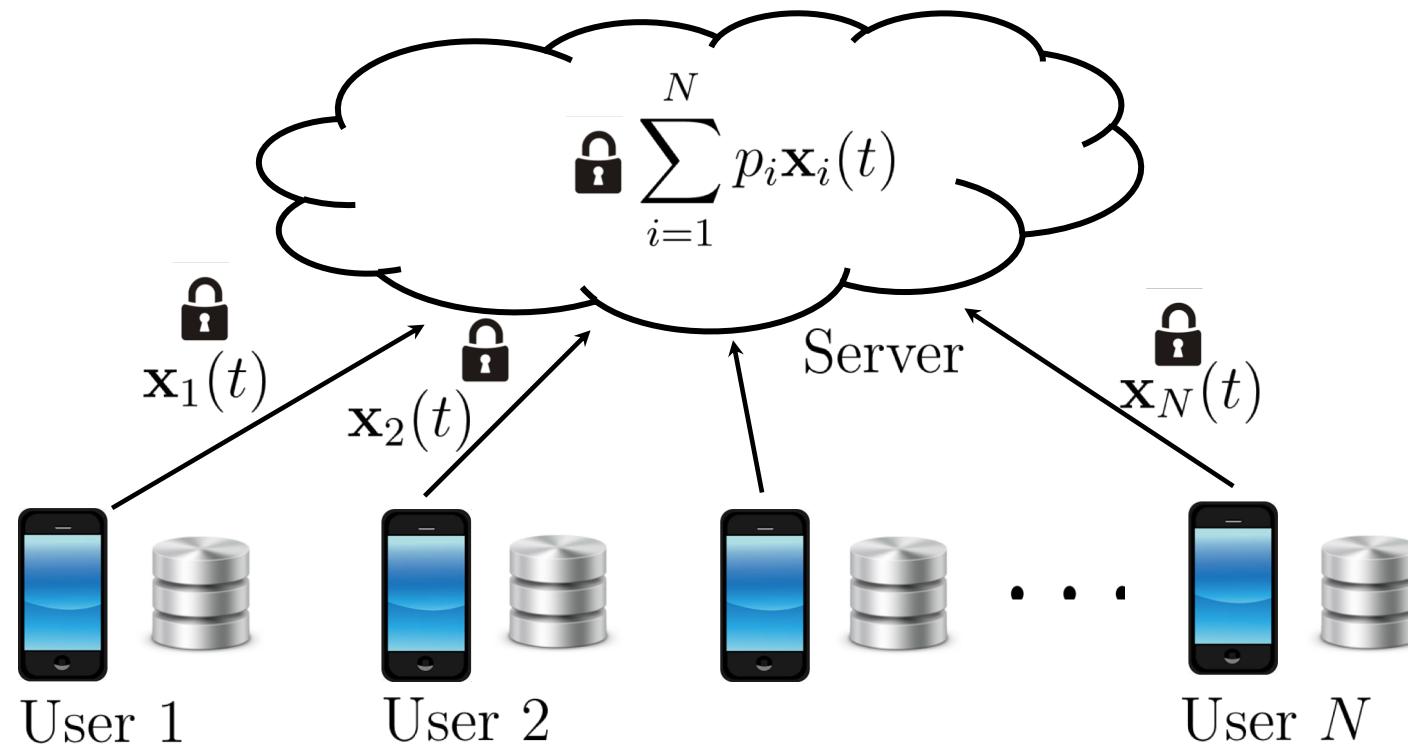


Problem: Individual model update can leak sensitive data.

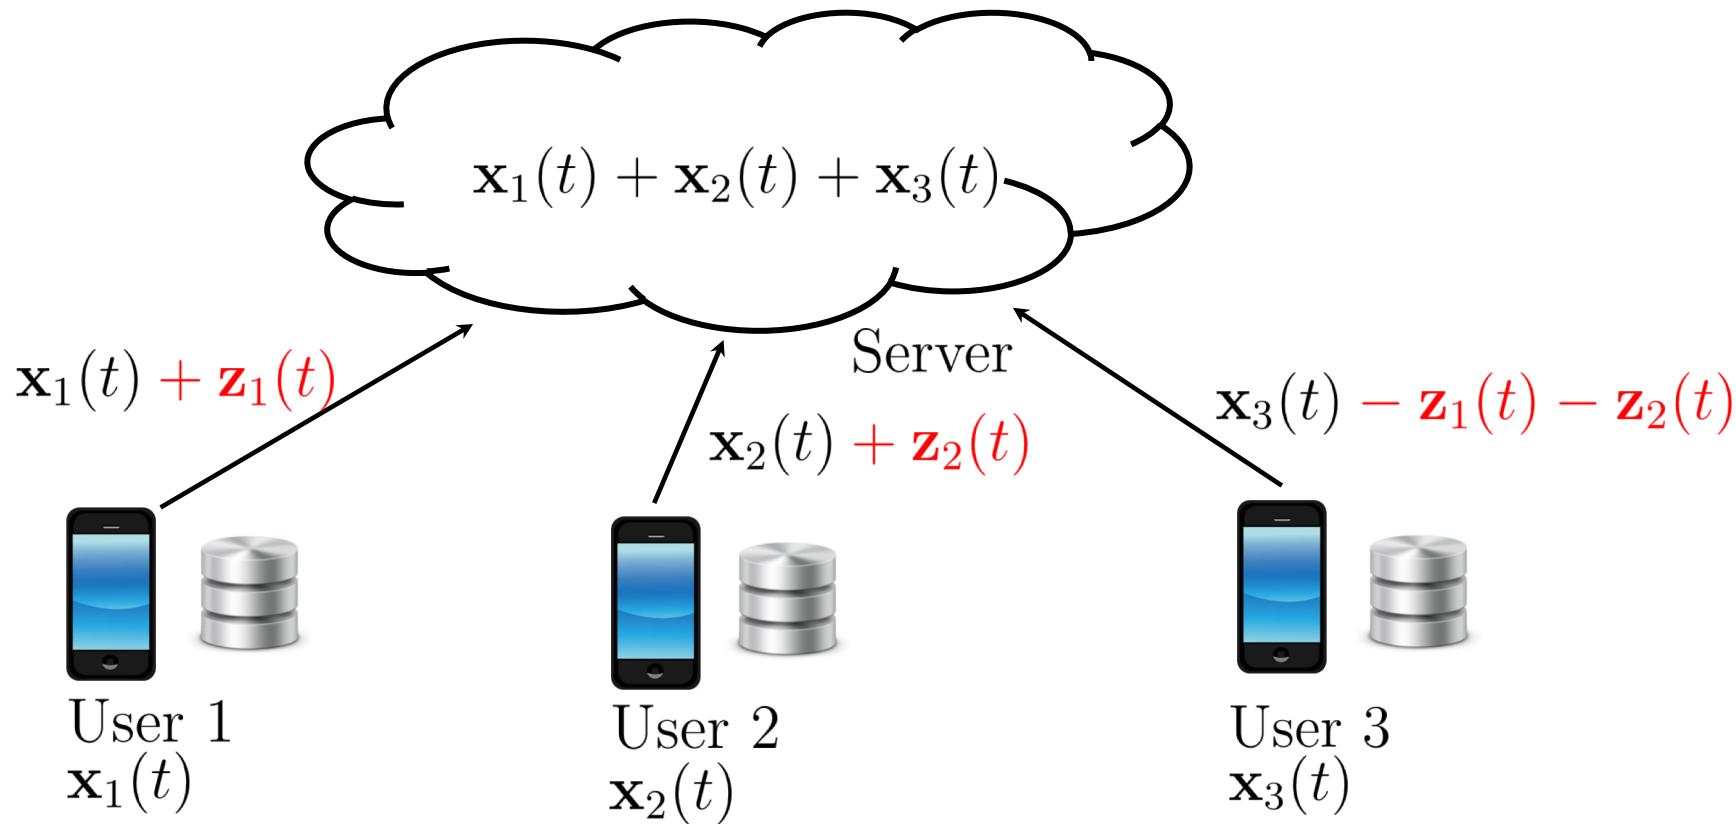
Key component: Secure Aggregation

Threat model:

- honest, but curious
- up to a fraction of users and/or server colluding



Key component: Secure Aggregation



Perfectly secure: the server only learns the aggregated model.

Key component: Secure Aggregation



$\mathbf{x}_1(t)$



User 1
 $\mathbf{x}_1(t)$

User 2
 $\mathbf{x}_2(t)$

User 3
 $\mathbf{x}_3(t)$

Practical Secure Aggregation for Privacy-Preserving Machine Learning

Keith Bonawitz*, Vladimir Ivanov*, Ben Kreuter*,
Antonio Marcedone^{††}, H. Brendan McMahan*, Sarvar Patel*,
Daniel Ramage*, Aaron Segal*, and Karn Seth*

*{bonawitz,vlivan,benkreuter,mcmahan,
sarvar,dramage,asegal,karn}@google.com

Google, Mountain View, CA 94043

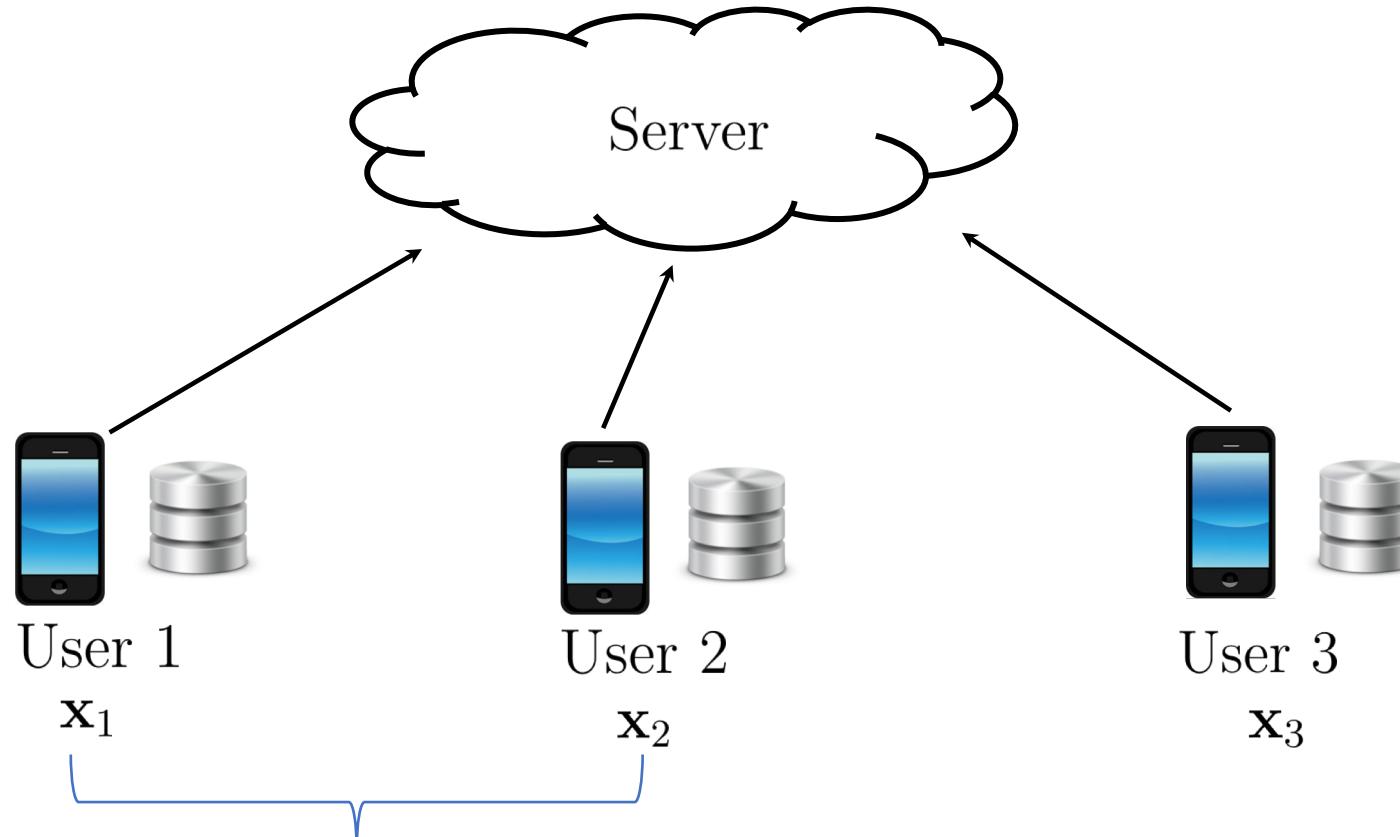
[†]marcedone@cs.cornell.edu

Cornell Tech, 2 West Loop Rd., New York, NY 10044



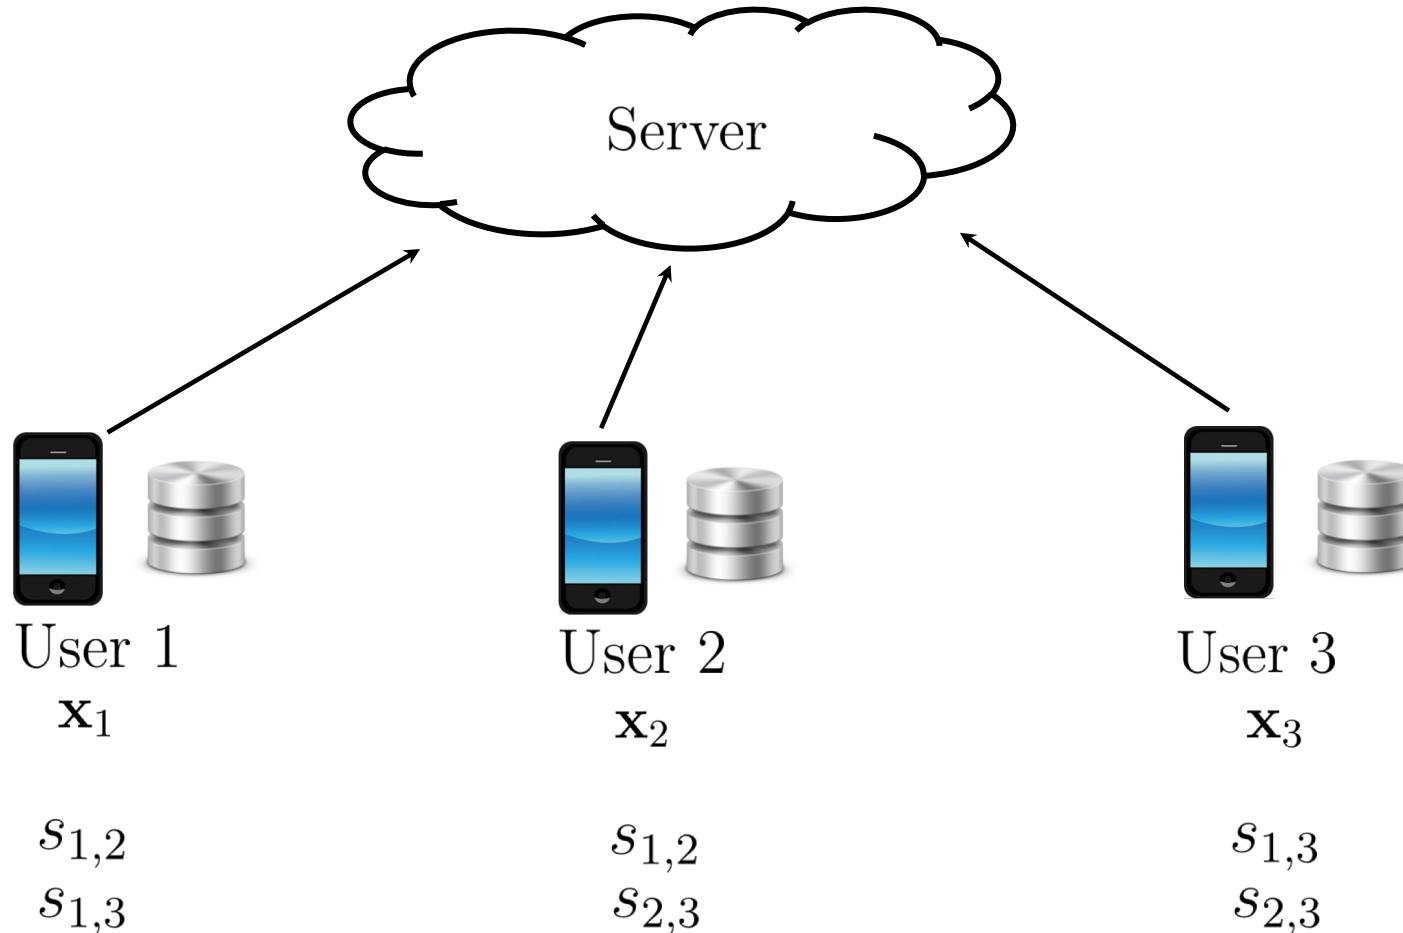
How to create an additive structure allowing global aggregation
amidst possible user dropouts?

Example (SecAgg)



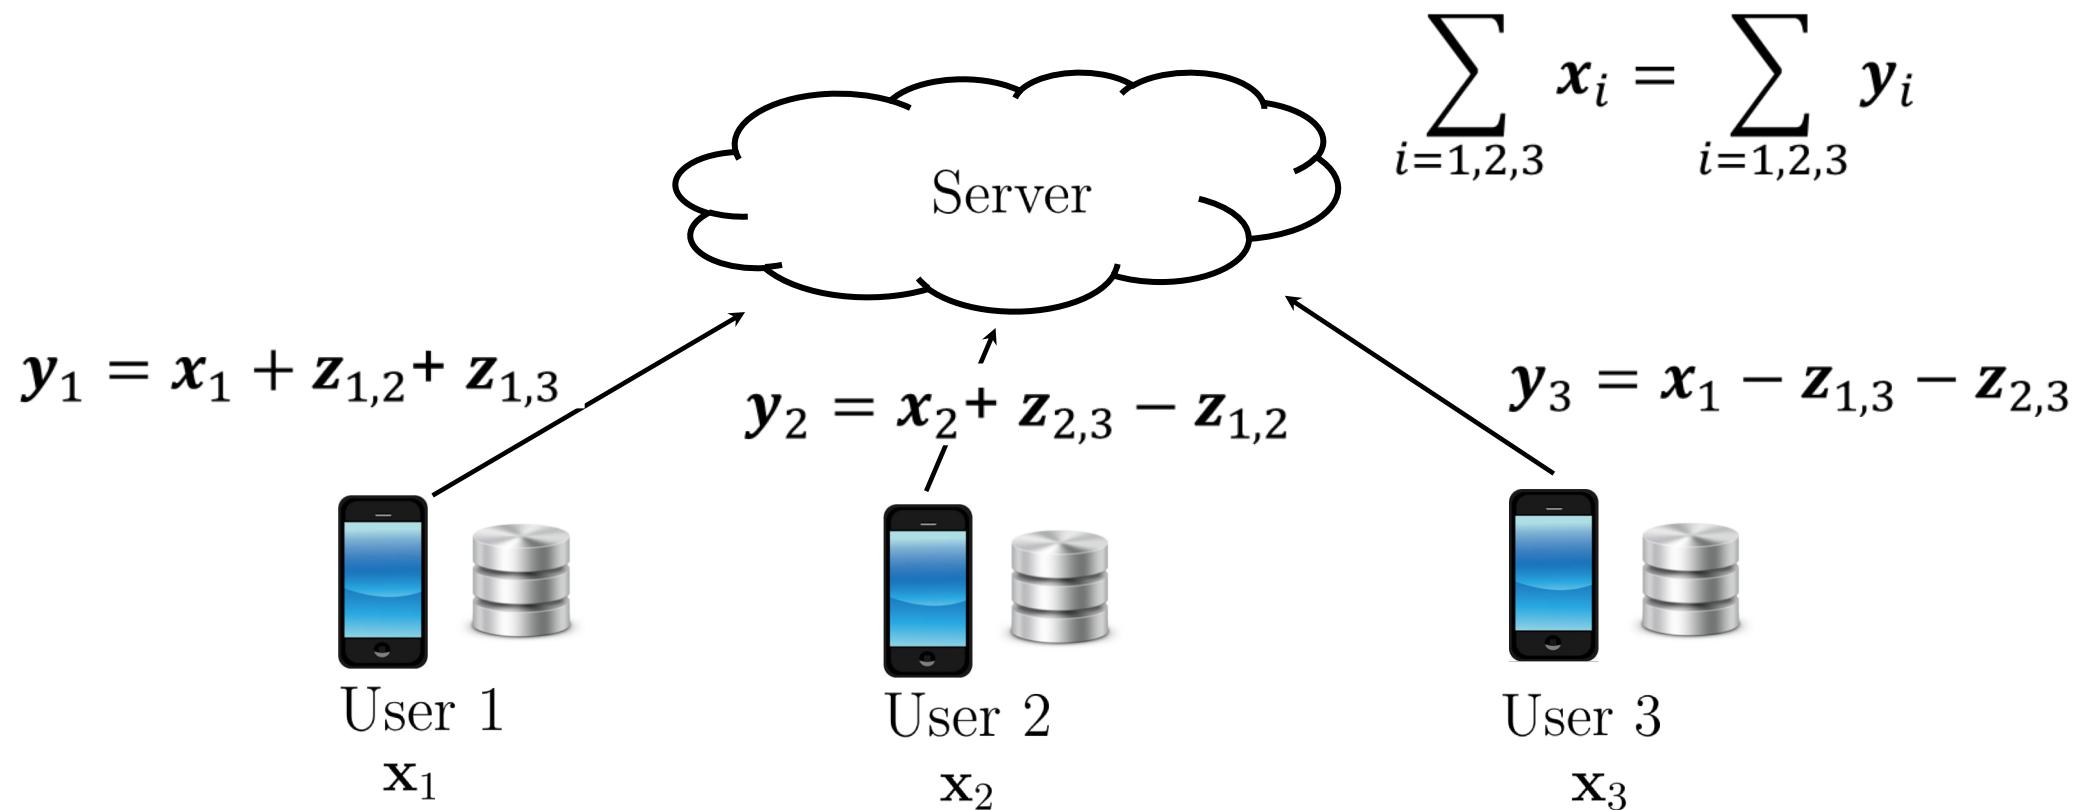
agree on pairwise random seed $s_{1,2}$

Example (SecAgg)



Pairwise random seed

Example (SecAgg)



$$z_{1,2} = \text{PRG}(s_{1,2})$$

$$z_{1,3} = \text{PRG}(s_{1,3})$$

$$z_{1,2} = \text{PRG}(s_{1,2})$$

$$z_{2,3} = \text{PRG}(s_{2,3})$$

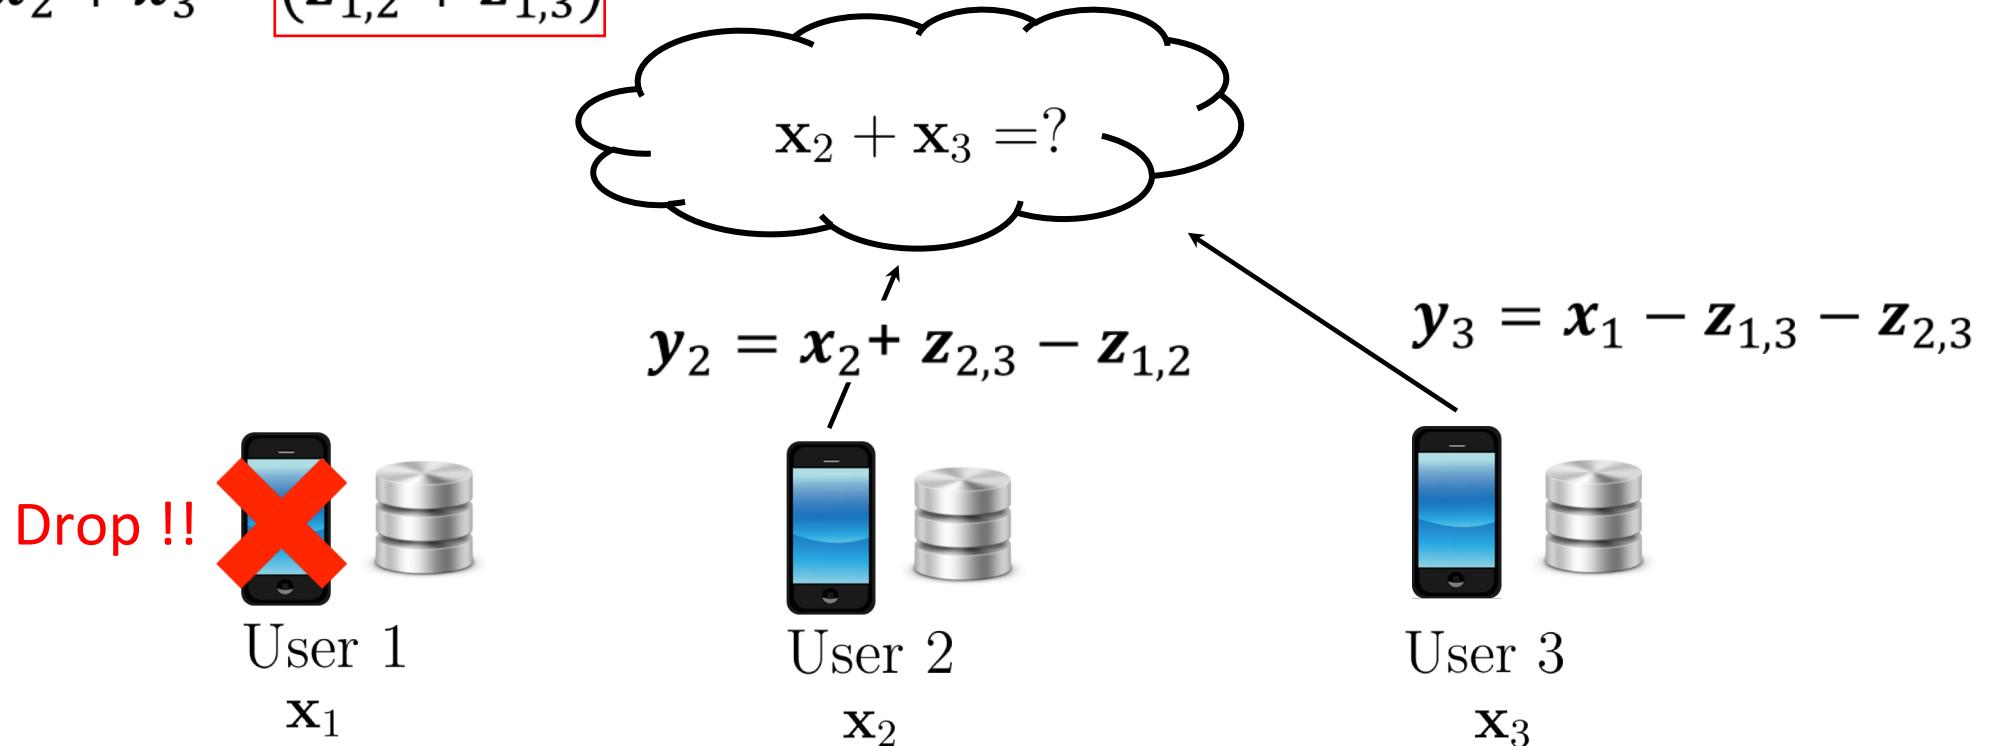
$$z_{1,3} = \text{PRG}(s_{1,3})$$

$$z_{2,3} = \text{PRG}(s_{2,3})$$

Example (SecAgg)

Remaining random vectors

$$y_2 + y_3 = x_2 + x_3 - (\mathbf{z}_{1,2} + \mathbf{z}_{1,3})$$



$$\mathbf{z}_{1,2} = \text{PRG}(s_{1,2})$$

$$\mathbf{z}_{1,2} = \text{PRG}(s_{1,2})$$

$$\mathbf{z}_{1,3} = \text{PRG}(s_{1,3})$$

$$\mathbf{z}_{1,3} = \text{PRG}(s_{1,3})$$

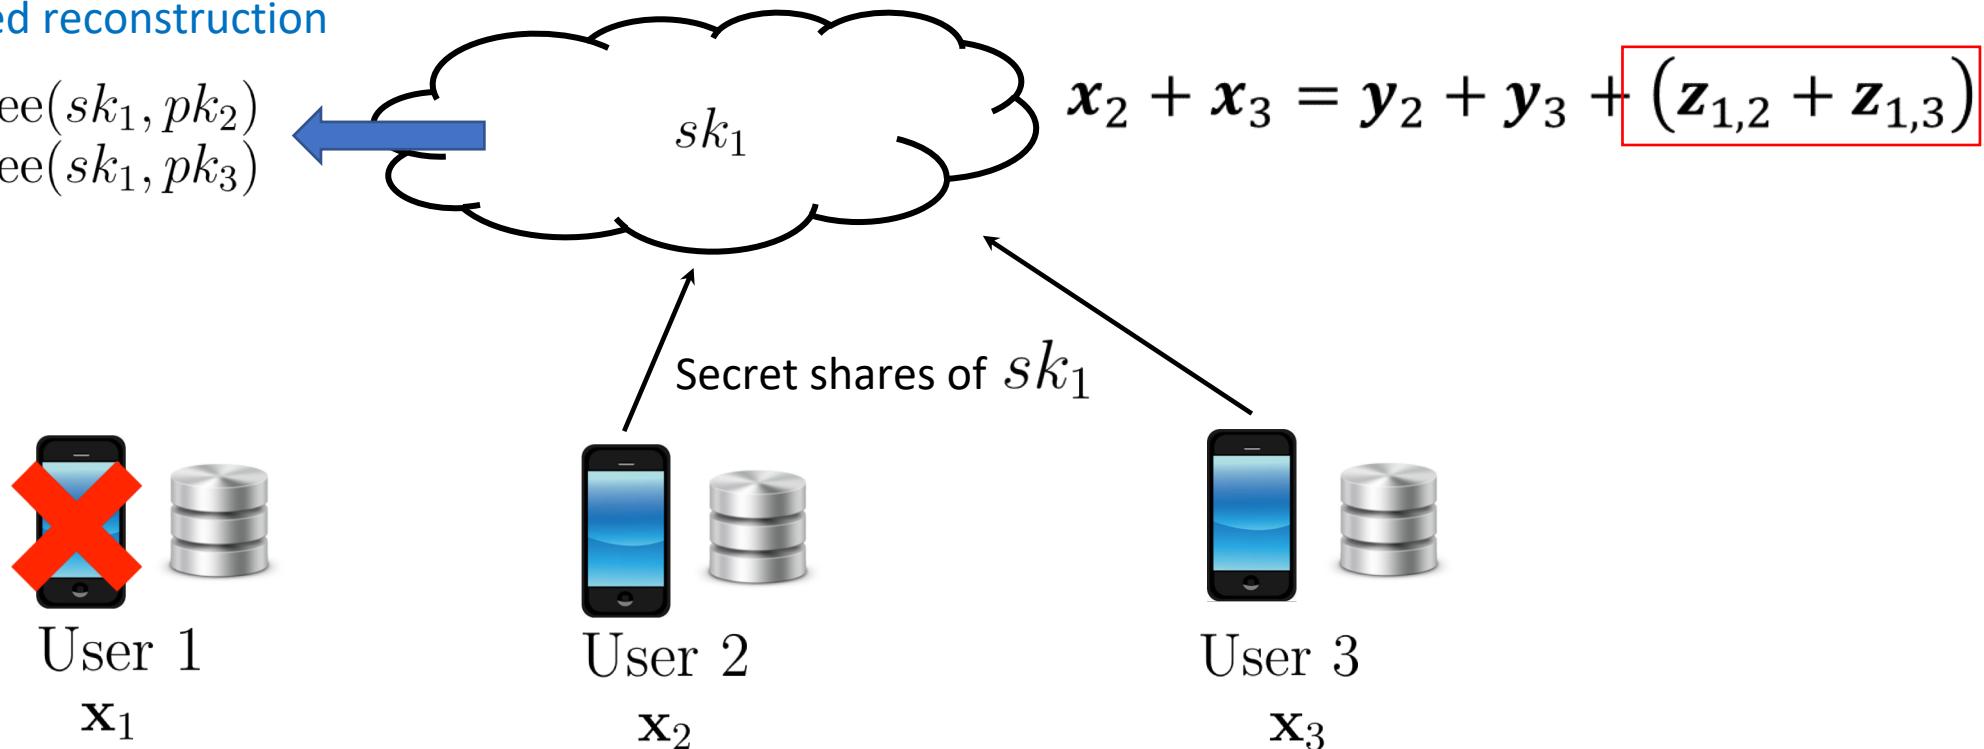
$$\mathbf{z}_{2,3} = \text{PRG}(s_{2,3})$$

$$\mathbf{z}_{2,3} = \text{PRG}(s_{2,3})$$

Example (SecAgg)

Pairwise random seed reconstruction

$$\begin{aligned}s_{1,2} &= \text{Key.Agree}(sk_1, pk_2) \\ s_{1,3} &= \text{Key.Agree}(sk_1, pk_3)\end{aligned}$$



$$z_{1,2} = \text{PRG}(s_{1,2})$$

$$z_{1,3} = \text{PRG}(s_{1,3})$$

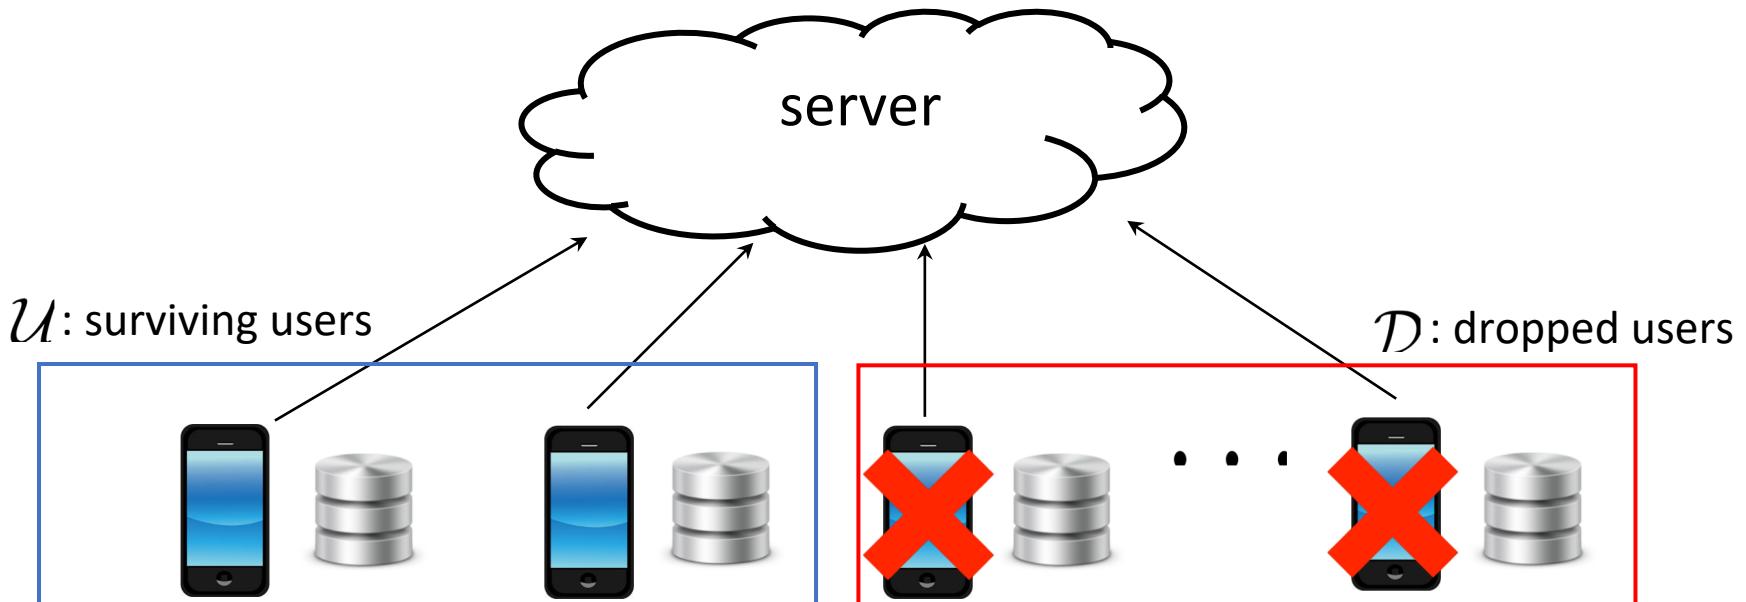
$$z_{1,2} = \text{PRG}(s_{1,2})$$

$$z_{2,3} = \text{PRG}(s_{2,3})$$

$$z_{1,3} = \text{PRG}(s_{1,3})$$

$$z_{2,3} = \text{PRG}(s_{2,3})$$

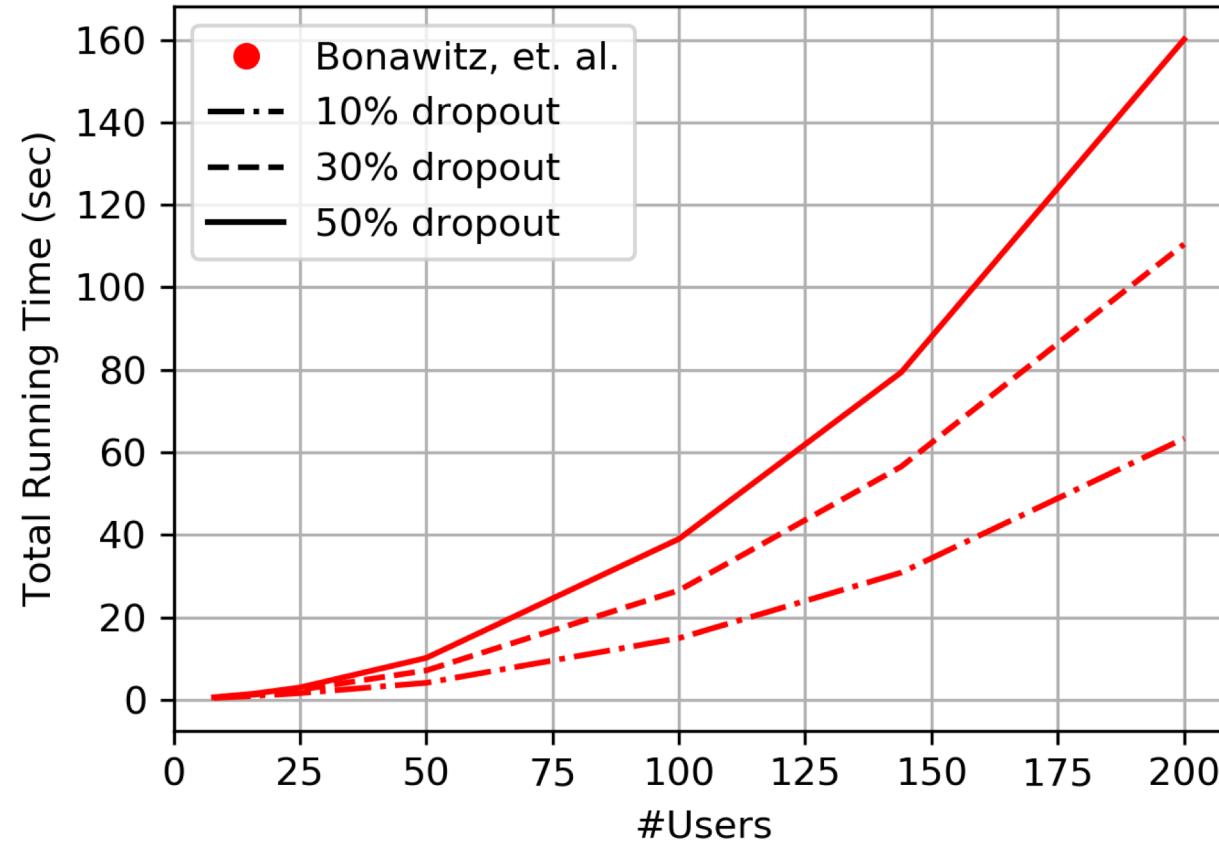
State-of-the-Art: SecAgg



$$\mathbf{z} = \sum_{u \in \mathcal{U}} (\mathbf{y}_u - \text{PRG}(b_u)) + \sum_{u \in \mathcal{D}} \left(\sum_{v: u < v} \text{PRG}(s_{u,v}) - \sum_{v: u > v} \text{PRG}(s_{v,u}) \right)$$

The number of **mask reconstructions** at the server substantially grows as more users are dropped, causing a major computational bottleneck.

State-of-the-Art: SecAgg



Individual model size of 100,000 with 32 bits entries- experiments over Amazon EC2

New Perspective

State-of-the-art:

$$\mathbf{z} = \sum_{u \in \mathcal{U}} (\mathbf{y}_u - \text{PRG}(b_u)) + \sum_{u \in \mathcal{D}} \left(\sum_{v: u < v} \text{PRG}(s_{u,v}) - \sum_{v: u > v} \text{PRG}(s_{v,u}) \right)$$



New Perspective:

$$\mathbf{z} = \sum_{u \in \mathcal{U}} \mathbf{x}_u + \sum_{u \in \mathcal{U}} \mathbf{z}_u$$

Our focus

We turn the focus from “random-seed reconstruction of the dropped users” to “one-shot aggregate-mask reconstruction of the surviving users”.

LightSecAgg: Key Components

Main components

1) Offline encoding and sharing of local masks

enables one-shot aggregate-mask reconstruction



2) Masking and uploading of local models

adds randomness that protects the local models

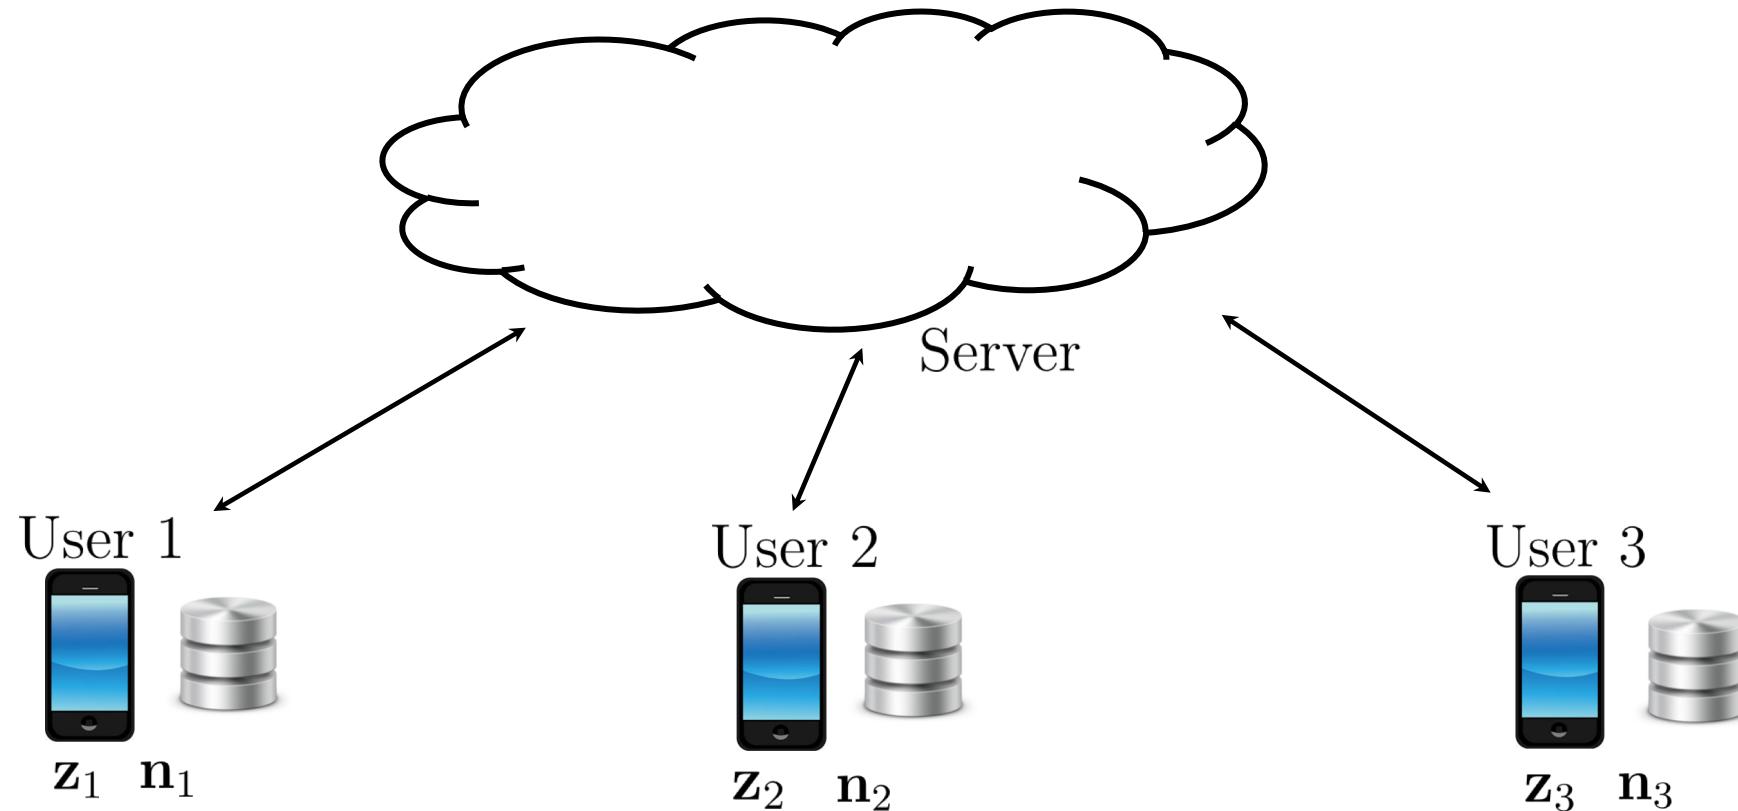


3) One-shot aggregate-model recovery

recovers the aggregate-model by canceling out the aggregate-mask

Example (LightSecAgg)

1) Offline **encoding** and **sharing** of local masks



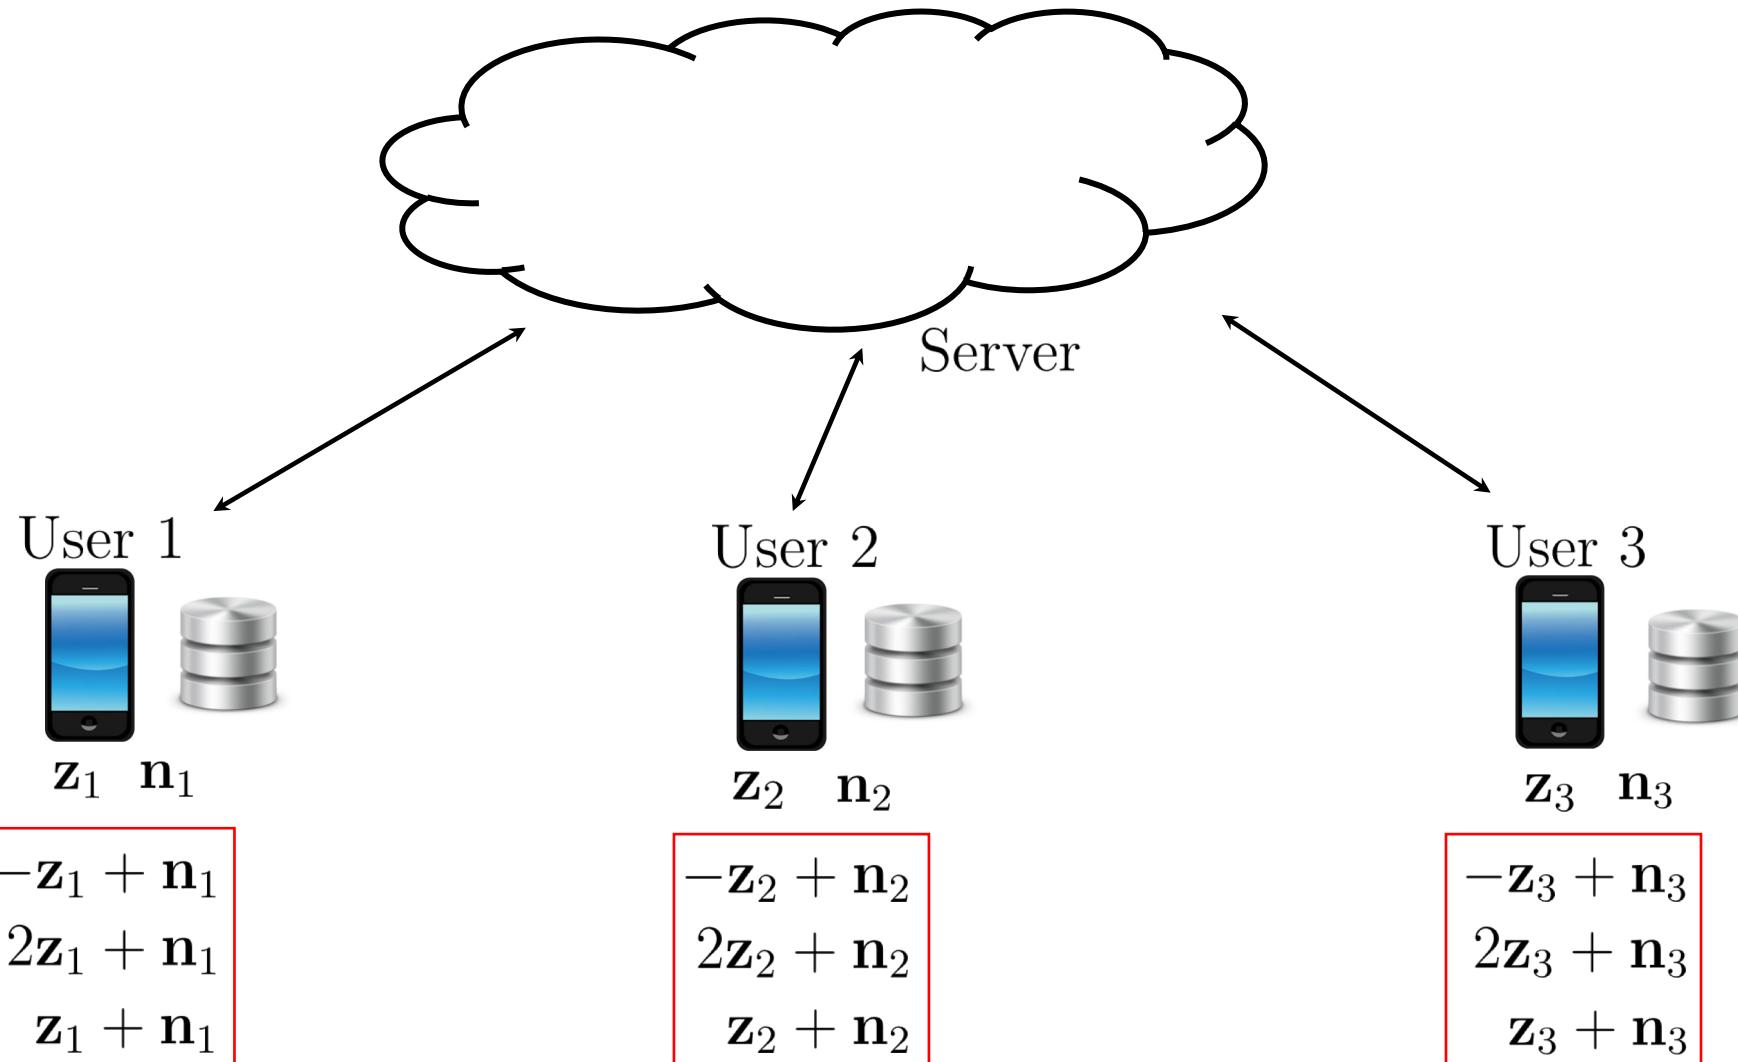
Encoding via MDS

$$\begin{array}{l} -z_1 + n_1 \\ 2z_1 + n_1 \\ z_1 + n_1 \end{array}$$

- 1) z_1 and n_1 can be recovered by any 2 out of 3 encoded masks.
=> robustness against dropped users
- 2) n_1 protects z_1 => privacy against 1 curious user

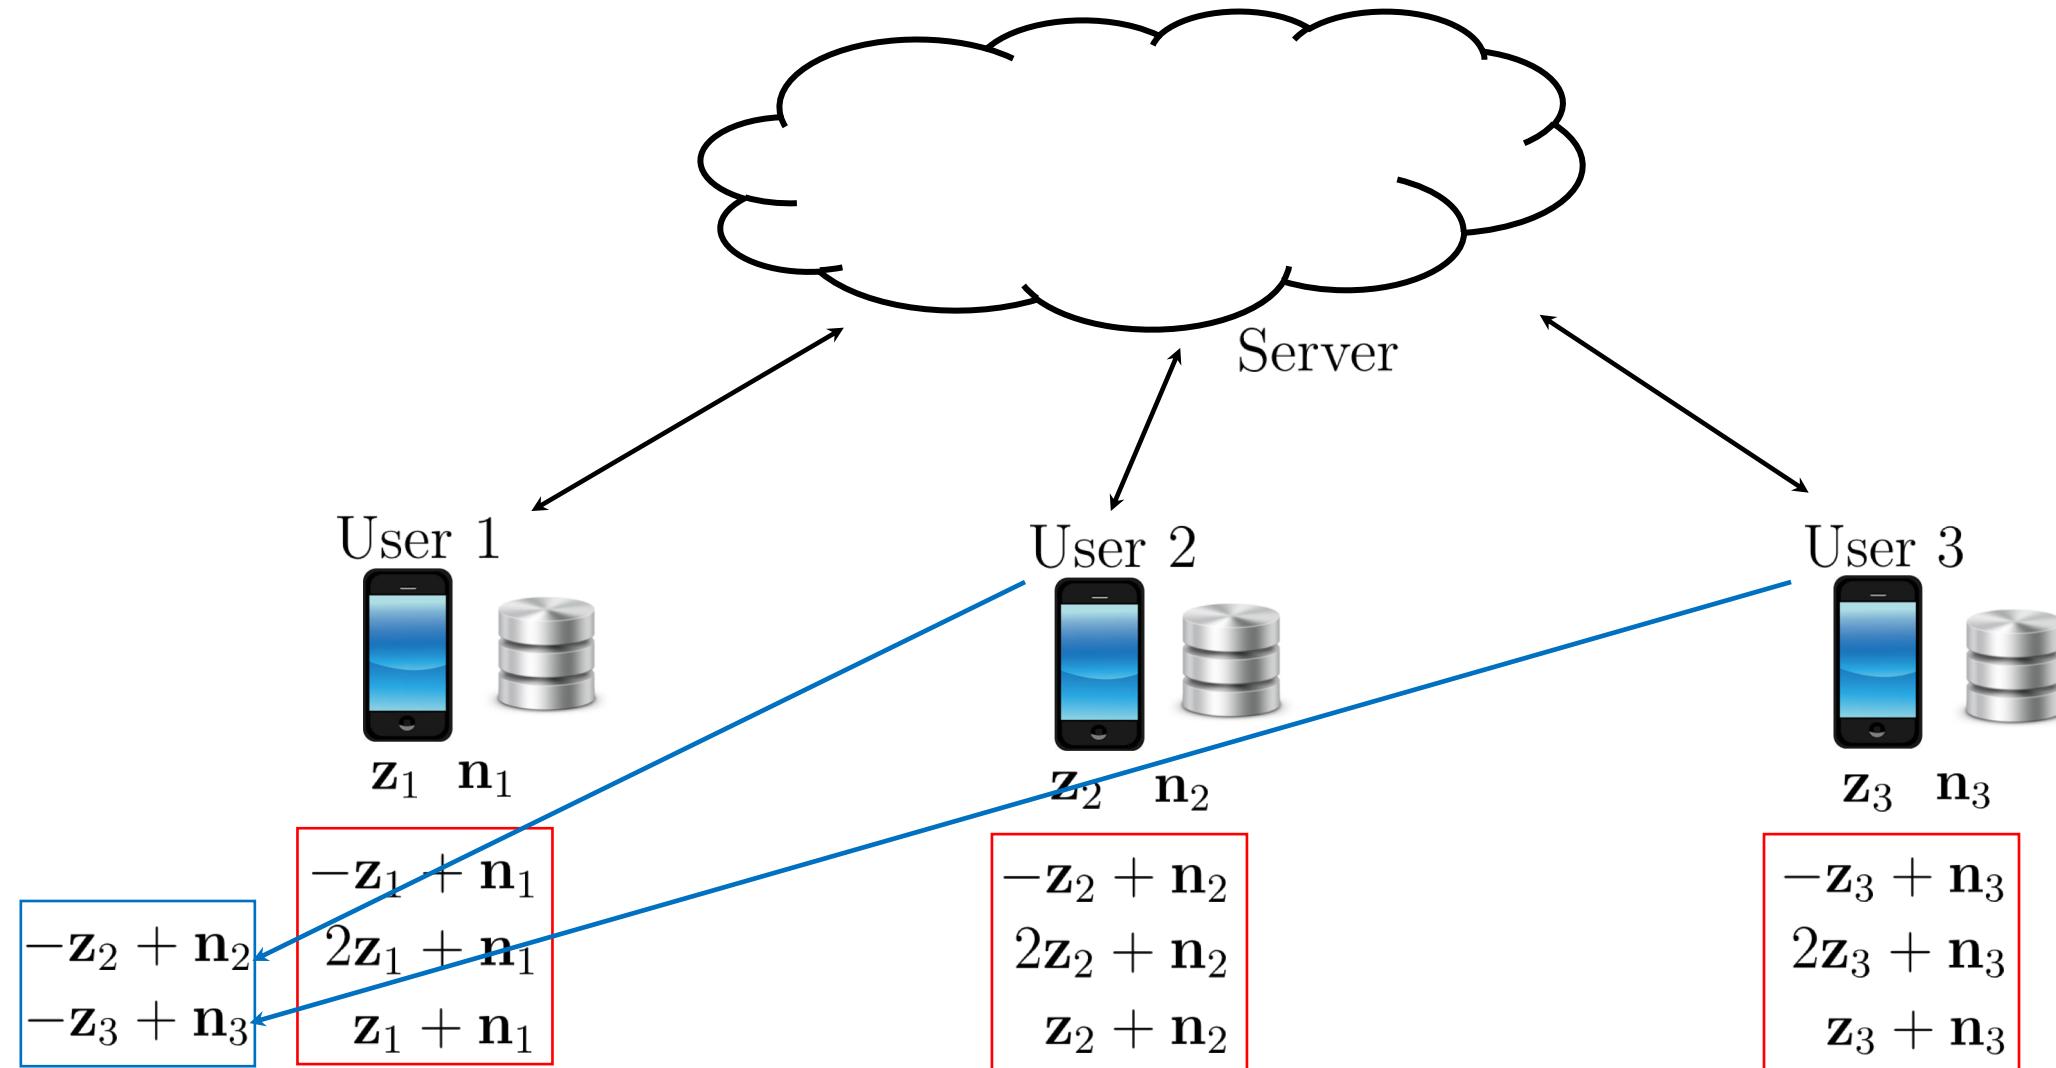
Example (LightSecAgg)

1) Offline **encoding** and **sharing** of local masks



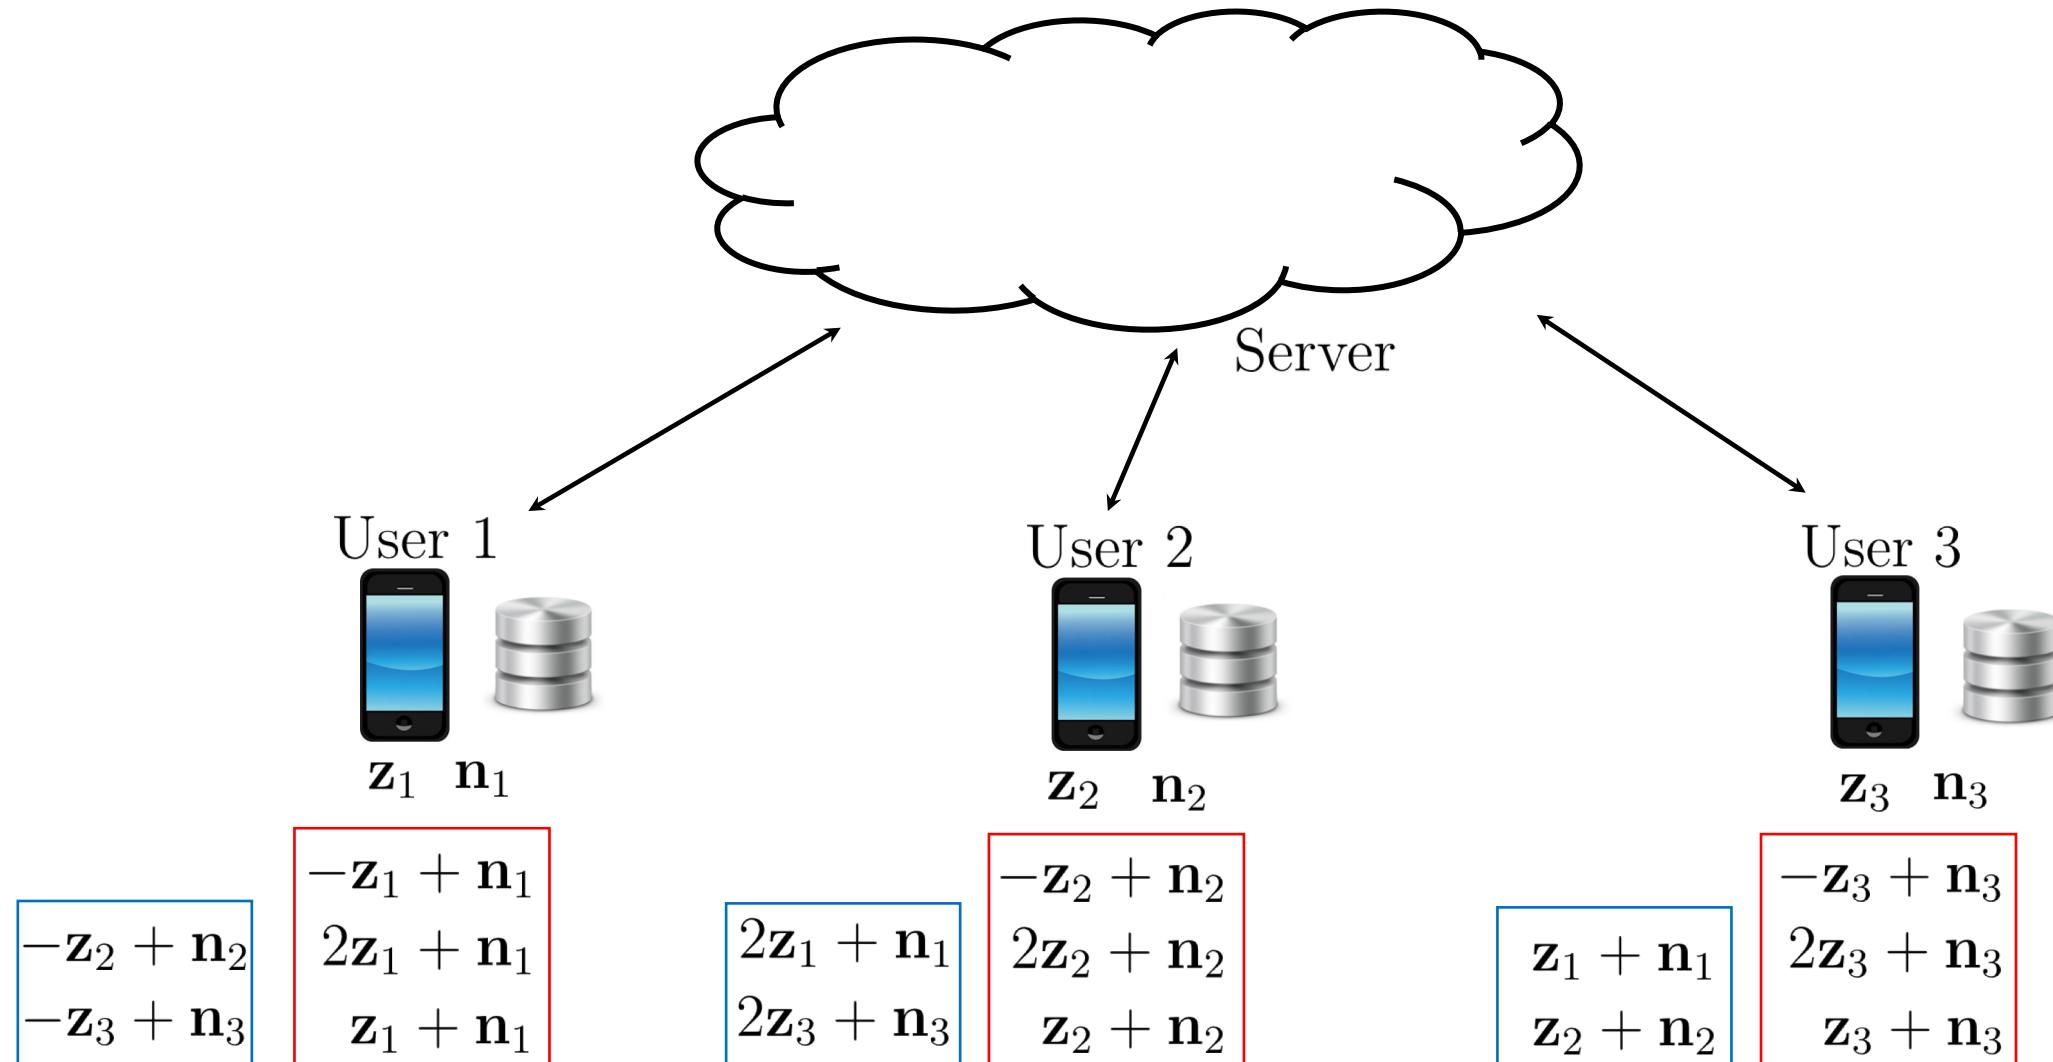
Example (LightSecAgg)

1) Offline **encoding** and **sharing** of local masks



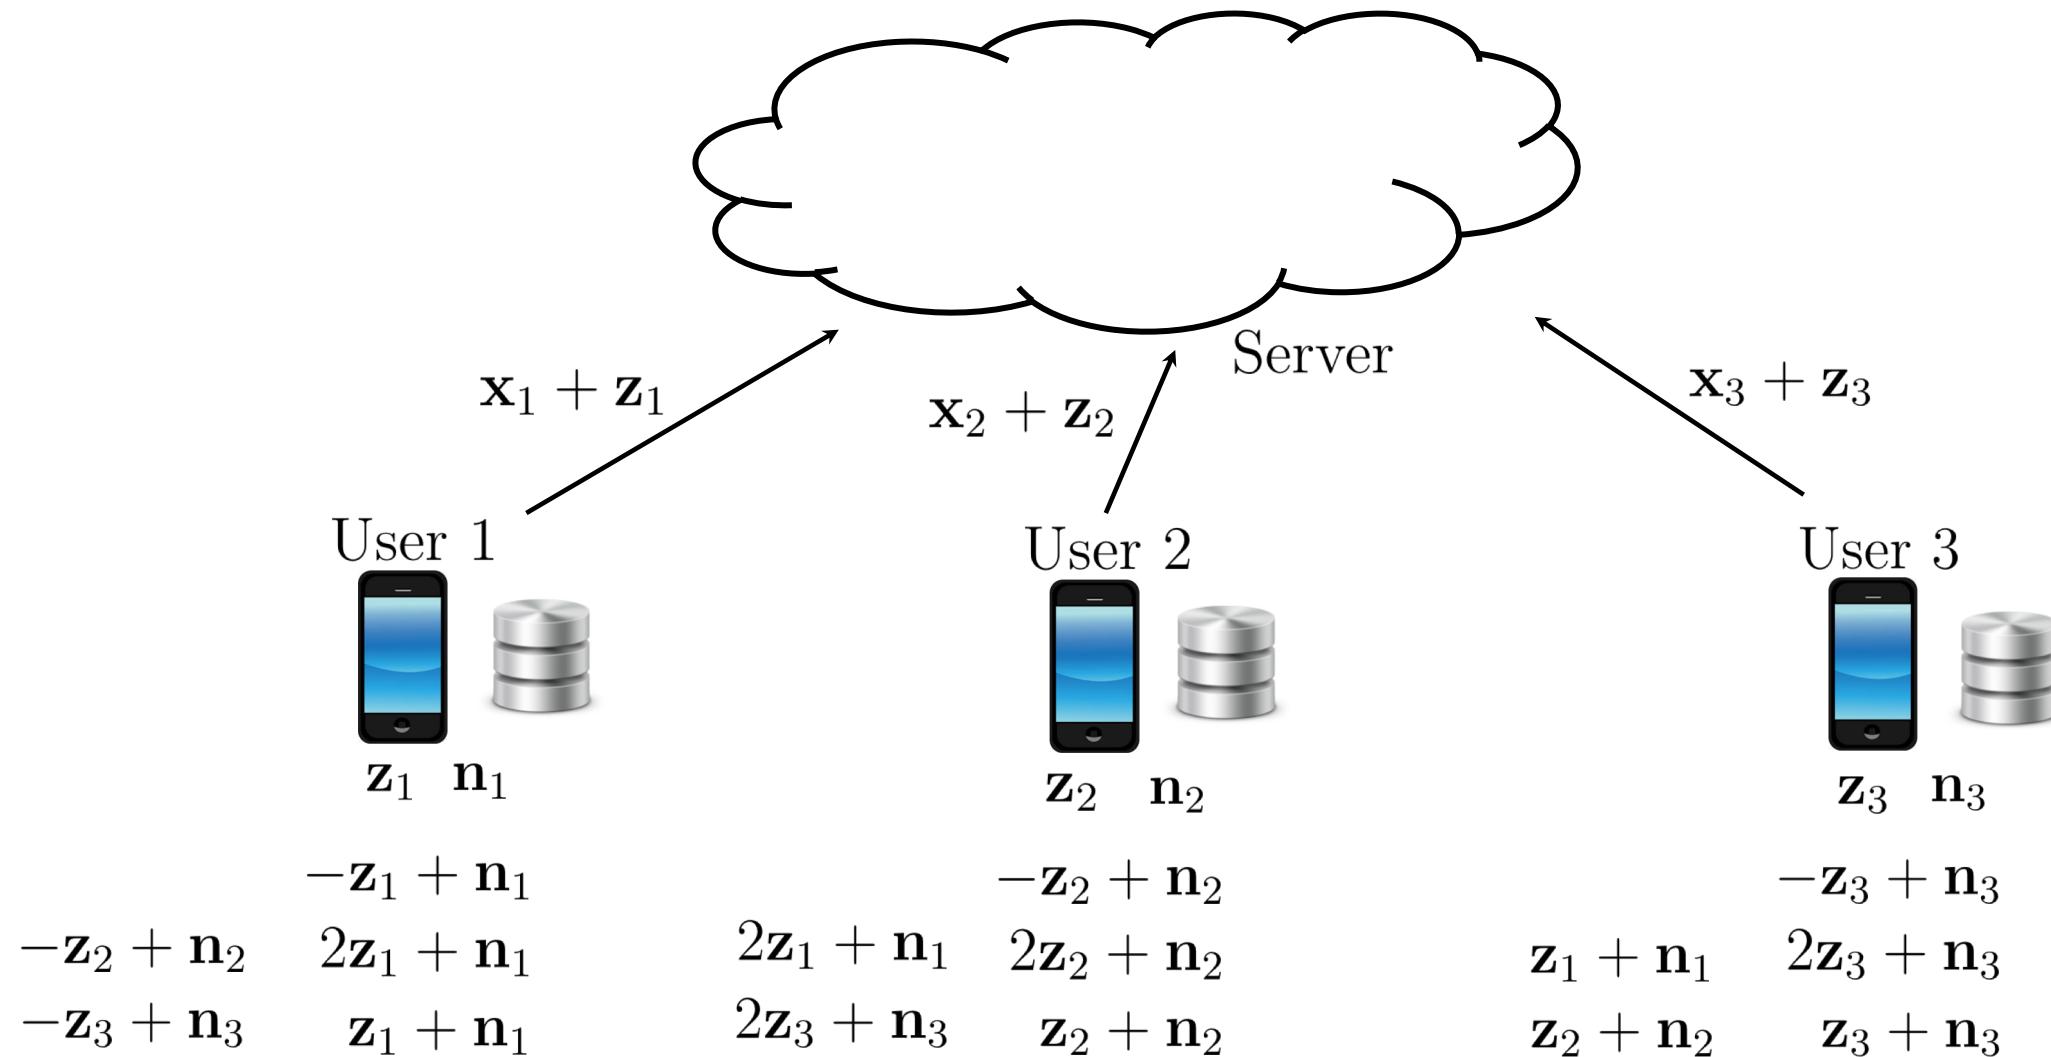
Example (LightSecAgg)

1) Offline **encoding** and **sharing** of local masks



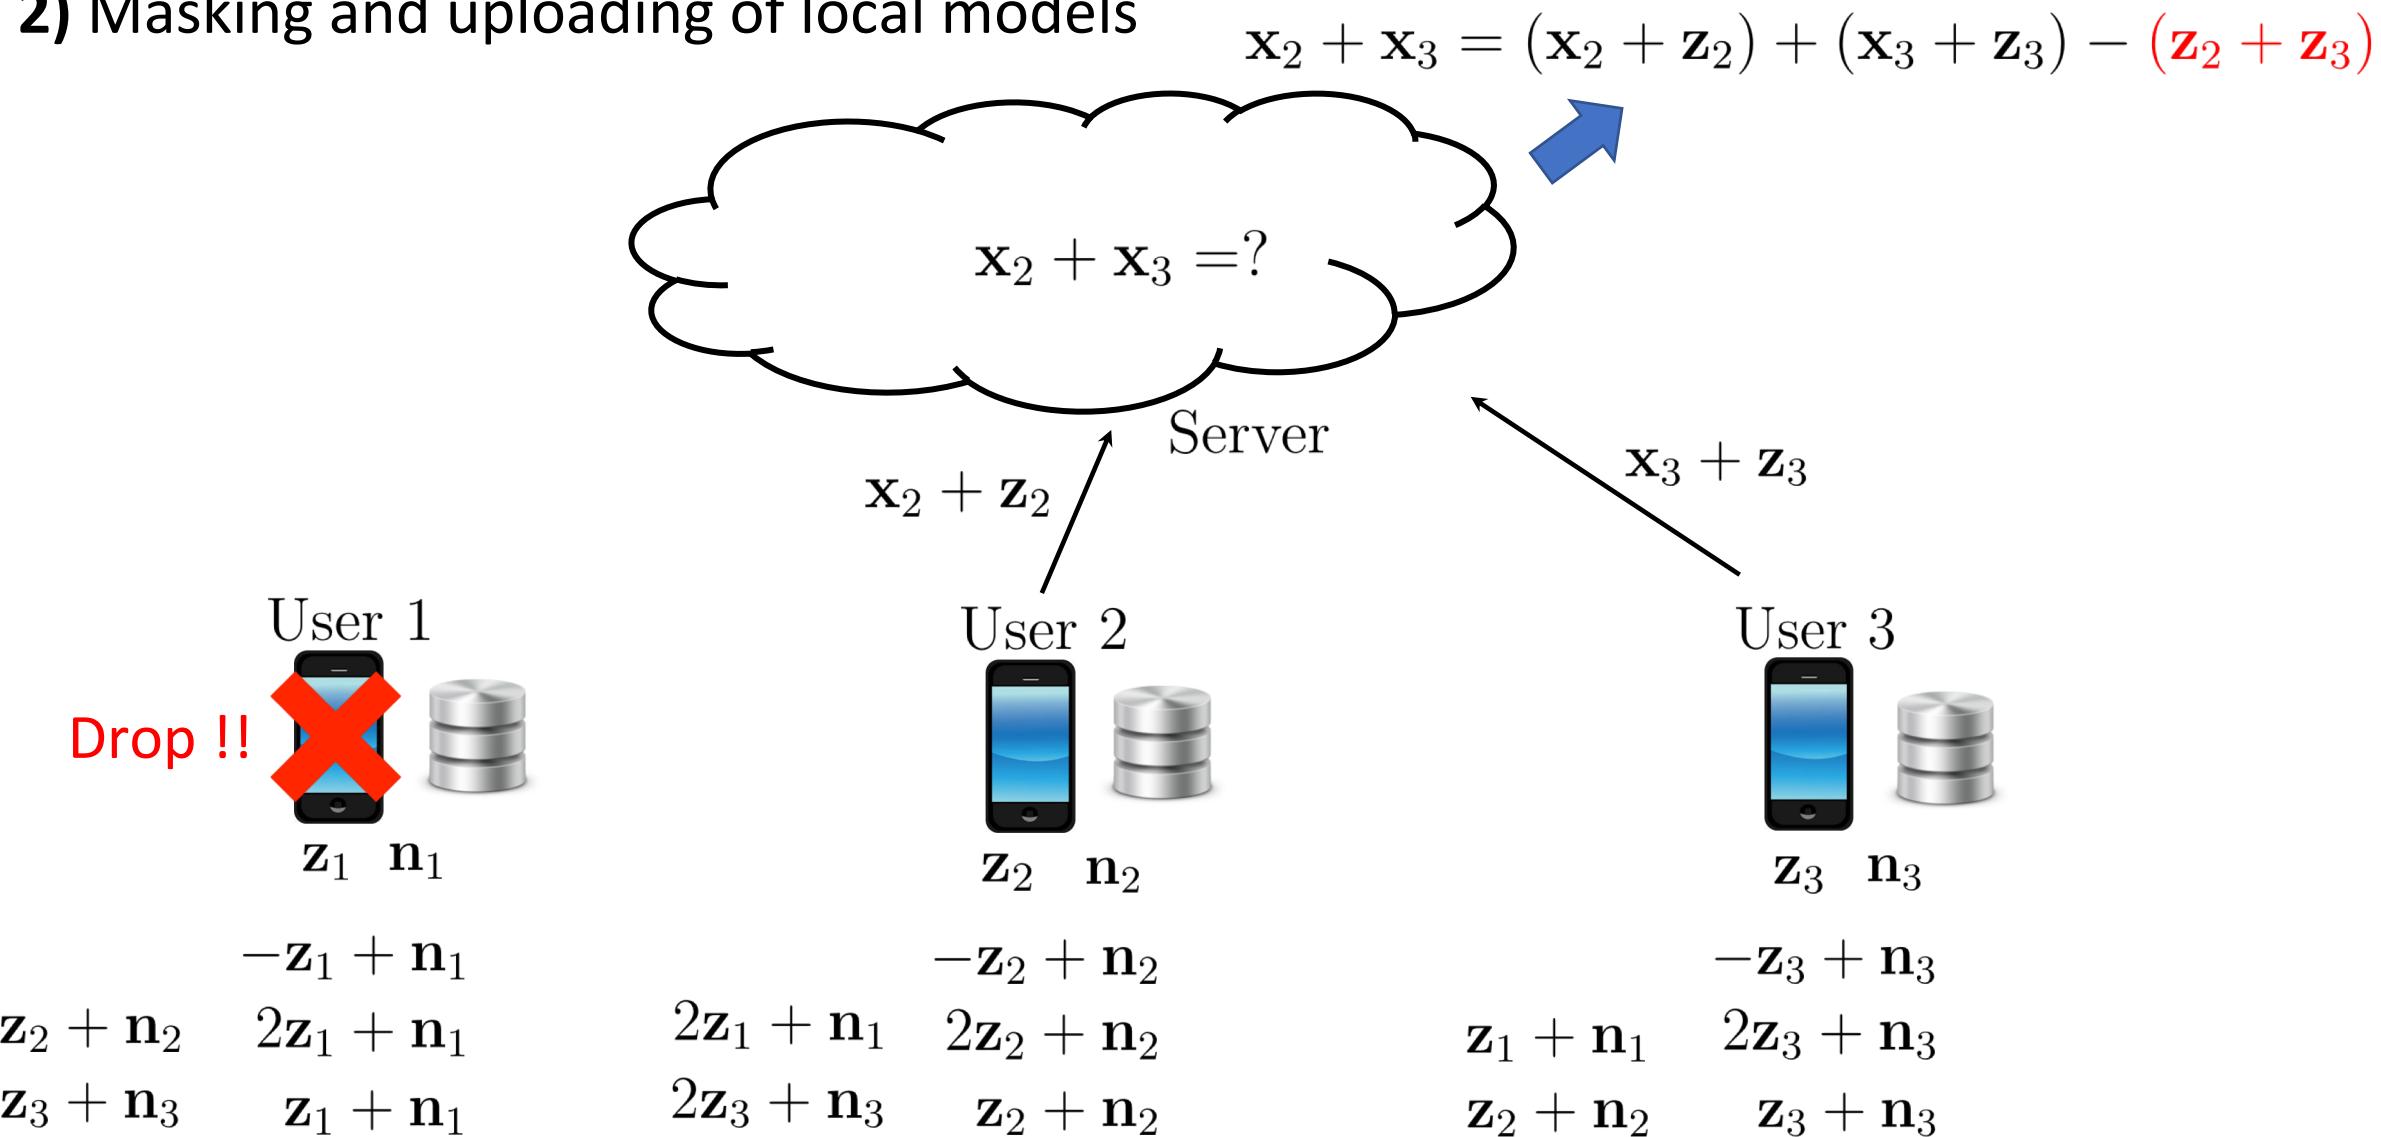
Example (LightSecAgg)

2) Masking and uploading of local models



Example (LightSecAgg)

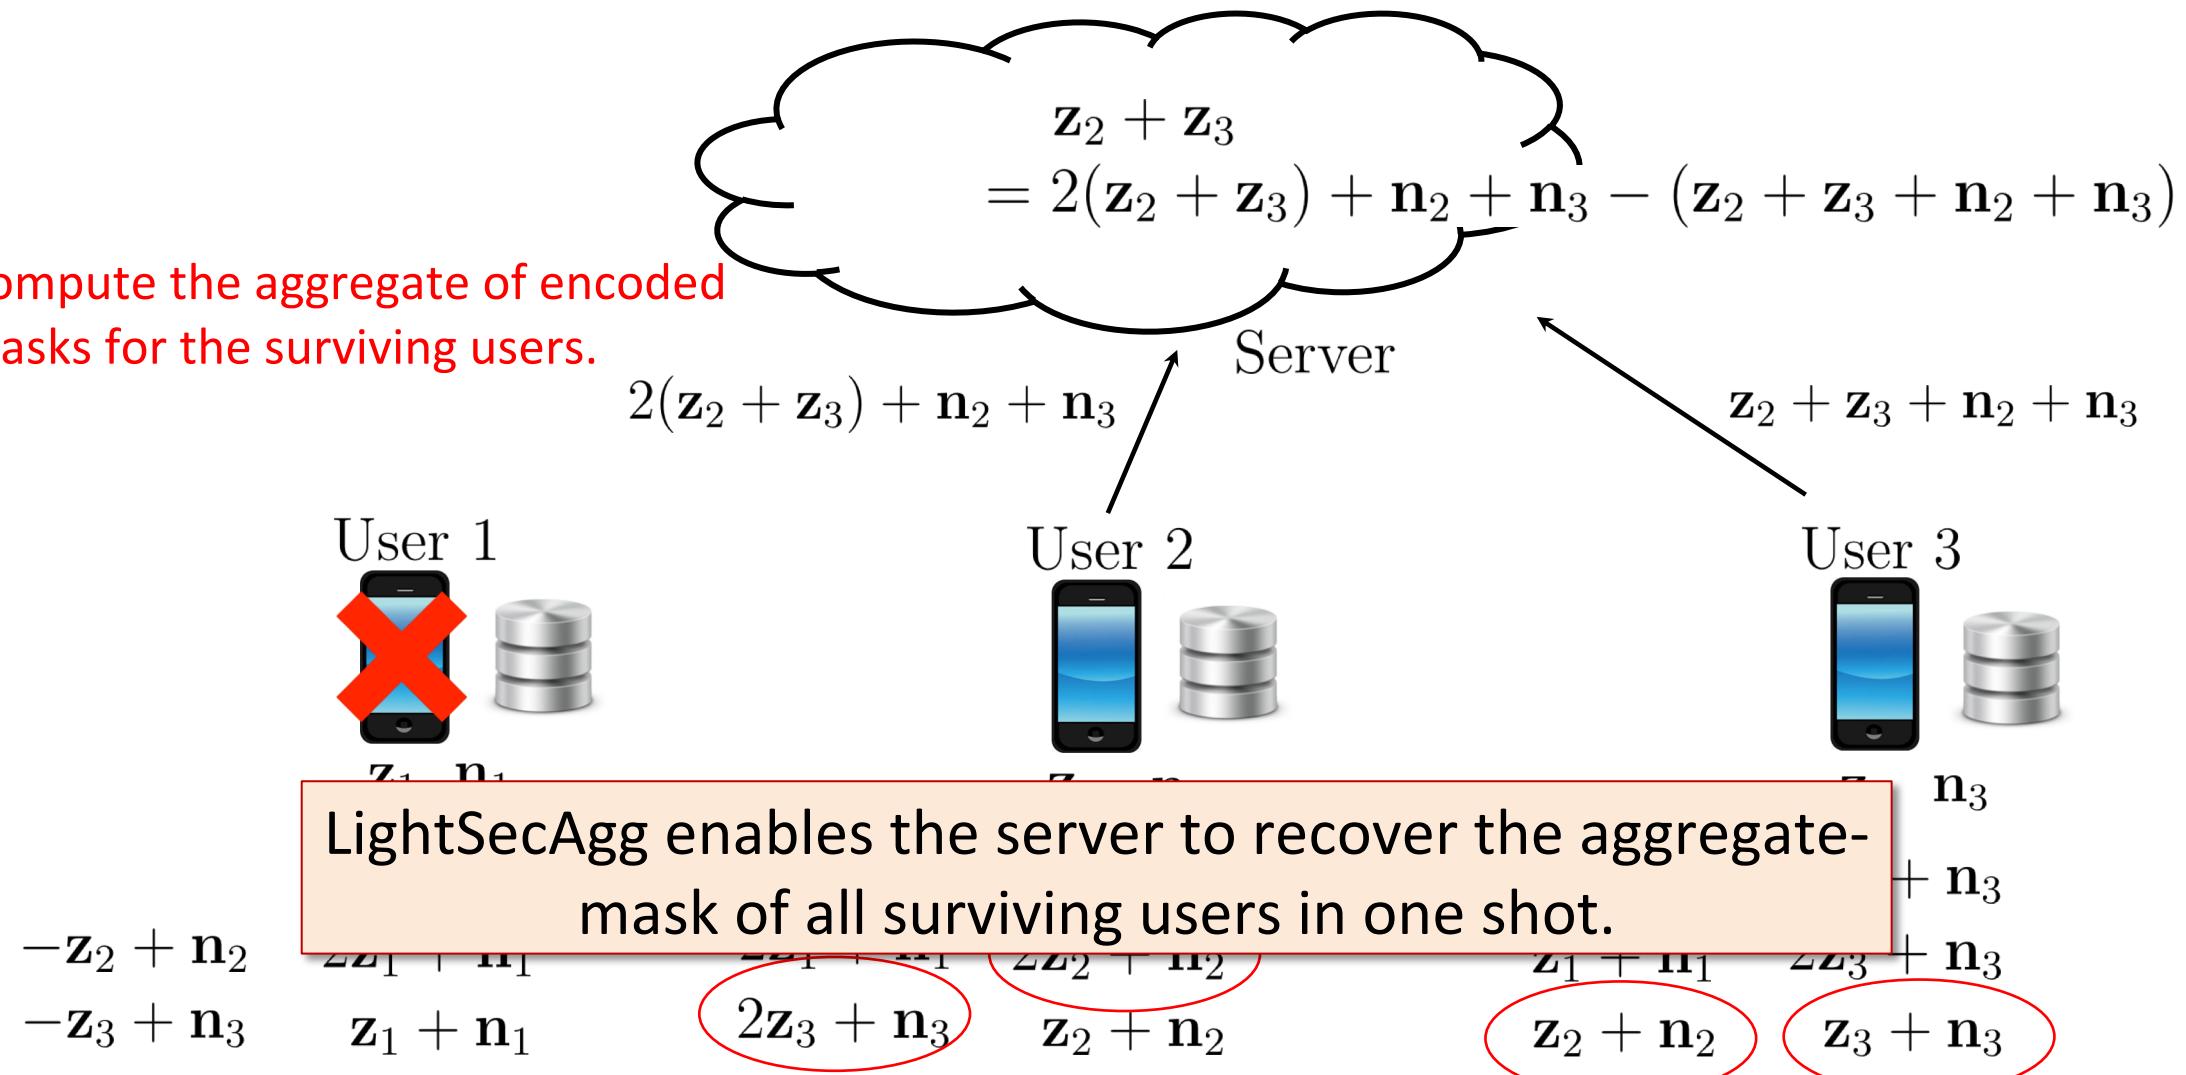
2) Masking and uploading of local models



Example (LightSecAgg)

3) One-shot aggregate-model recovery

Compute the aggregate of encoded masks for the surviving users.

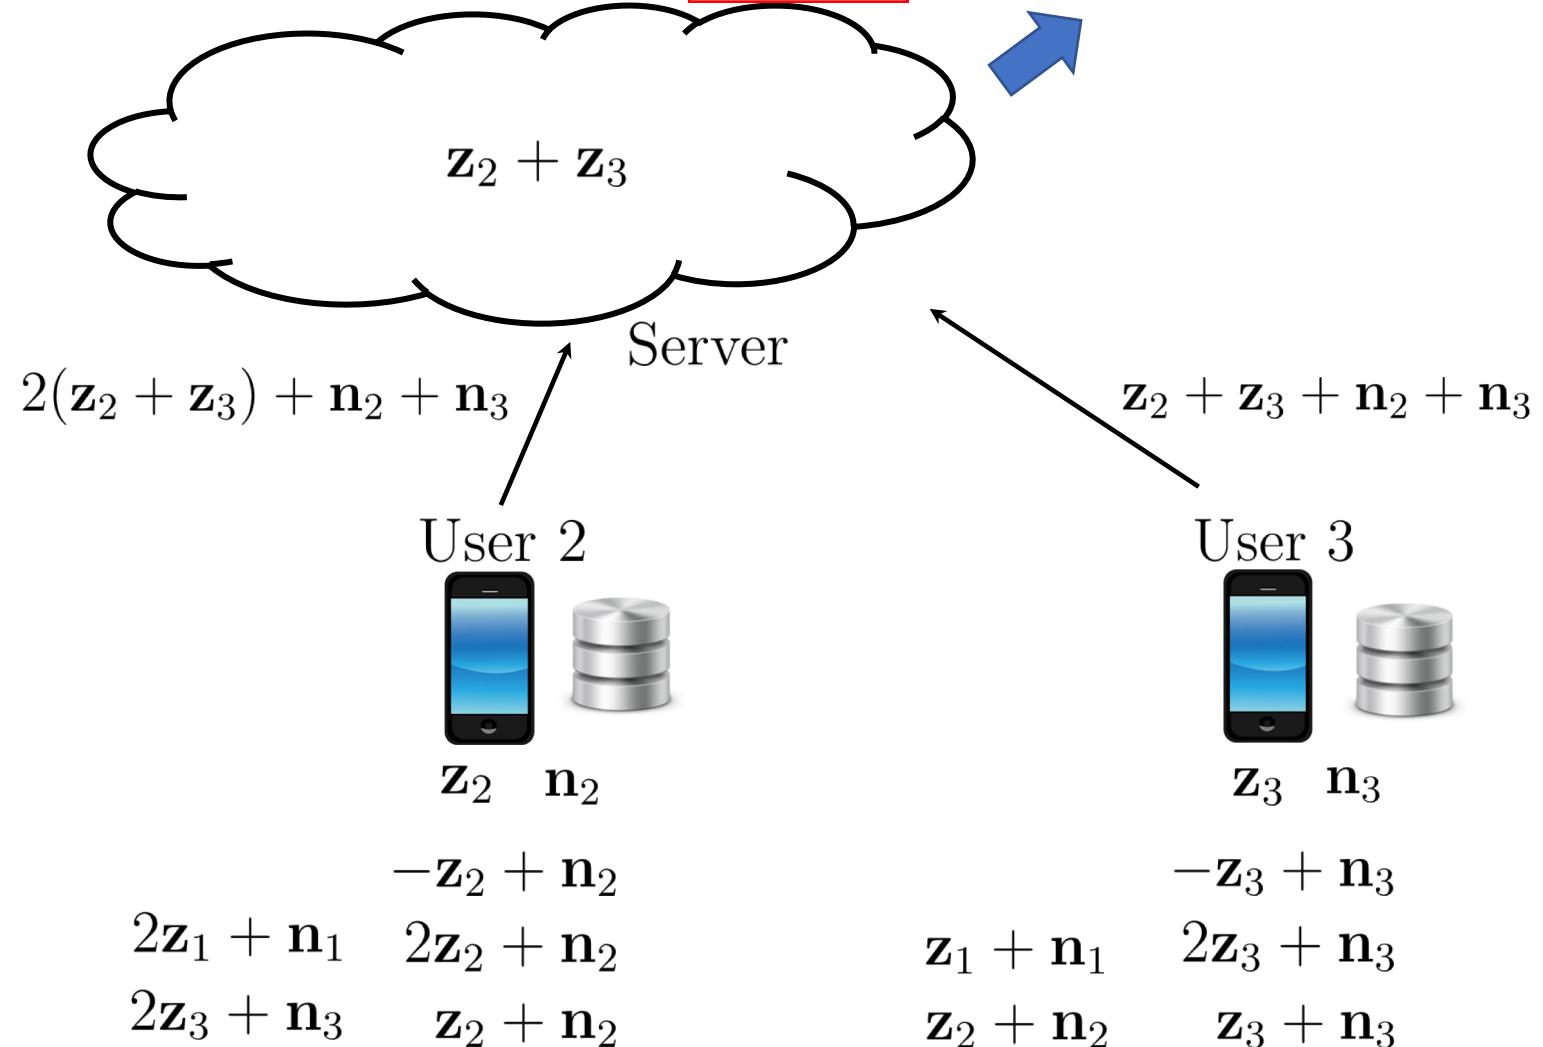


Example (LightSecAgg)

3) One-shot aggregate-model recovery

Aggregate-model

$$\boxed{x_2 + x_3} = (x_2 + z_2) + (x_3 + z_3) - (z_2 + z_3)$$



Phase 1: Offline Encoding and Sharing of Local Masks

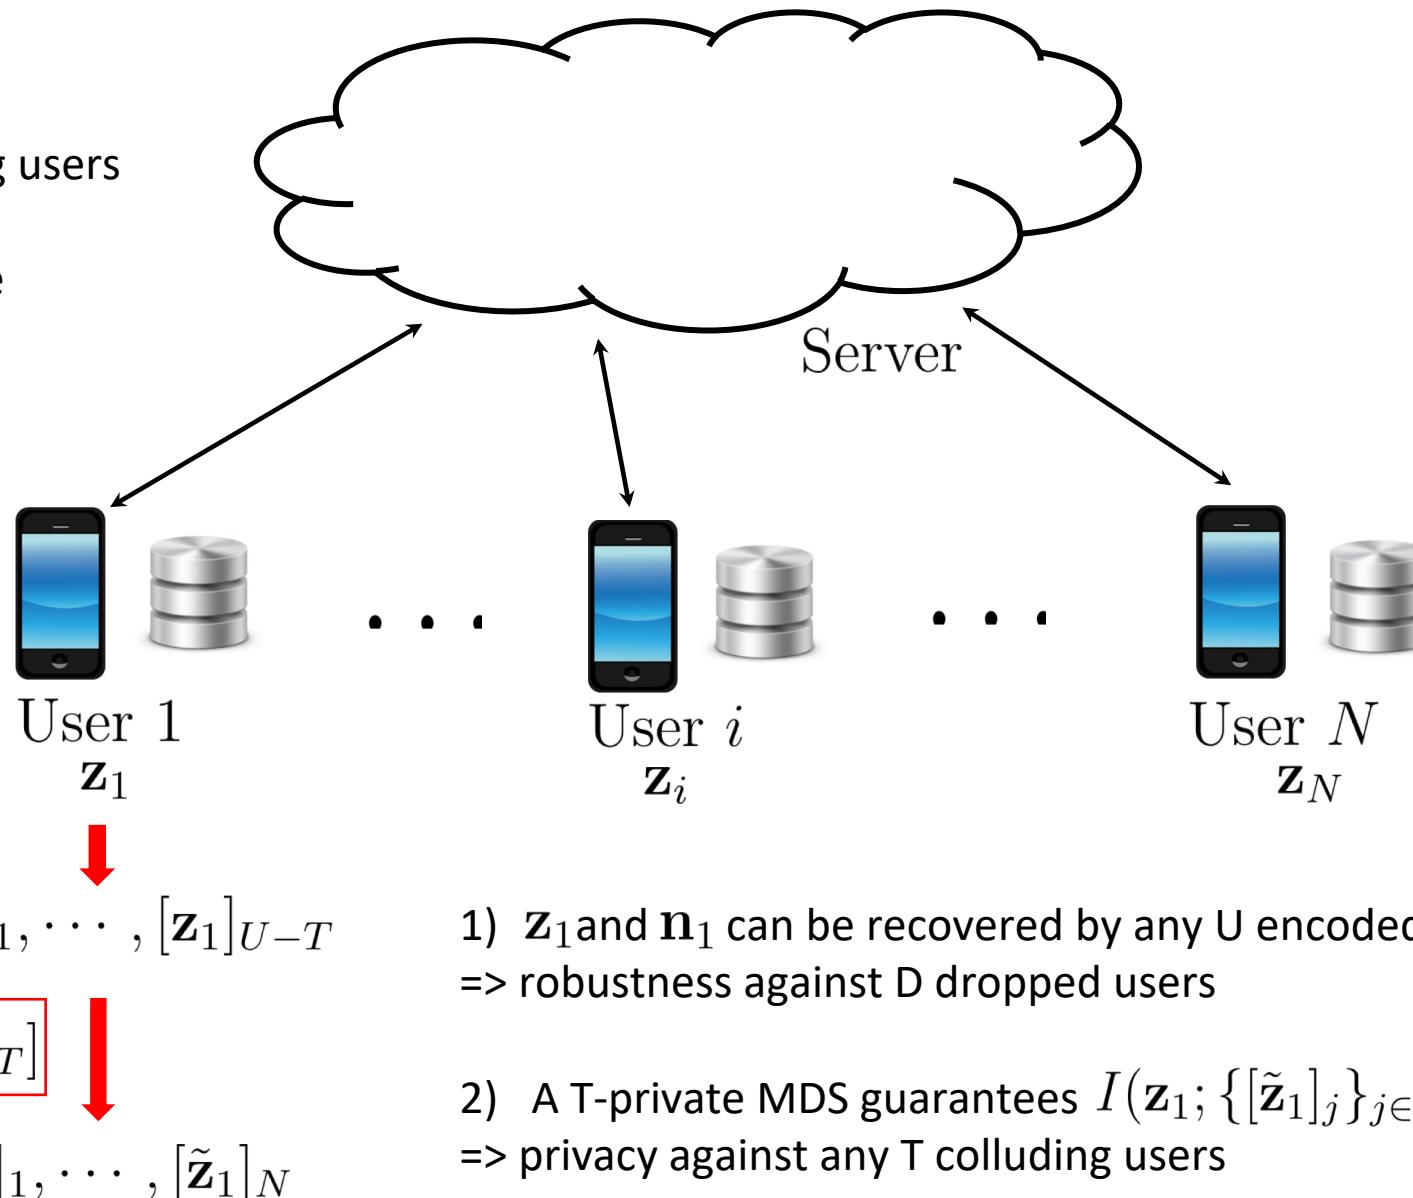
Design Parameters:

U: targeted number of surviving users

T: privacy guarantee

D: dropout-resiliency guarantee

$$N-D \geq U > T \geq 0$$



Local mask partitioning

$$[\mathbf{z}_1]_1, \dots, [\mathbf{z}_1]_{U-T}$$

$$\mathbf{n}_1 = [[\mathbf{n}_1]_1, \dots, [\mathbf{n}_1]_T]$$

1) \mathbf{z}_1 and \mathbf{n}_1 can be recovered by any U encoded masks.
=> robustness against D dropped users

2) A T-private MDS guarantees $I(\mathbf{z}_1; \{\tilde{\mathbf{z}}_1\}_{j \in \mathcal{T}}) = 0$
=> privacy against any T colluding users

Encoding via T-private MDS

$$[\tilde{\mathbf{z}}_1]_1, \dots, [\tilde{\mathbf{z}}_1]_N$$

Phase 1: Offline Encoding and Sharing of Local Masks

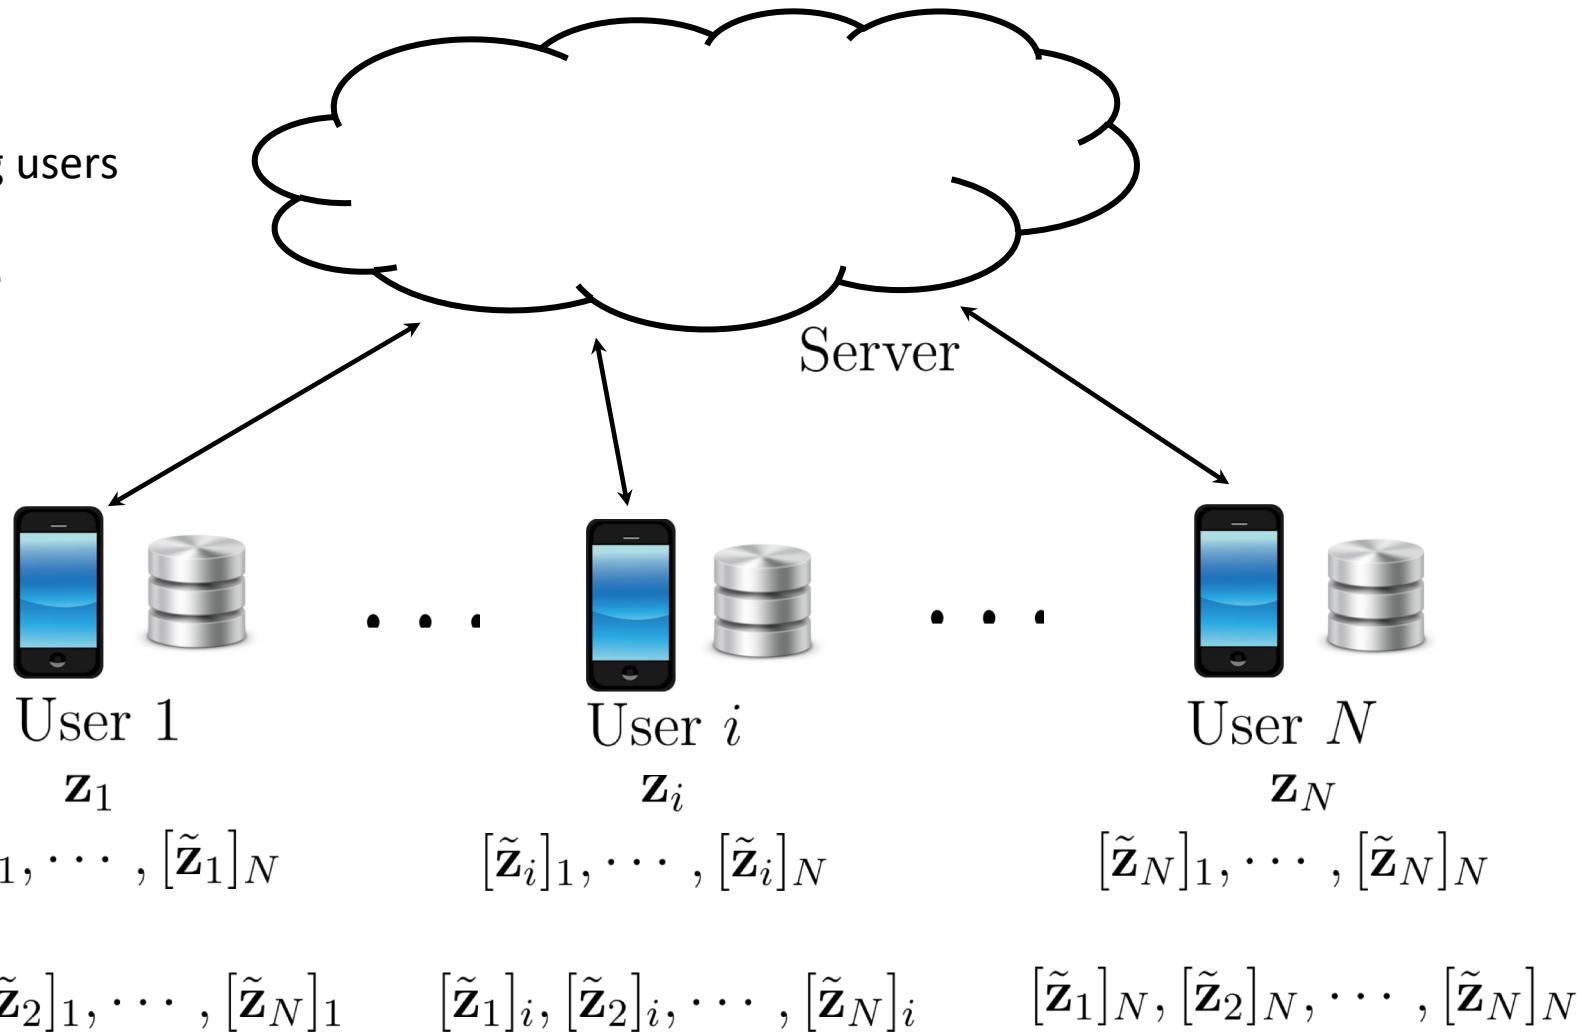
Design Parameters:

U: targeted number of surviving users

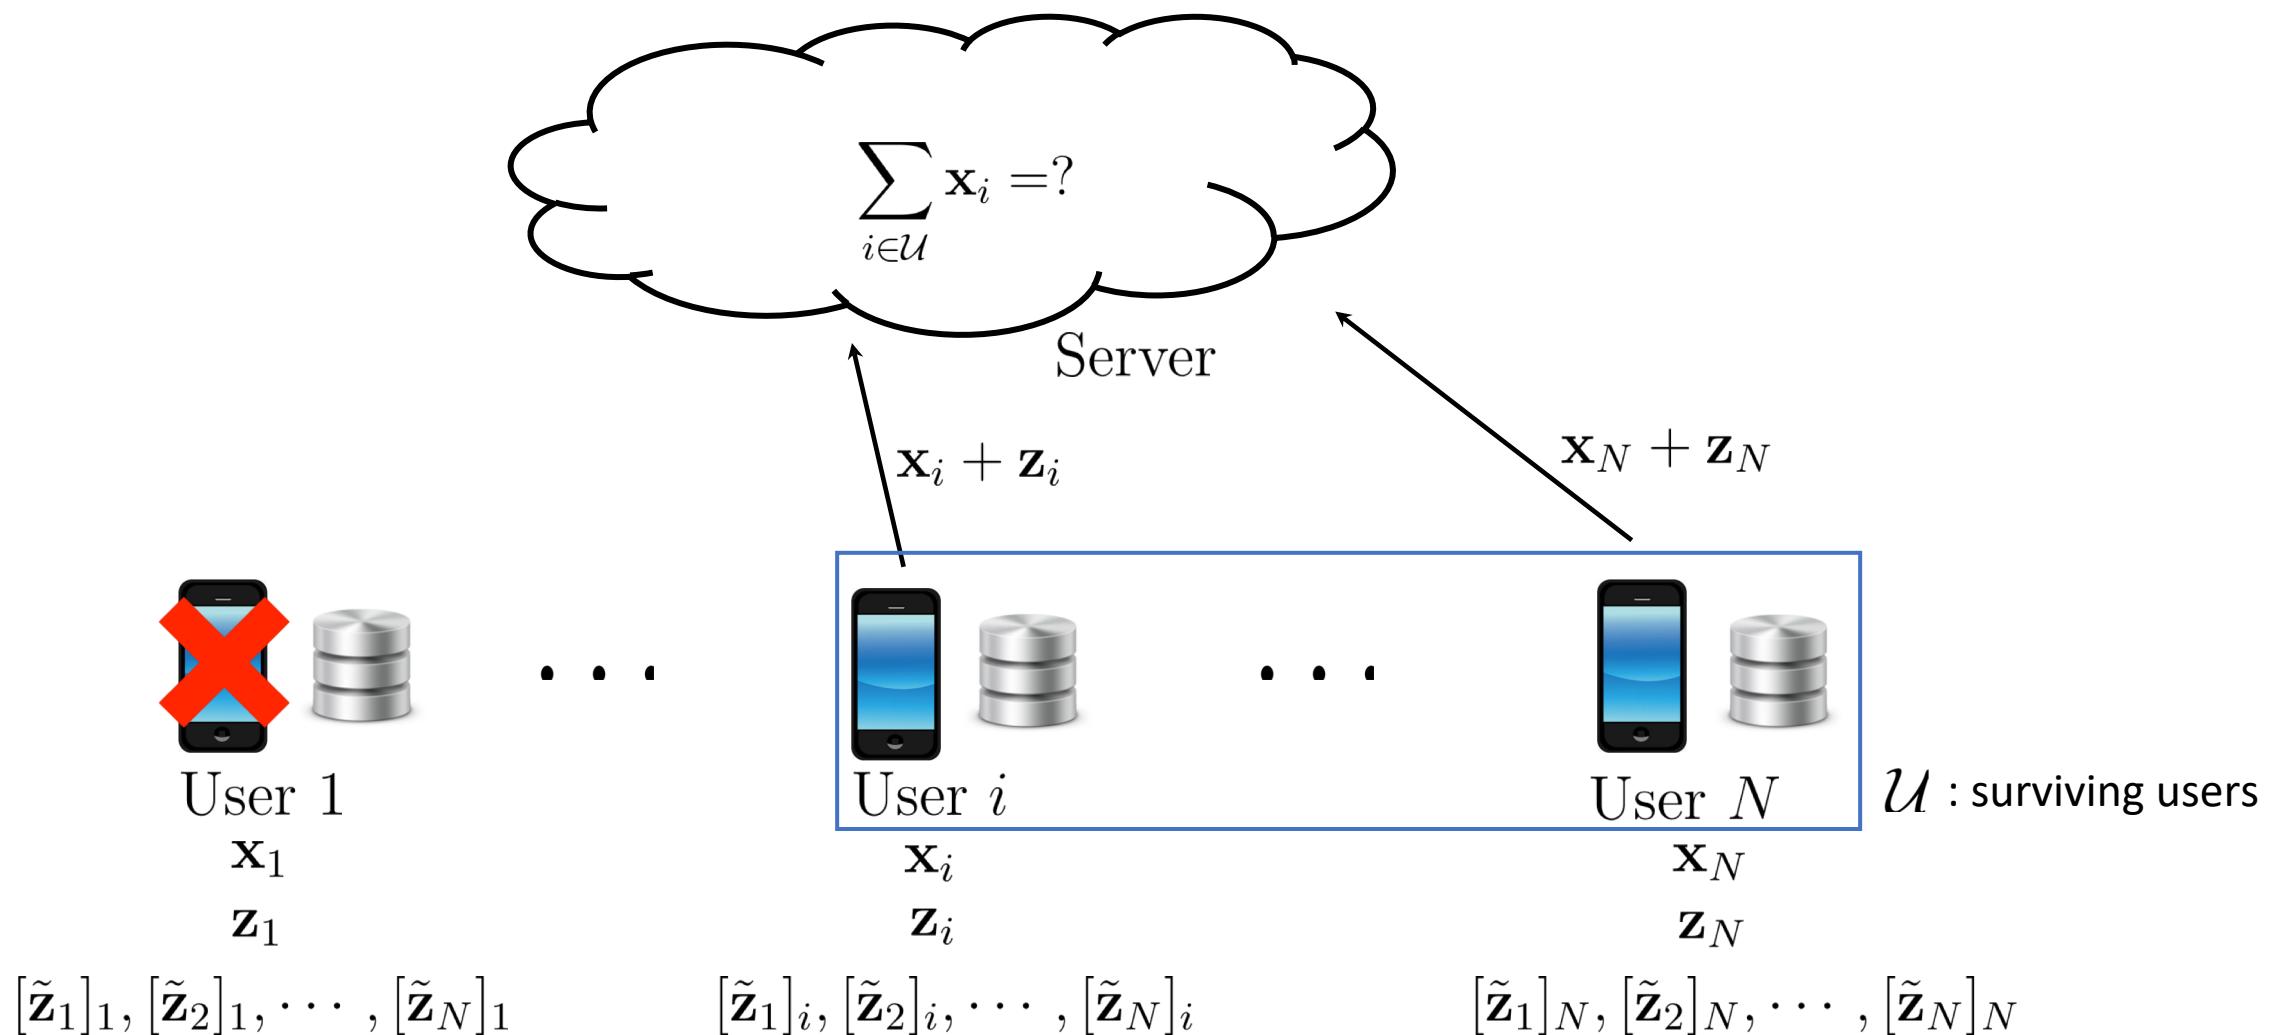
T: privacy guarantee

D: dropout-resiliency guarantee

$$N - D \geq U > T \geq 0$$

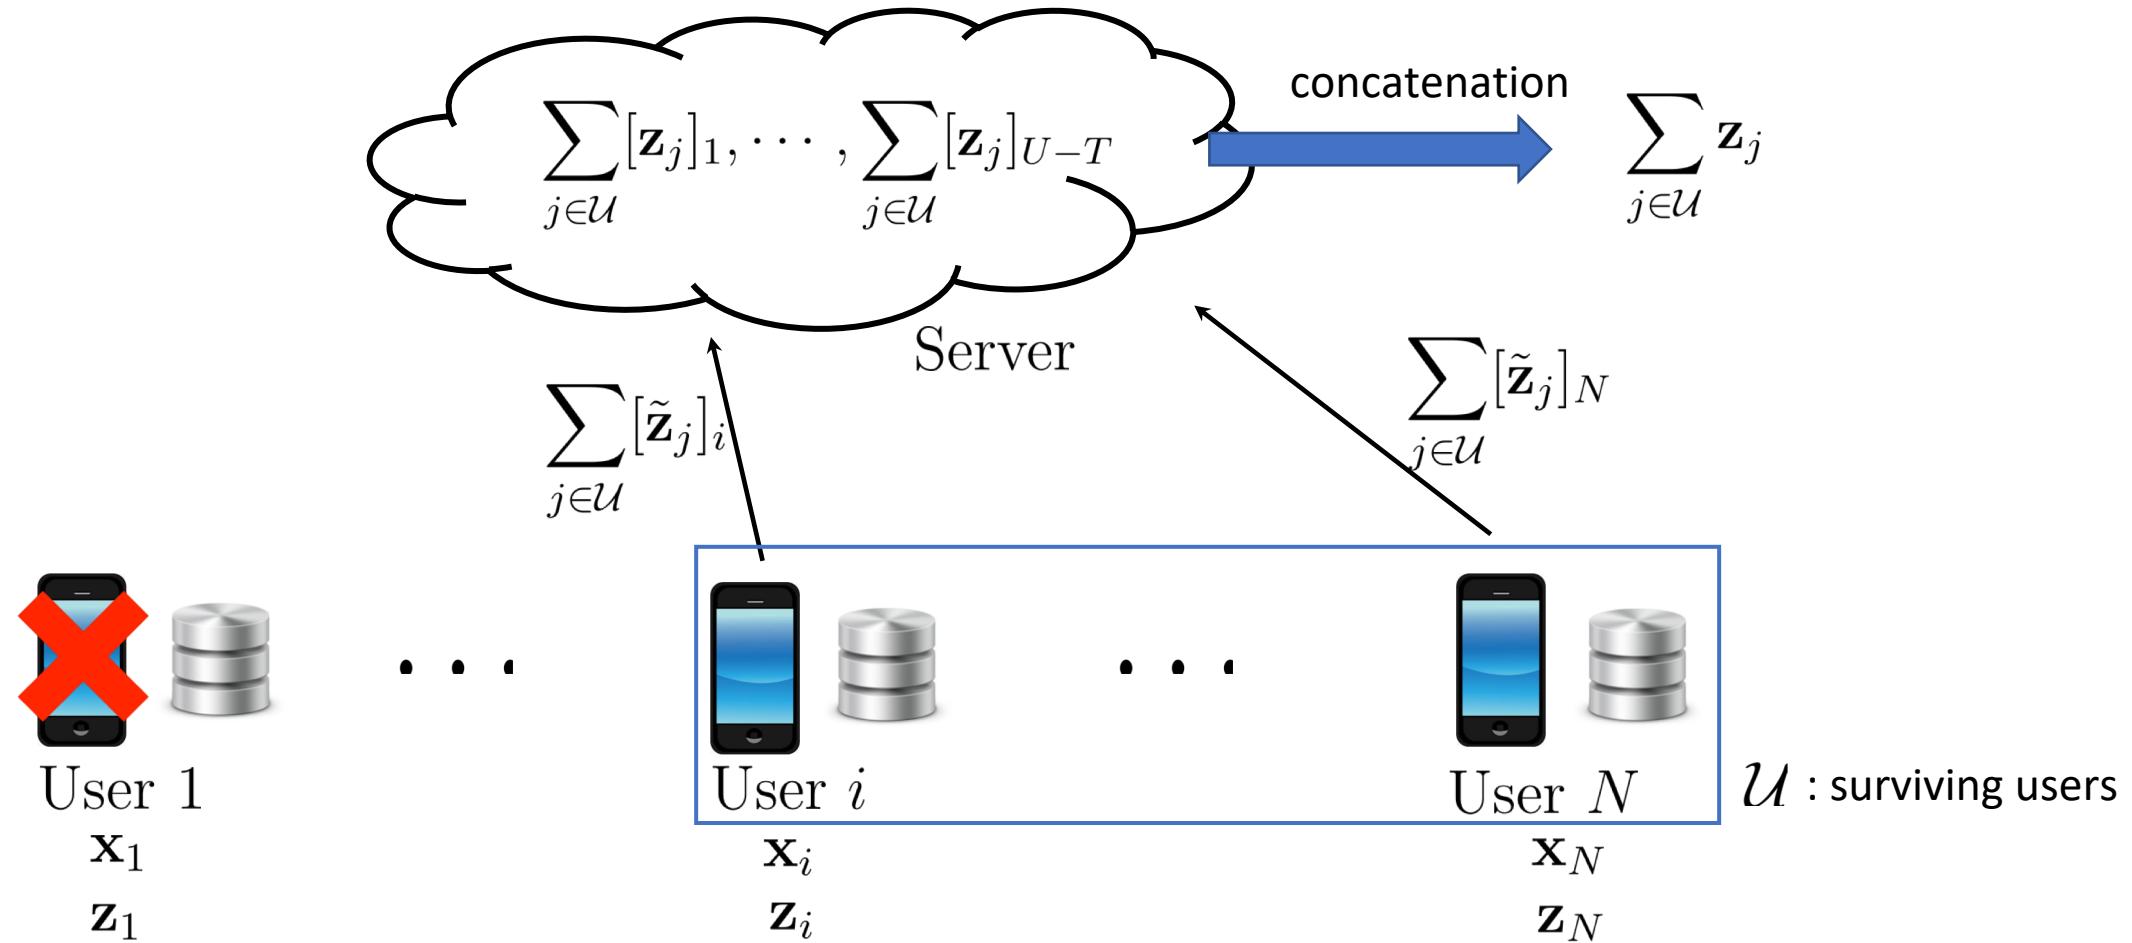


Phase 2: Masking and Uploading of Local Models



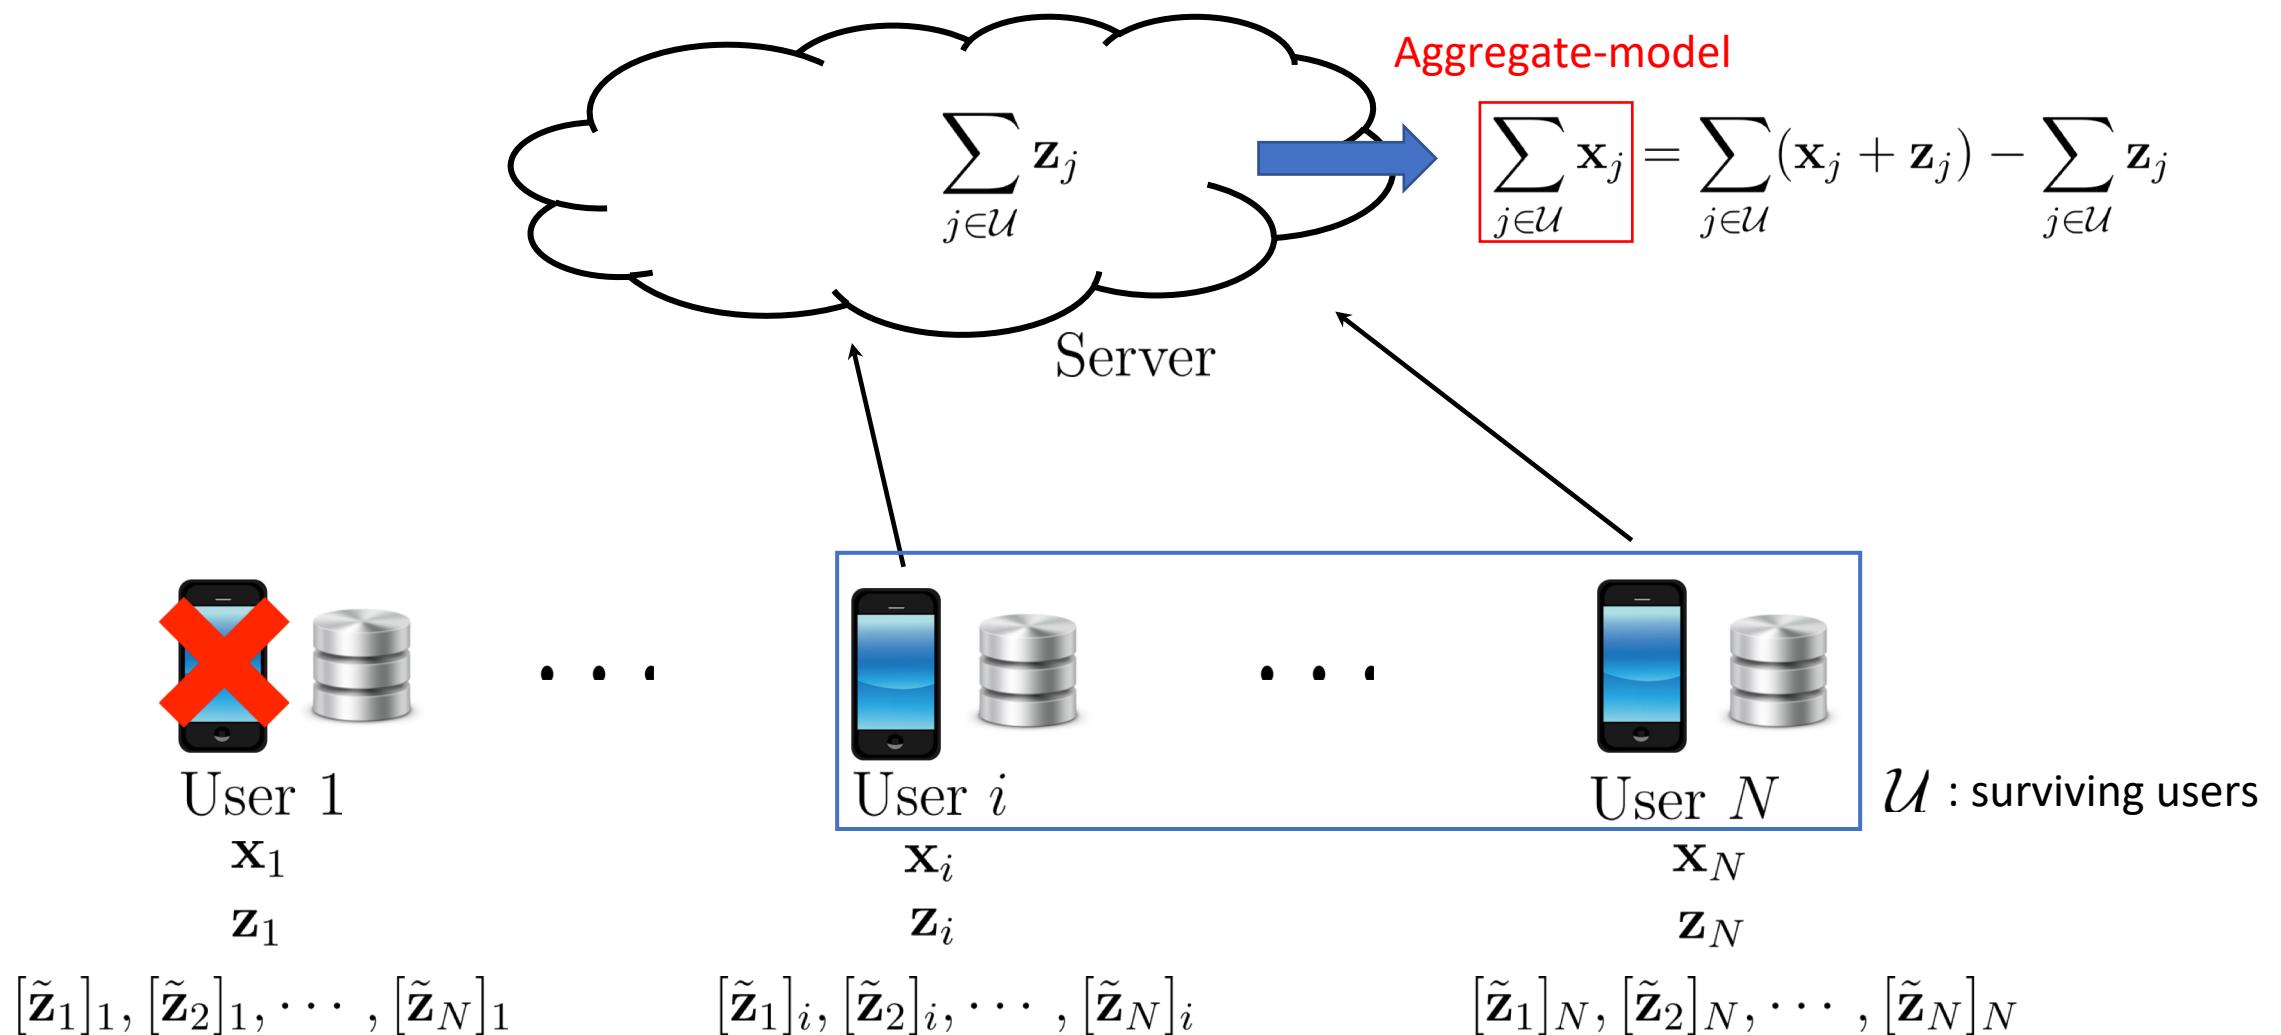
Phase 3: One-Shot Aggregate-Model Recovery

The server recovers the aggregate-mask after receiving any **U** messages via MDS decoding.



$[\tilde{\mathbf{z}}_1]_1, \dots, [\tilde{\mathbf{z}}_N]_N$: The server only needs to reconstruct **one mask** in the recovery phase, independent of the number of dropped users.

Phase 3: One-Shot Aggregate-Model Recovery



Theoretical Guarantees

Theorem: *LightSecAgg can simultaneously achieve*

- (1) *privacy guarantee against up to any T colluding users, and*
- (2) *dropout-resiliency guarantee against up to any D dropped users,*

for any pair of privacy guarantee T and dropout-resiliency guarantee D such that $T + D < N$.

Theoretical Guarantees

- Complexity comparison between SecAgg, SecAgg+ and LightSecAgg:
 - d: model size.
 - s: length of the secret keys.

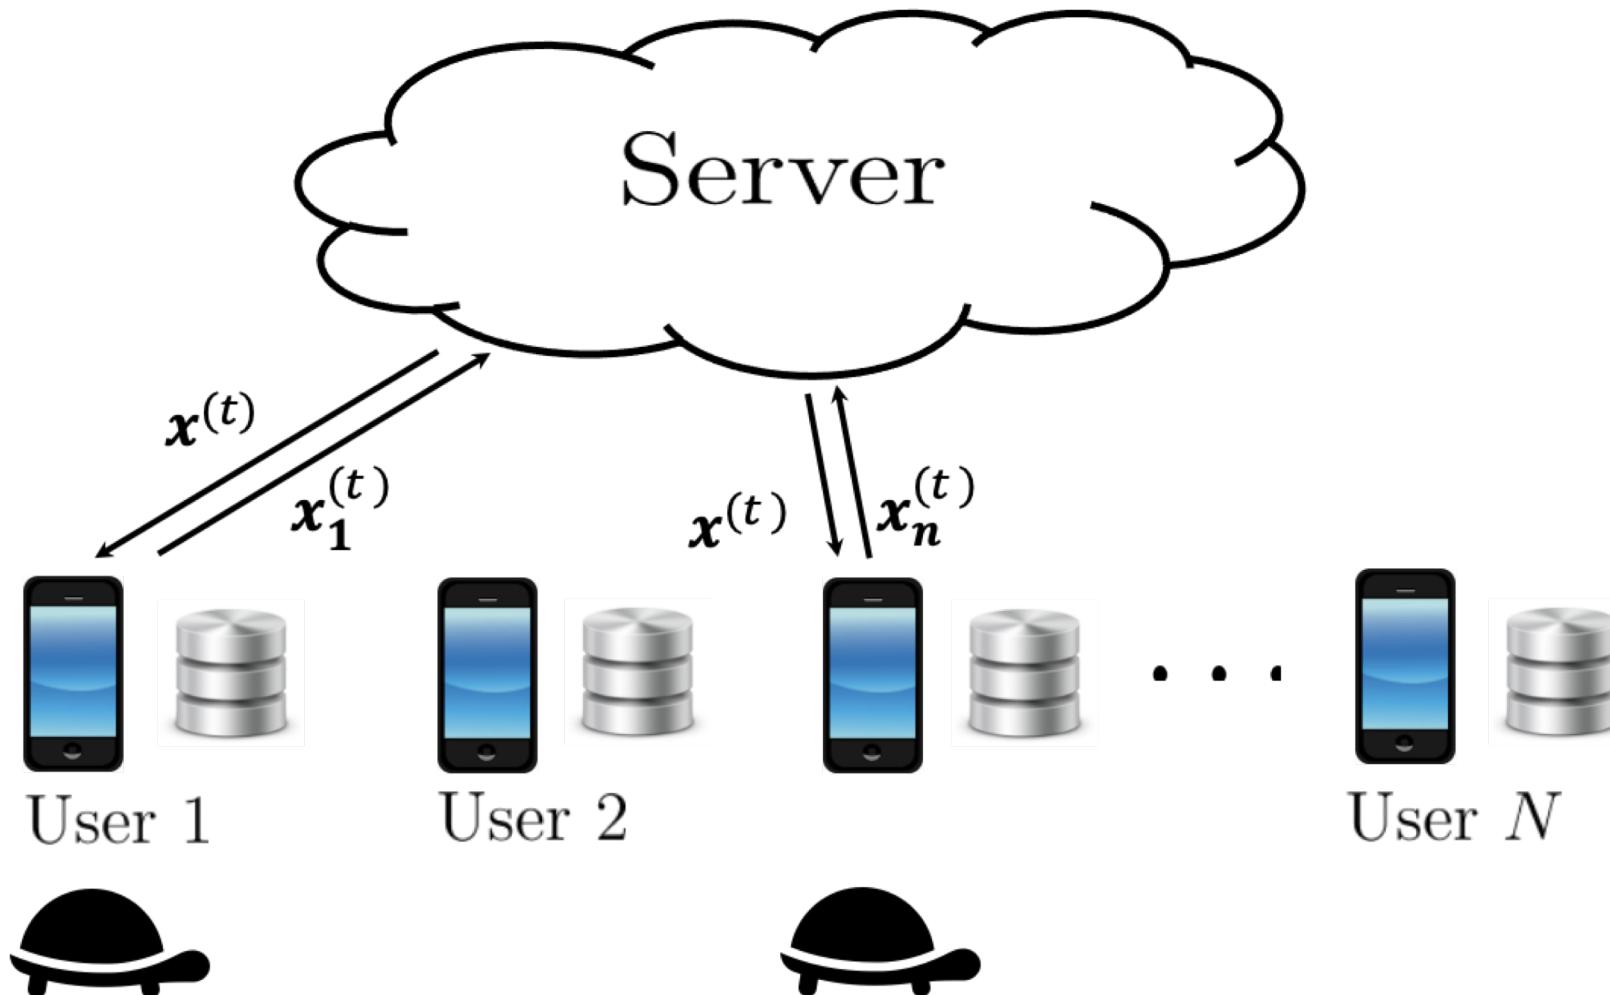
	SecAgg	SecAgg+	LightSecAgg
Offline communication per user	$O(sN)$	$O(s \log N)$	$O(d)$
Offline computation per user	$O(dN + sN^2)$	$O(d \log N + s \log^2 N)$	$O(d \log N)$
Online communication per user	$O(d + sN)$	$O(d + s \log N)$	$O(d)$
Online communication at server	$O(dN + sN^2)$	$O(dN + sN \log N)$	$O(dN)$
Online computation per user	$O(d)$	$O(d)$	$O(d)$
Reconstruction complexity at server	$O(dN^2)$	$O(dN \log N)$	$O(d \log N)$

LightSecAgg significantly improves the computation efficiency
at the server during aggregation.

Part 2. LightSecAgg for Asynchronous Federated Learning

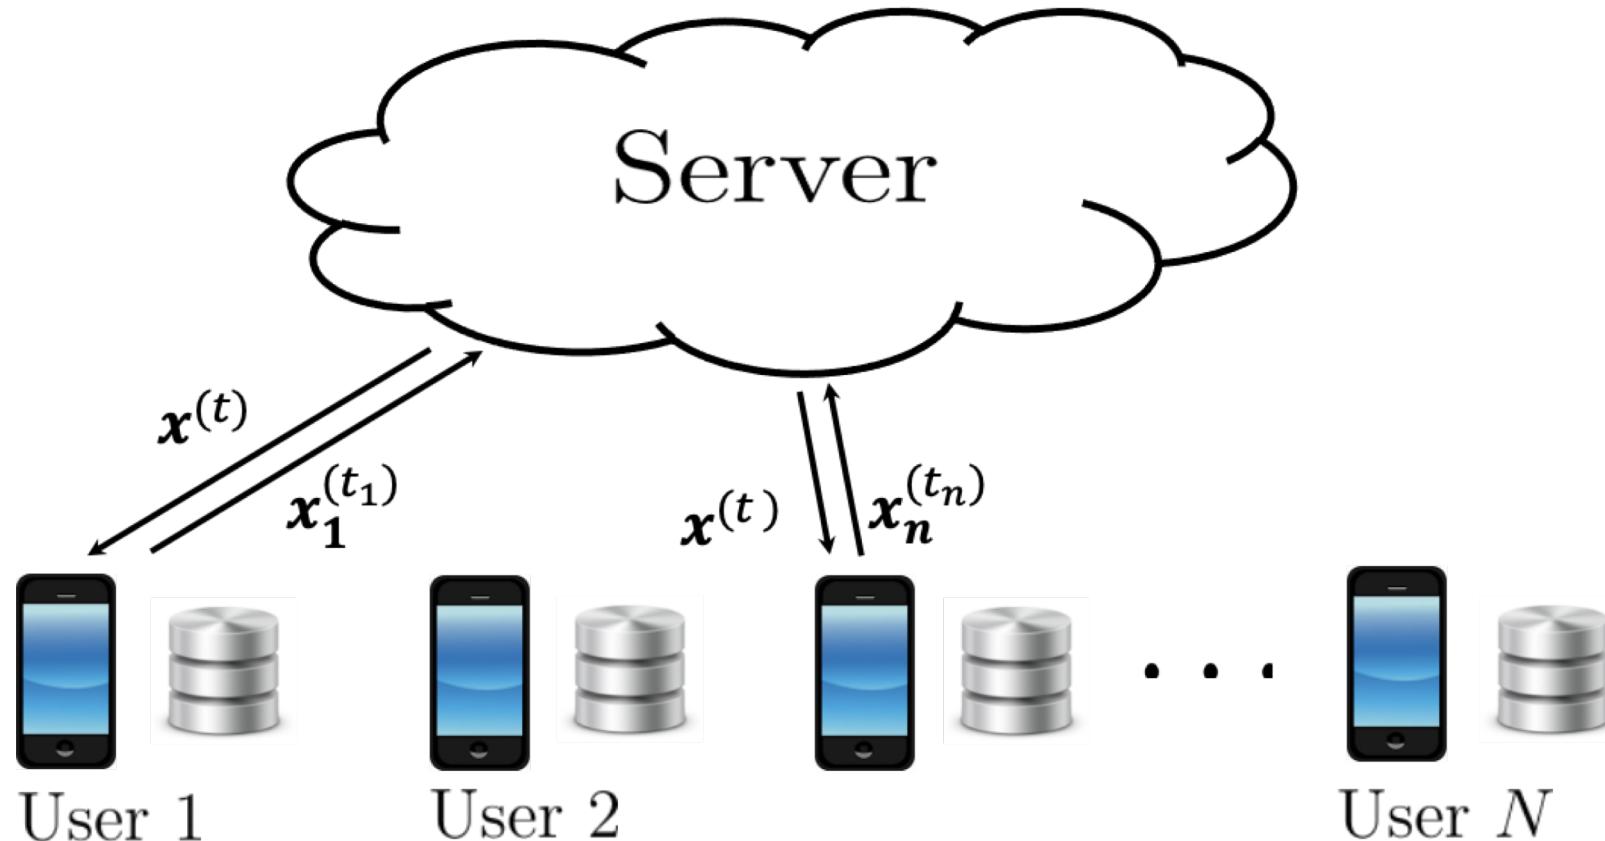
Asynchronous FL

- Synchronous FL suffers from stragglers!



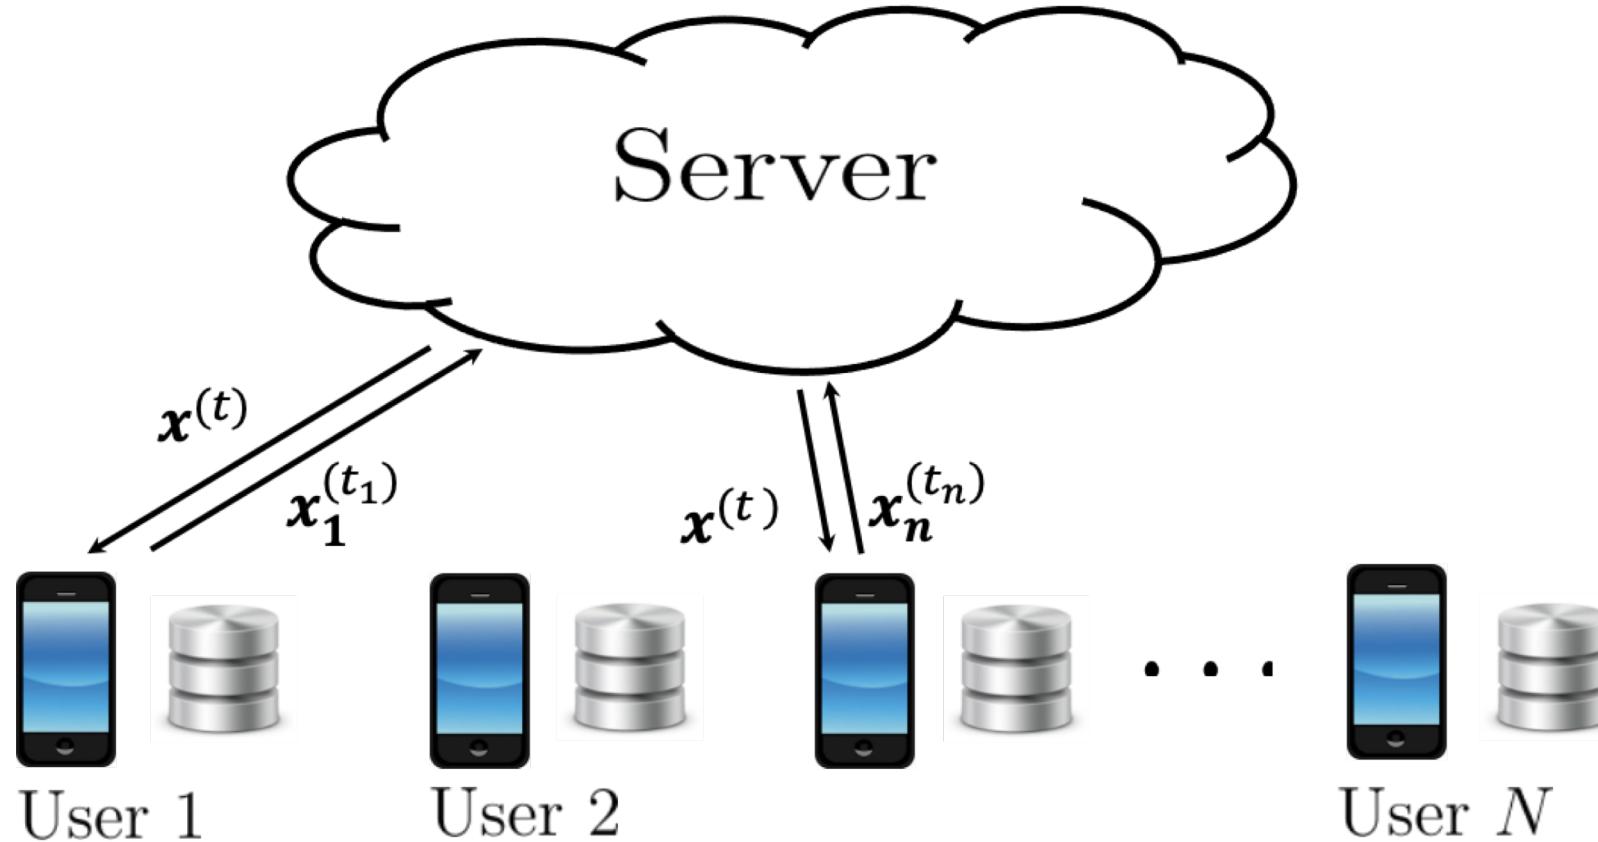
Asynchronous FL

- Updates are not synchronized.
- Each local model received updates the global model.



Asynchronous FL

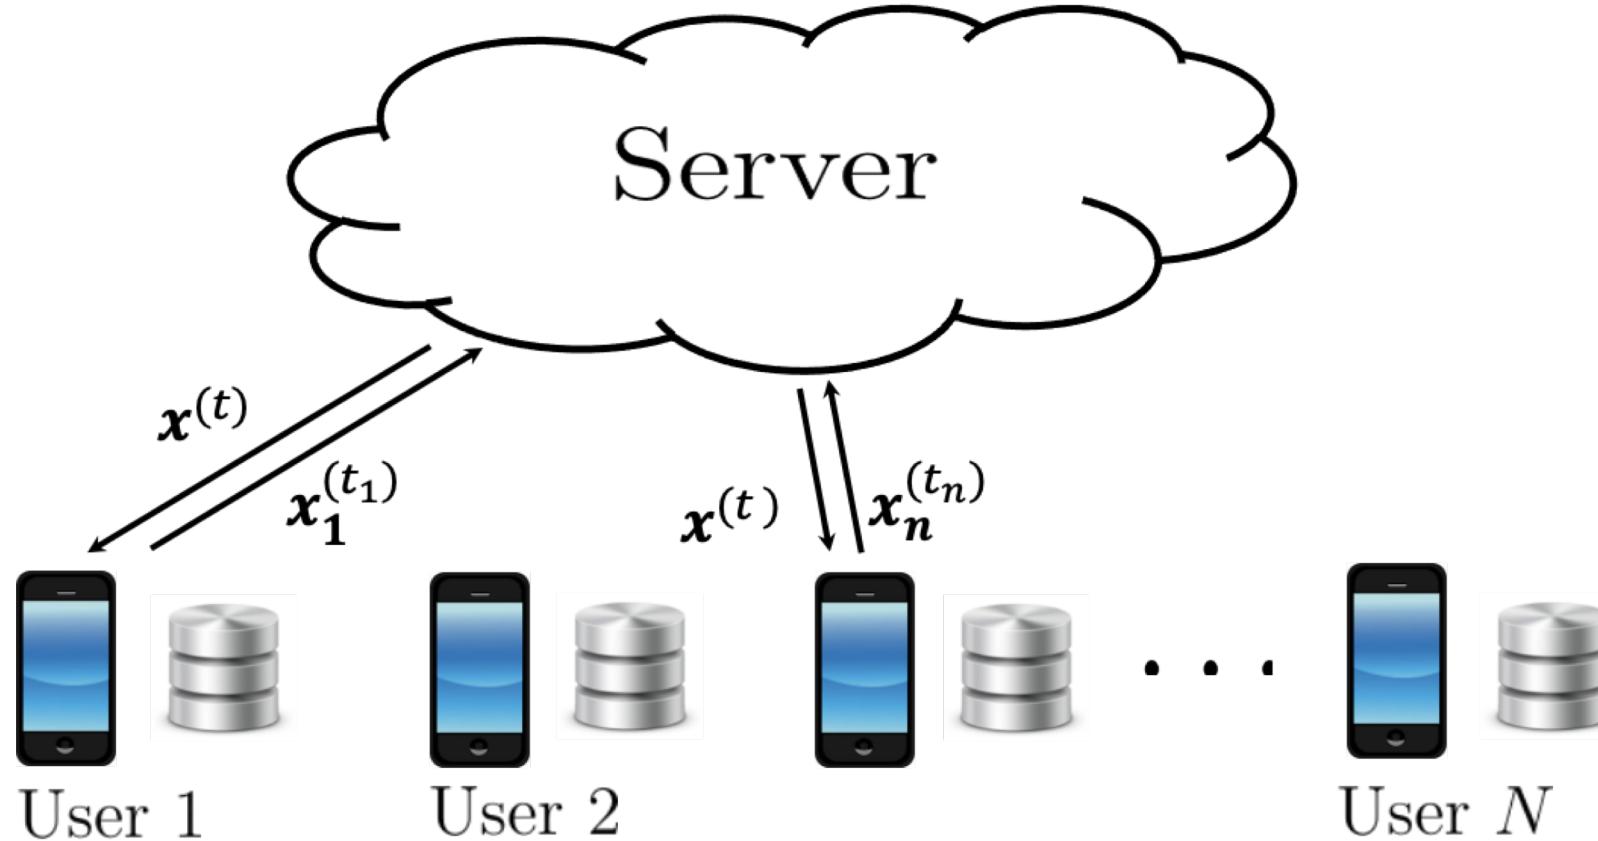
- Updates are not synchronized.
- Each local model received updates the global model.



Not compatible with secure aggregation!

Asynchronous FL

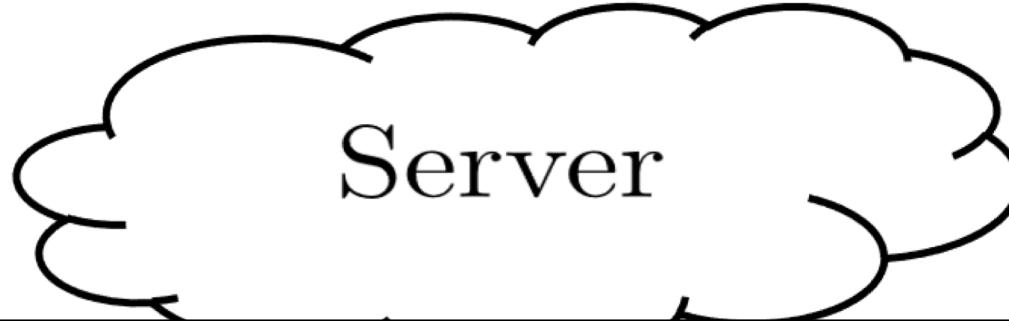
- Updates are not synchronized.
- Each local model received updates the global model.



How we enable secure aggregation in asynchronous FL?

Asynchronous FL

- Updates are not synchronized.
- Each local model received updates the global model.



Federated Learning with Buffered Asynchronous Aggregation



User

John Nguyen Kshitiz Malik Hongyuan Zhan Ashkan Yousefpour
Michael Rabbat Mani Malek Dzmitry Huba

Facebook

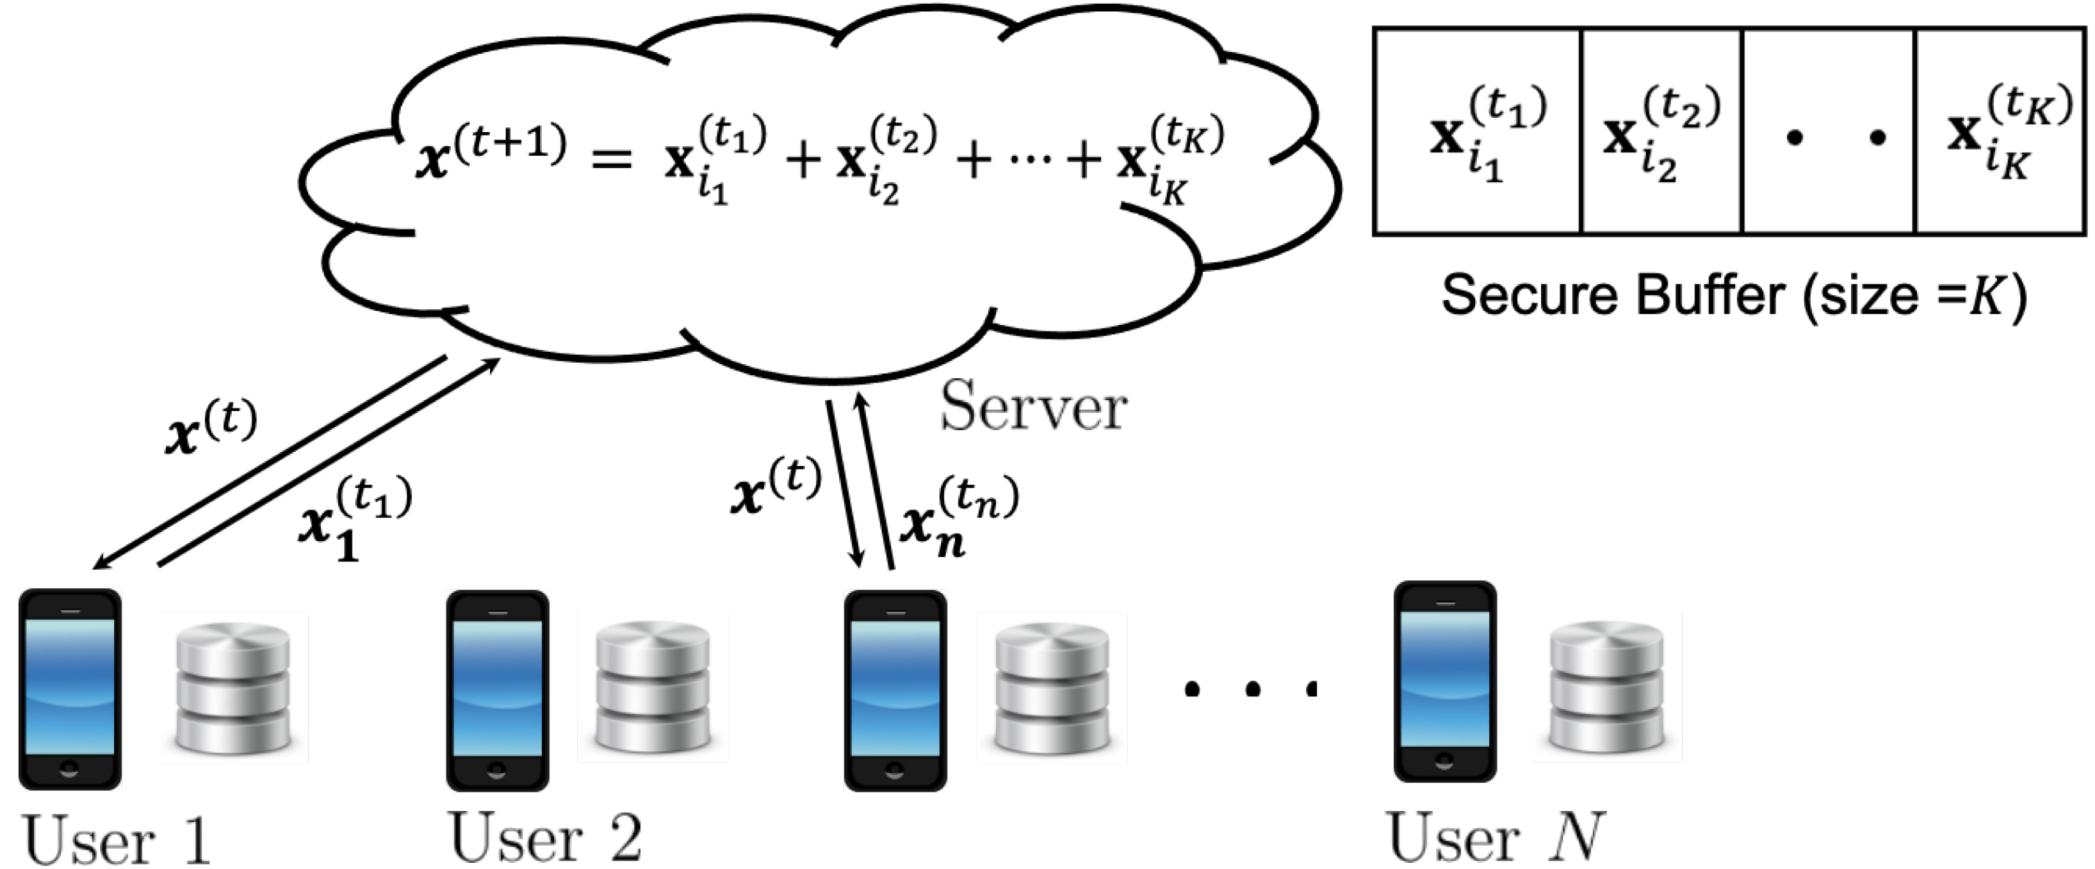
{ngjhn,kmalik2,hyzhan,yousefpour,mikerabbat,manimalek,huba}@fb.com

How we enable secure aggregation in asynchronous FL?



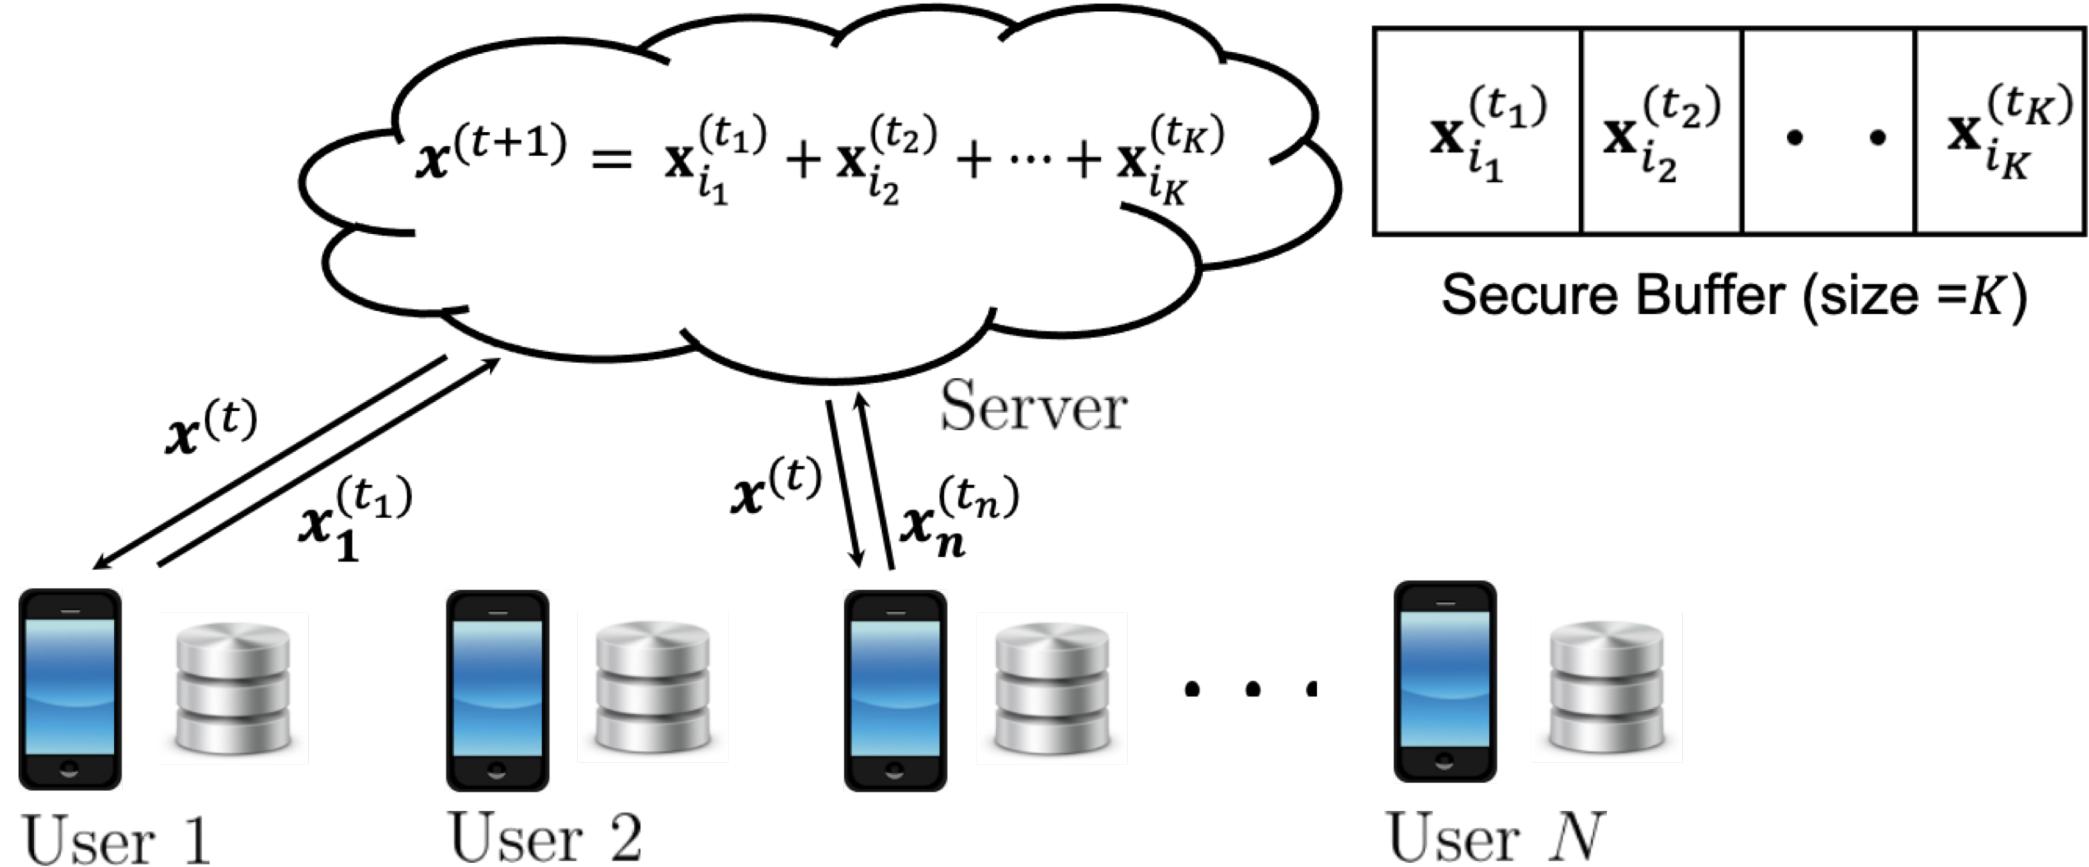
Asynchronous FL

- Typical Asynchronous FL: $K=1$ (not compatible with secure aggregation)
- **Buffered Asynchronous FL (FedBuff): $K>1$**



Asynchronous FL

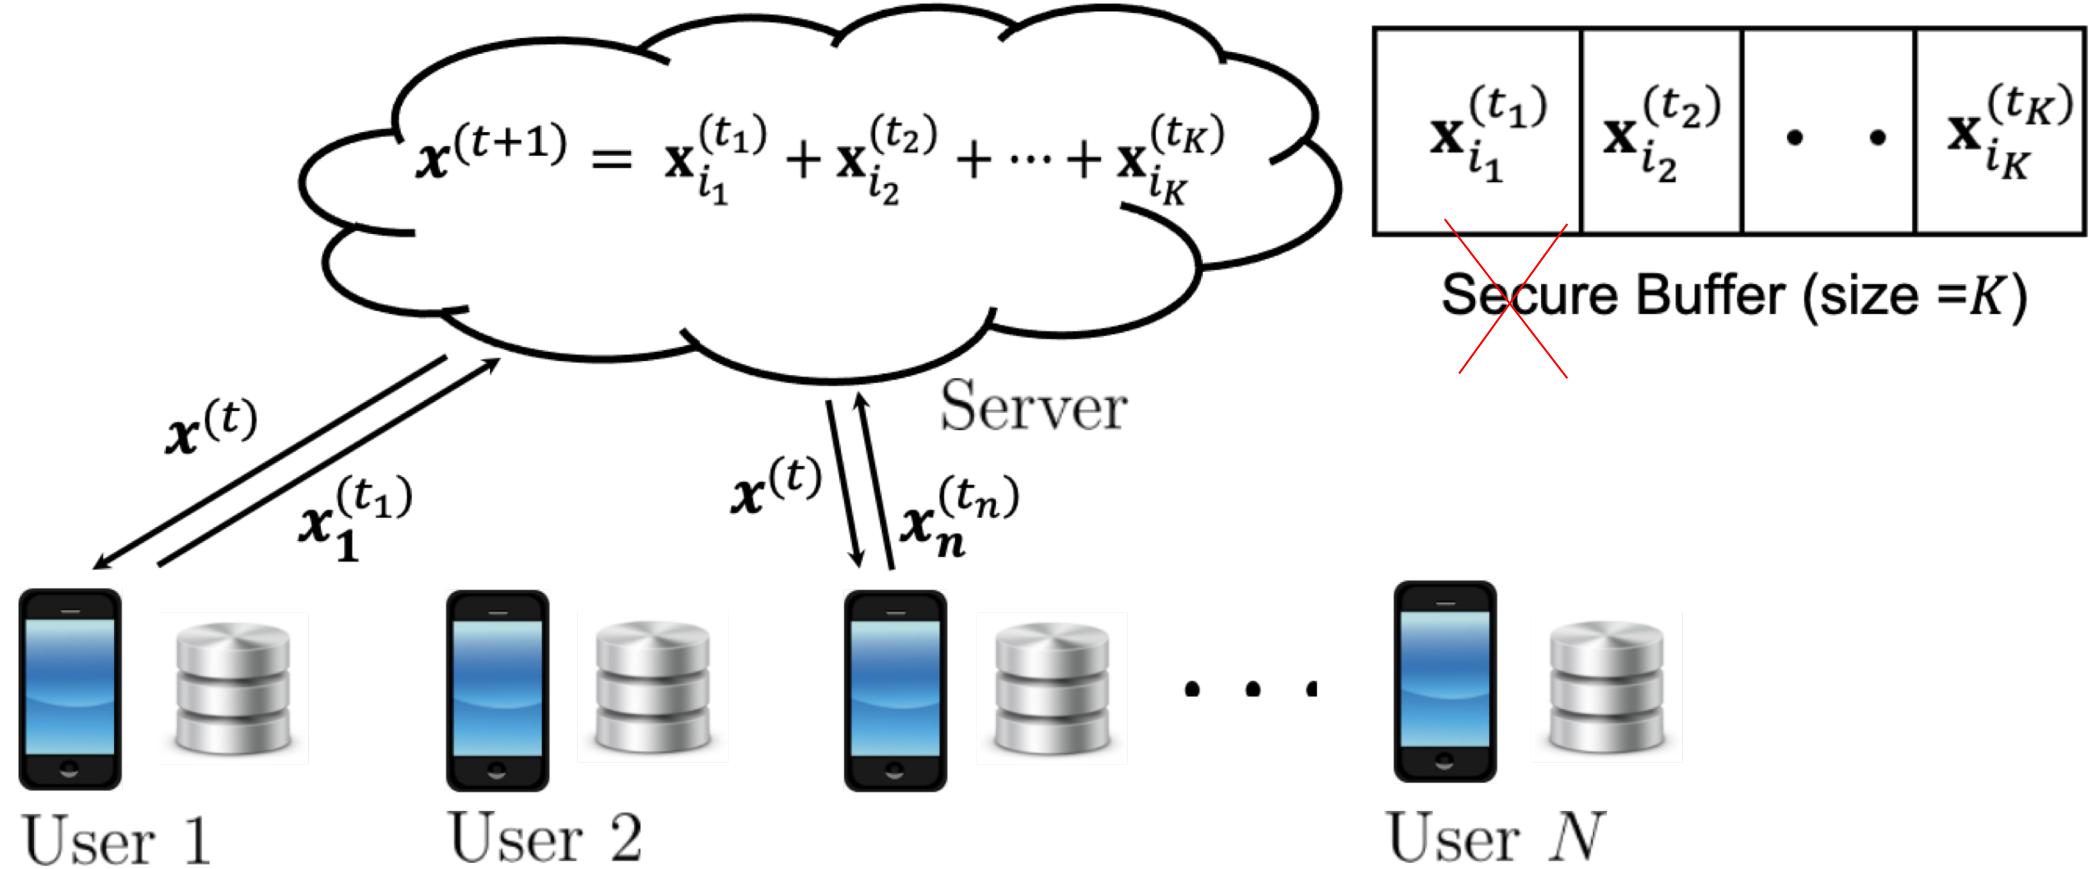
- Typical Asynchronous FL: $K=1$ (not compatible with secure aggregation)
- **Buffered Asynchronous FL (FedBuff): $K>1$**



The TEE-enabled secure buffer, however, has limited memory.
(**K must be small!**)

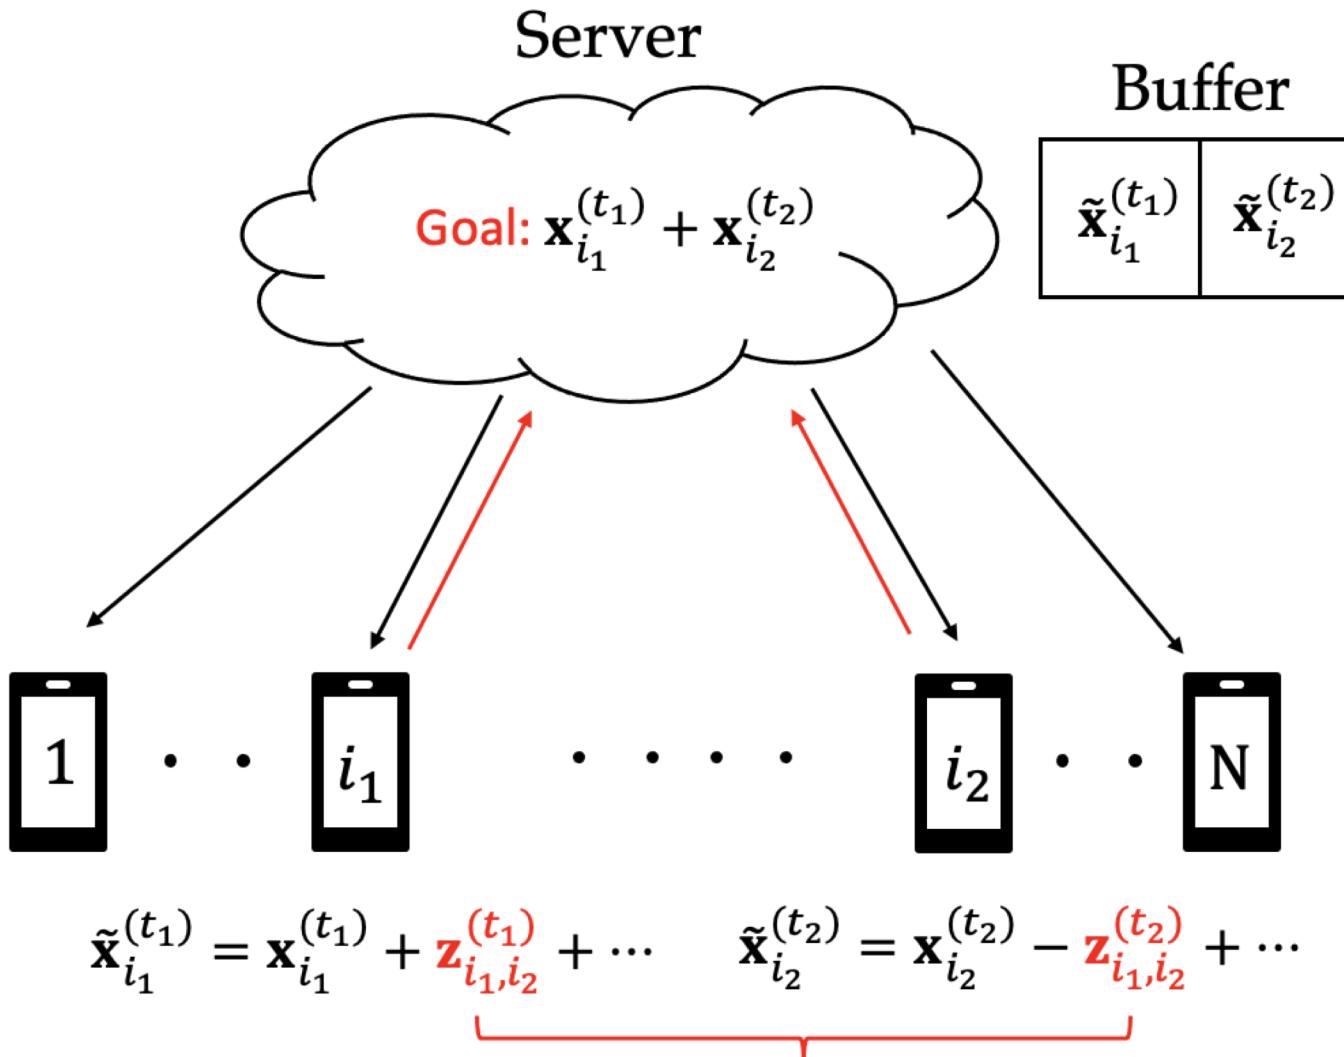
Asynchronous FL

- Typical Asynchronous FL: K=1 (not compatible with secure aggregation)
- Buffered Asynchronous FL (FedBuff): K>1



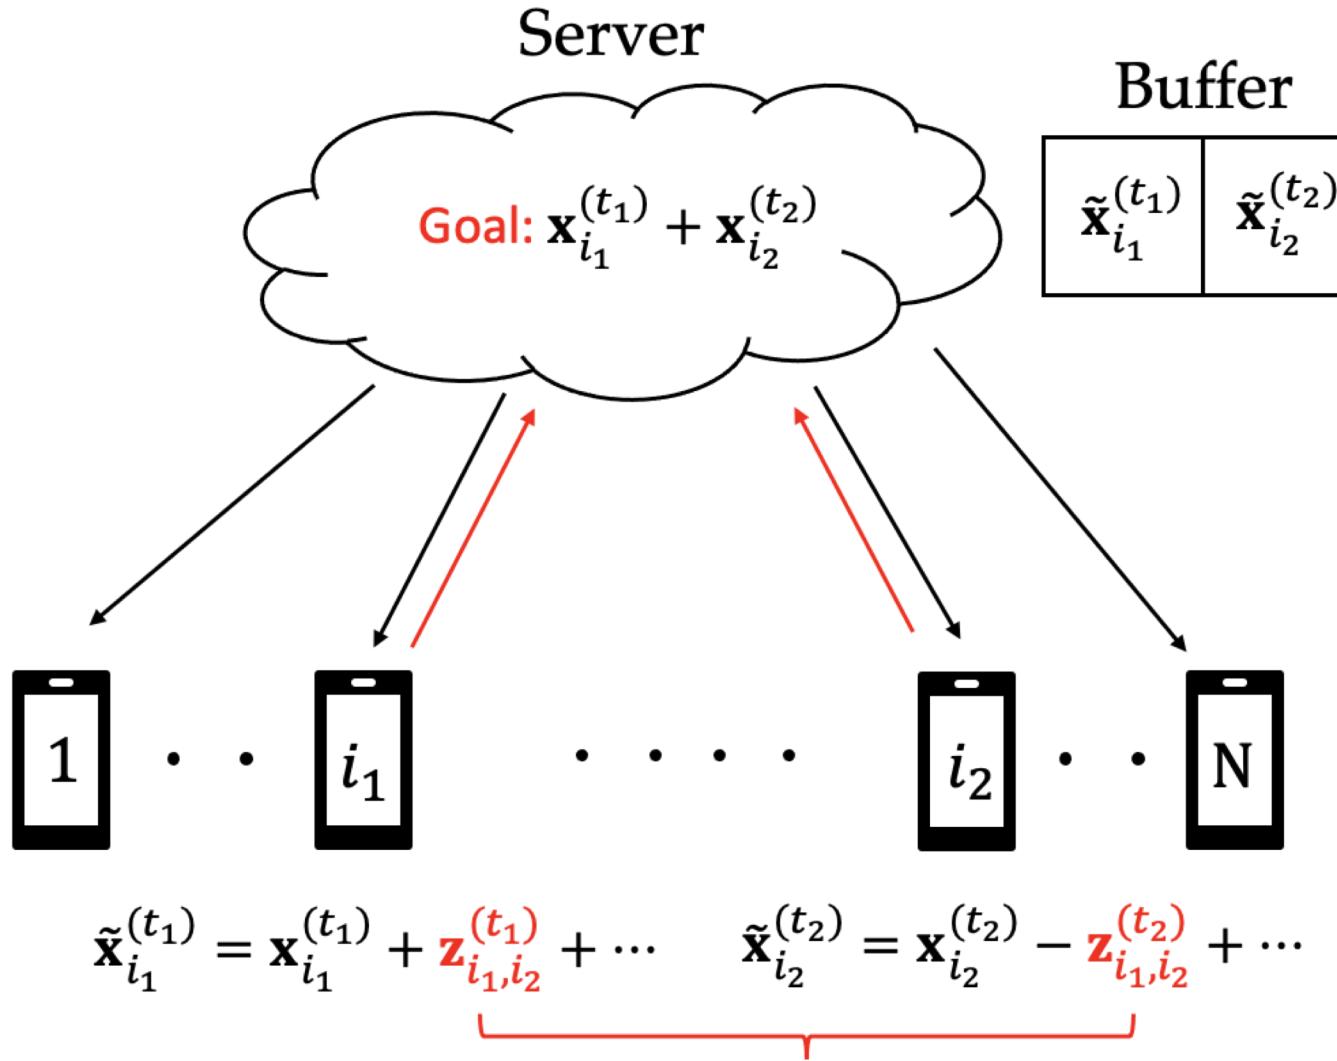
LightSecAgg does not require a secure buffer!

Incompatibility of SecAgg with Asynchronous FL



The masks do not cancel out due to the **mismatch in staleness!**

Incompatibility of SecAgg with Asynchronous FL



How to design the masks to cancel out even if they belong to different rounds?

Asynchronous LightSecAgg

Key objective: Design the masks such that they cancel out even if they belong to different training rounds.

$$\boldsymbol{x}^{(t+1)} = \boldsymbol{x}^{(t)} - \eta_g \boldsymbol{g}^{(t)} \text{ where } \boldsymbol{g}^{(t)} = \sum_{i \in S^{(t)}} \tilde{\boldsymbol{x}}_i^{(t; t_i)} = \sum_{i \in S^{(t)}} \bar{\boldsymbol{x}}_i^{(t; t_i)} + \boxed{\sum_{i \in S^{(t)}} \boldsymbol{z}_i^{(t_i)}}$$

Our focus

LightSecAgg is compatible as it does not use **pair-wise** masking!

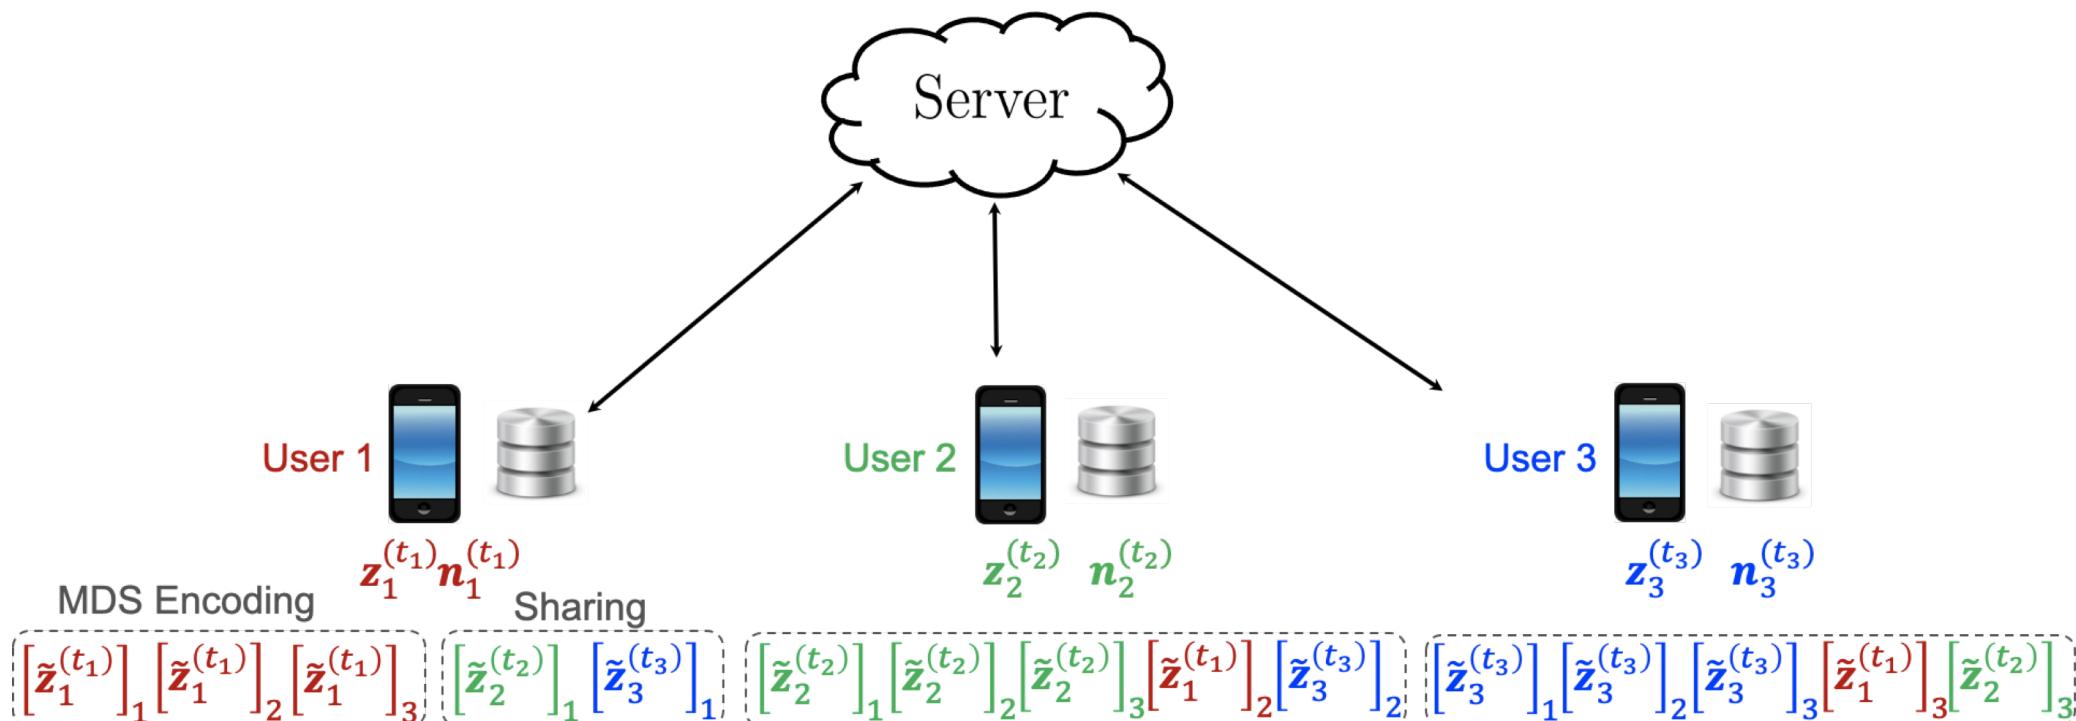
Asynchronous LightSecAgg

Key objective: Design the masks such that they cancel out even if they belong to different training rounds.

$$\boldsymbol{x}^{(t+1)} = \boldsymbol{x}^{(t)} - \eta_g \boldsymbol{g}^{(t)} \text{ where } \boldsymbol{g}^{(t)} = \sum_{i \in S^{(t)}} \tilde{\boldsymbol{x}}_i^{(t; t_i)} = \sum_{i \in S^{(t)}} \bar{\boldsymbol{x}}_i^{(t; t_i)} + \boxed{\sum_{i \in S^{(t)}} \boldsymbol{z}_i^{(t_i)}}$$

Our focus

Step 1. Offline encoding and sharing of local masks



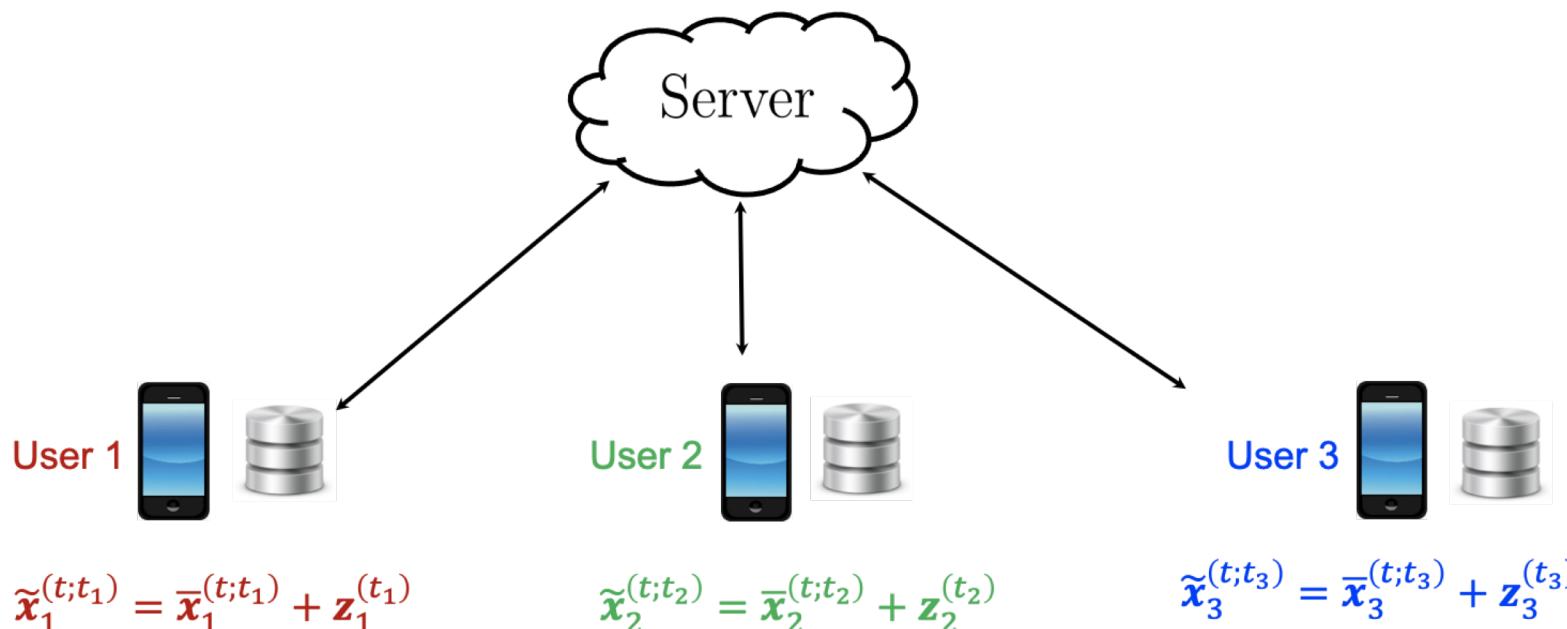
Asynchronous LightSecAgg

Key objective: Design the masks such that they cancel out even if they belong to different training rounds.

$$\boldsymbol{x}^{(t+1)} = \boldsymbol{x}^{(t)} - \eta_g \boldsymbol{g}^{(t)} \text{ where } \boldsymbol{g}^{(t)} = \sum_{i \in S^{(t)}} \tilde{\boldsymbol{x}}_i^{(t; t_i)} = \sum_{i \in S^{(t)}} \bar{\boldsymbol{x}}_i^{(t; t_i)} + \boxed{\sum_{i \in S^{(t)}} \boldsymbol{z}_i^{(t_i)}}$$

Our focus

Step 2. Quantization & Masking



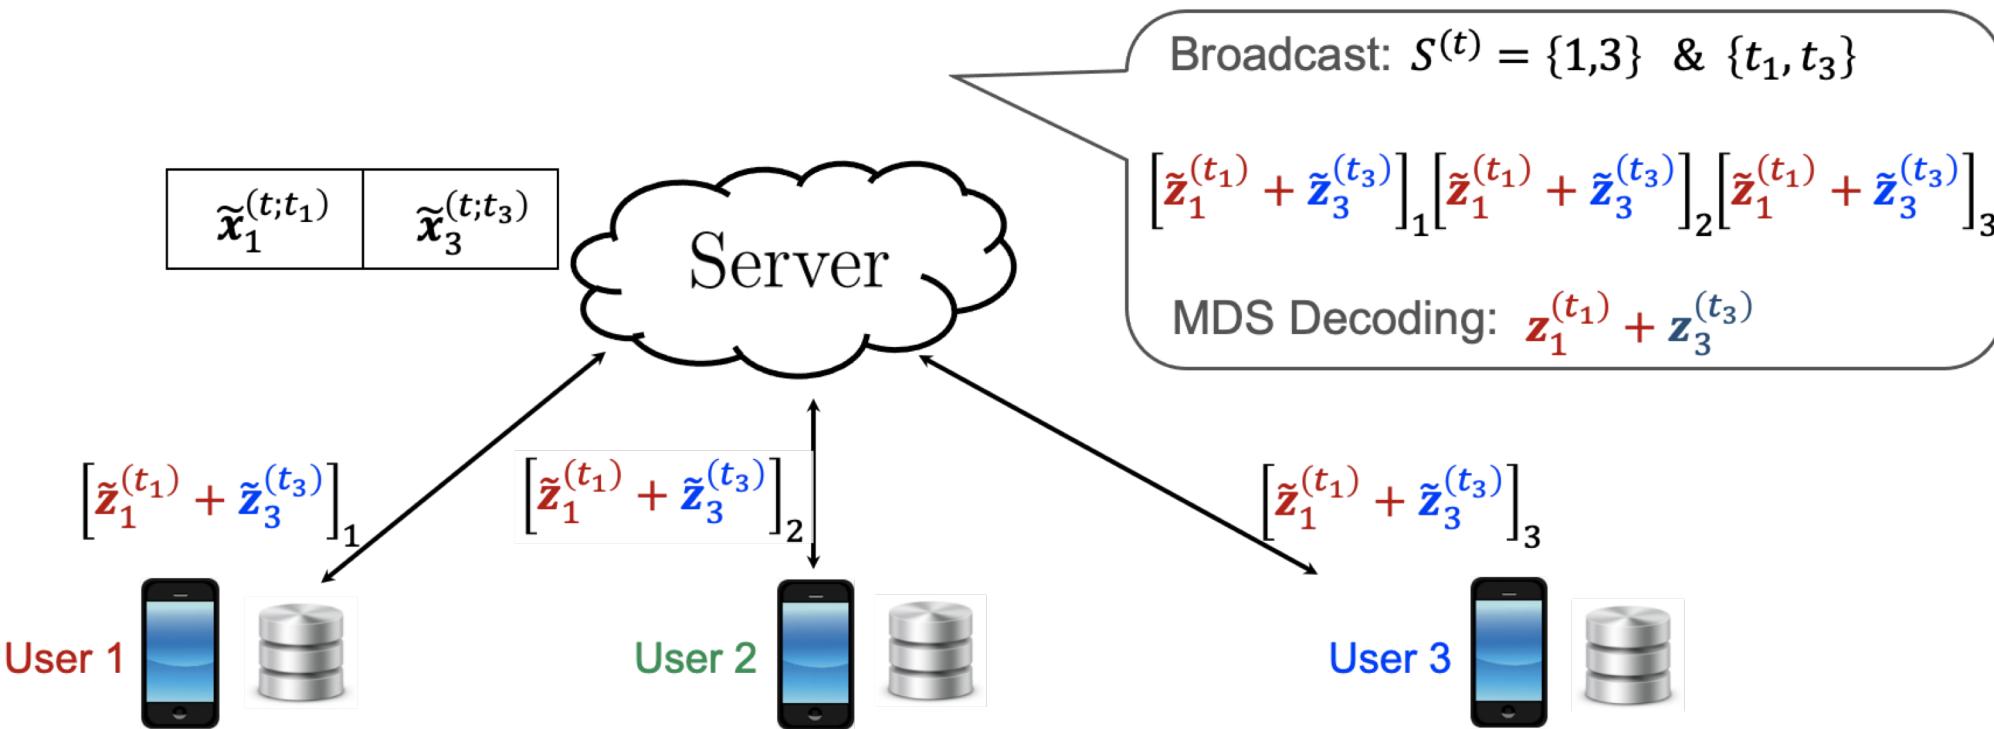
Asynchronous LightSecAgg

Key objective: Design the masks such that they cancel out even if they belong to different training rounds.

$$\mathbf{x}^{(t+1)} = \mathbf{x}^{(t)} - \eta_g \mathbf{g}^{(t)} \text{ where } \mathbf{g}^{(t)} = \sum_{i \in S^{(t)}} \tilde{\mathbf{x}}_i^{(t;t_i)} = \sum_{i \in S^{(t)}} \bar{\mathbf{x}}_i^{(t;t_i)} + \boxed{\sum_{i \in S^{(t)}} \mathbf{z}_i^{(t_i)}}$$

Our focus

Step 3. One-Shot Recovery of Aggregate Masks



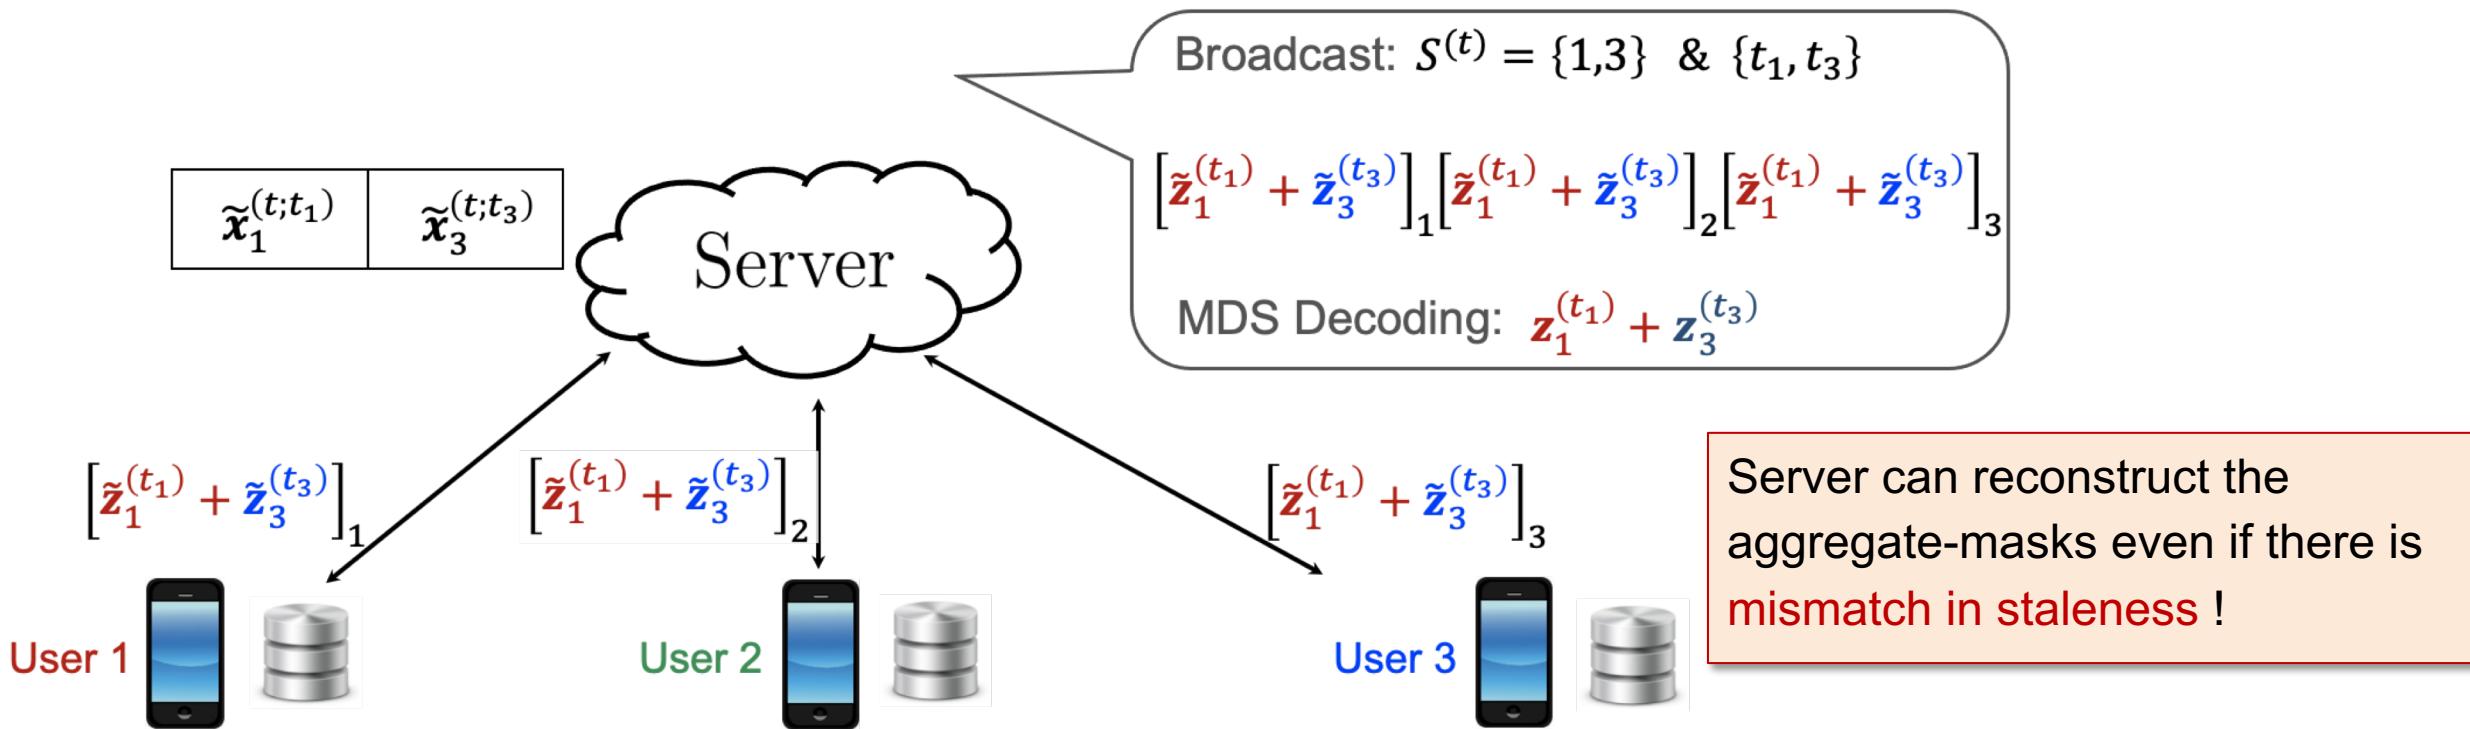
Asynchronous LightSecAgg

Key objective: Design the masks such that they cancel out even if they belong to different training rounds.

$$\mathbf{x}^{(t+1)} = \mathbf{x}^{(t)} - \eta_g \mathbf{g}^{(t)} \text{ where } \mathbf{g}^{(t)} = \sum_{i \in S^{(t)}} \tilde{\mathbf{x}}_i^{(t;t_i)} = \sum_{i \in S^{(t)}} \bar{\mathbf{x}}_i^{(t;t_i)} + \boxed{\sum_{i \in S^{(t)}} \mathbf{z}_i^{(t_i)}}$$

Our focus

Step 3. One-Shot Recovery of Aggregate Masks



Convergence Analysis

- Convergence rate of LightSecAgg and FedBuff are the same except for the increased variance of the local updates due to the quantization noise.

Theorem: *The global model of LightSecAgg achieves the following ergodic convergence rate*

$$\frac{1}{J} \sum_{t=0}^{J-1} \mathbb{E} \left[|\nabla F(\mathbf{x}^{(t)})|^2 \right] \leq \frac{2F^*}{\eta_g \eta_l E K J} + \frac{L \eta_g \eta_l \sigma_{c_l}^2}{2} + 3L^2 E^2 \eta_l^2 (\eta_g^2 K^2 \tau_{\max}^2) \sigma^2,$$

where $F^* = F(\mathbf{x}^{(0)}) - F(\mathbf{x}^*)$, $\sigma^2 = G + \sigma_g^2 + \sigma_{c_l}^2$, and $\sigma_{c_l}^2 = \boxed{\frac{d}{4c_l^2}} + \sigma_l^2$.

τ_{\max} : bound on delay

$\sigma_{c_l}^2$: bound on local variance

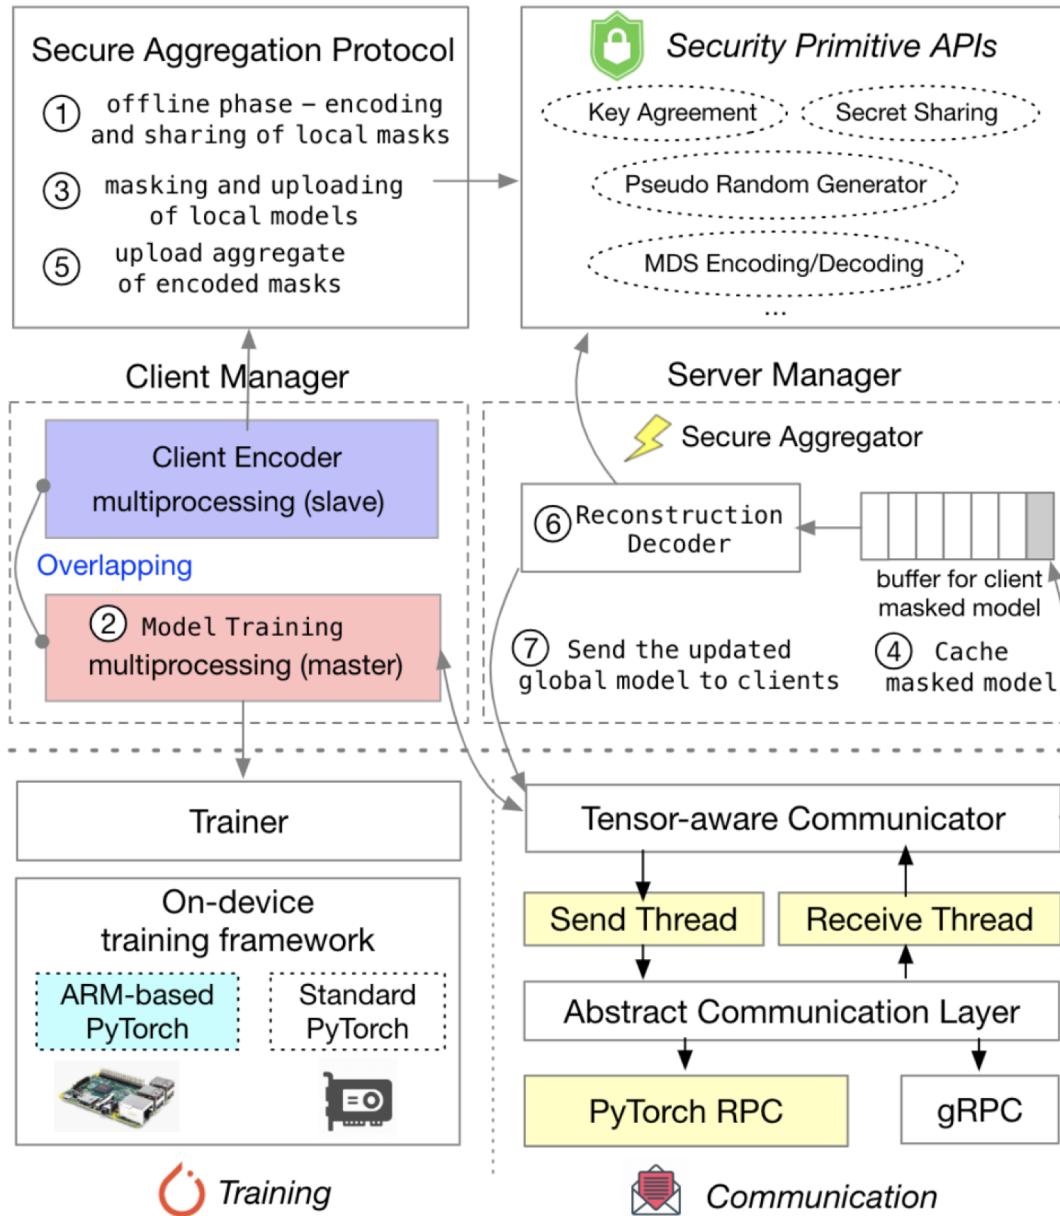
σ_g^2 : bound on global variance

L-smooth gradients

Negligible for large c_l

Part 3. System Design and Experiments

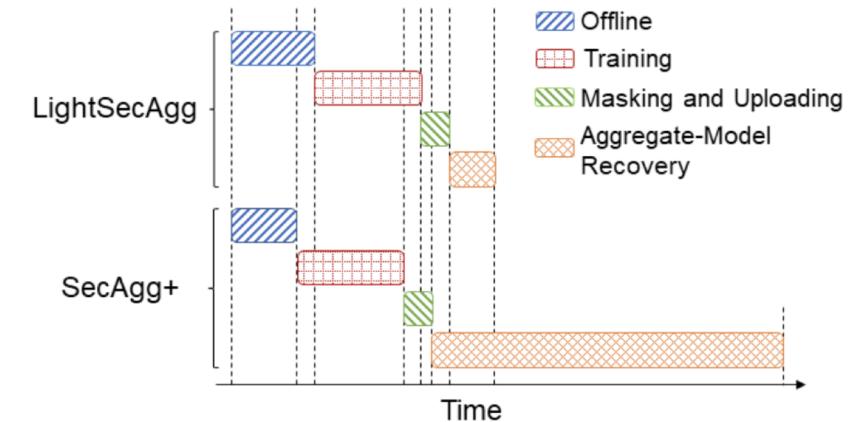
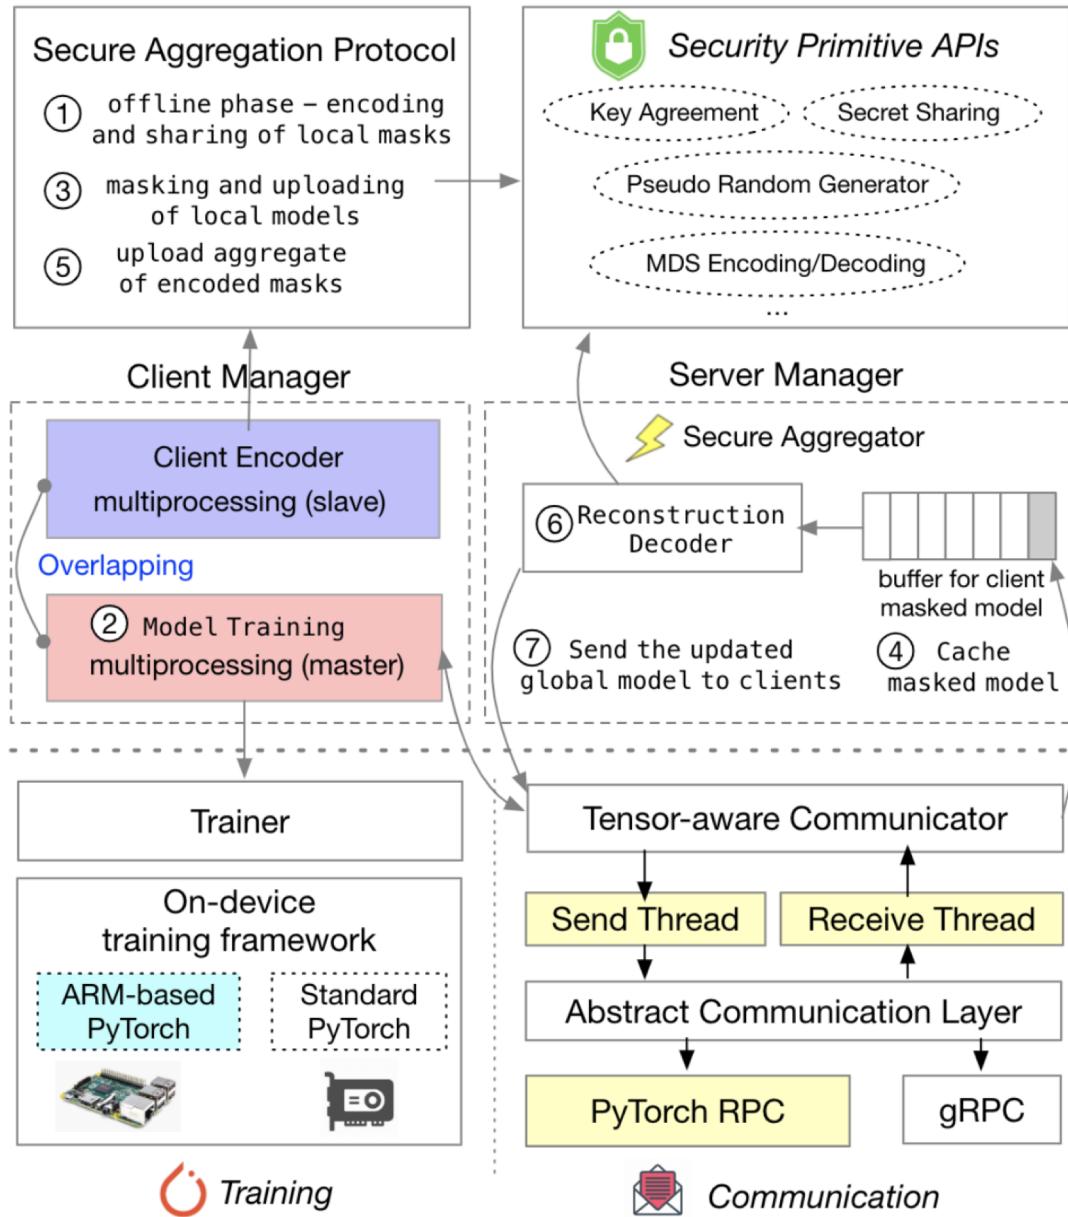
Overview of the System Design



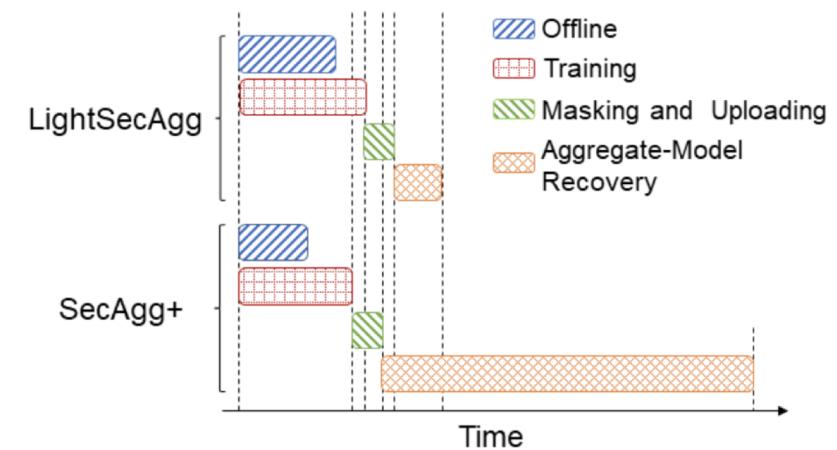
Design Goals:

1. Reduce the cost of encoding and decoding at the edge
2. Optimize the communication backend, making it Torch Tensor-aware
3. Make the system API friendly to pure ML researchers who may not have expertise in SA/Security.

1. Parallelization of offline phase and model training

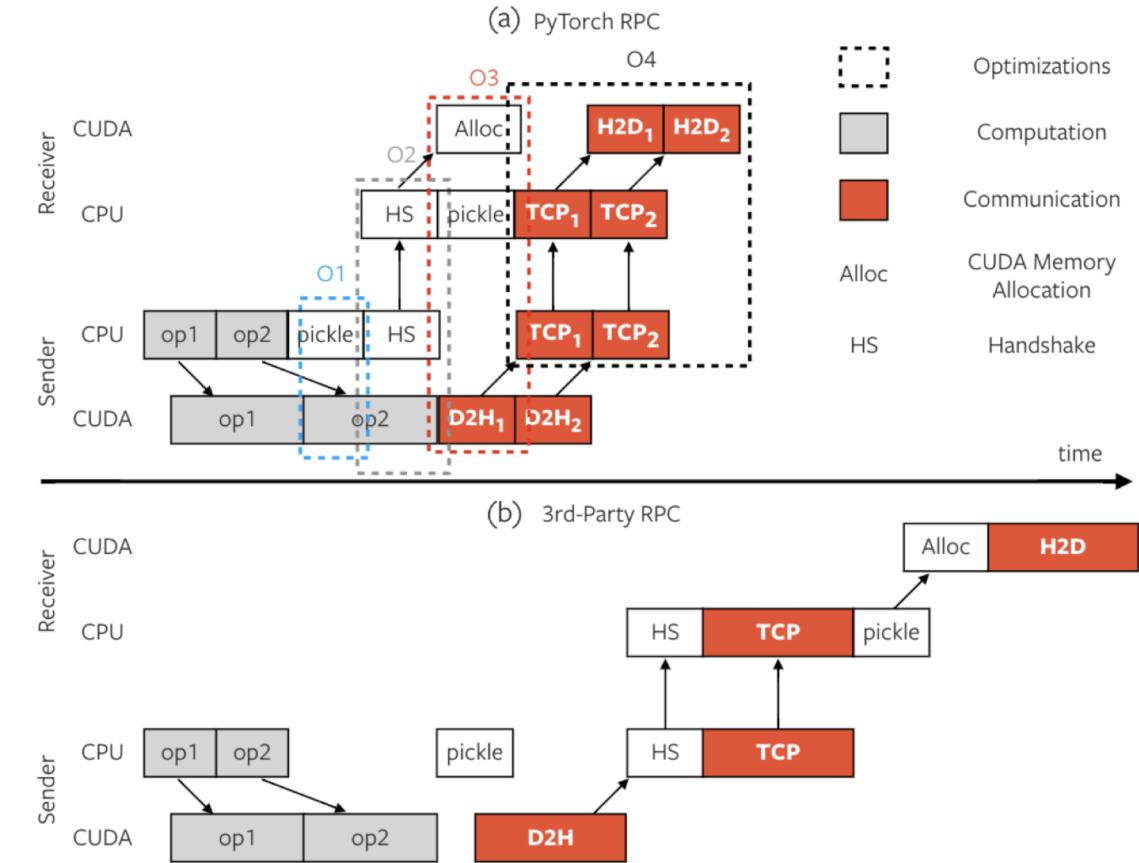
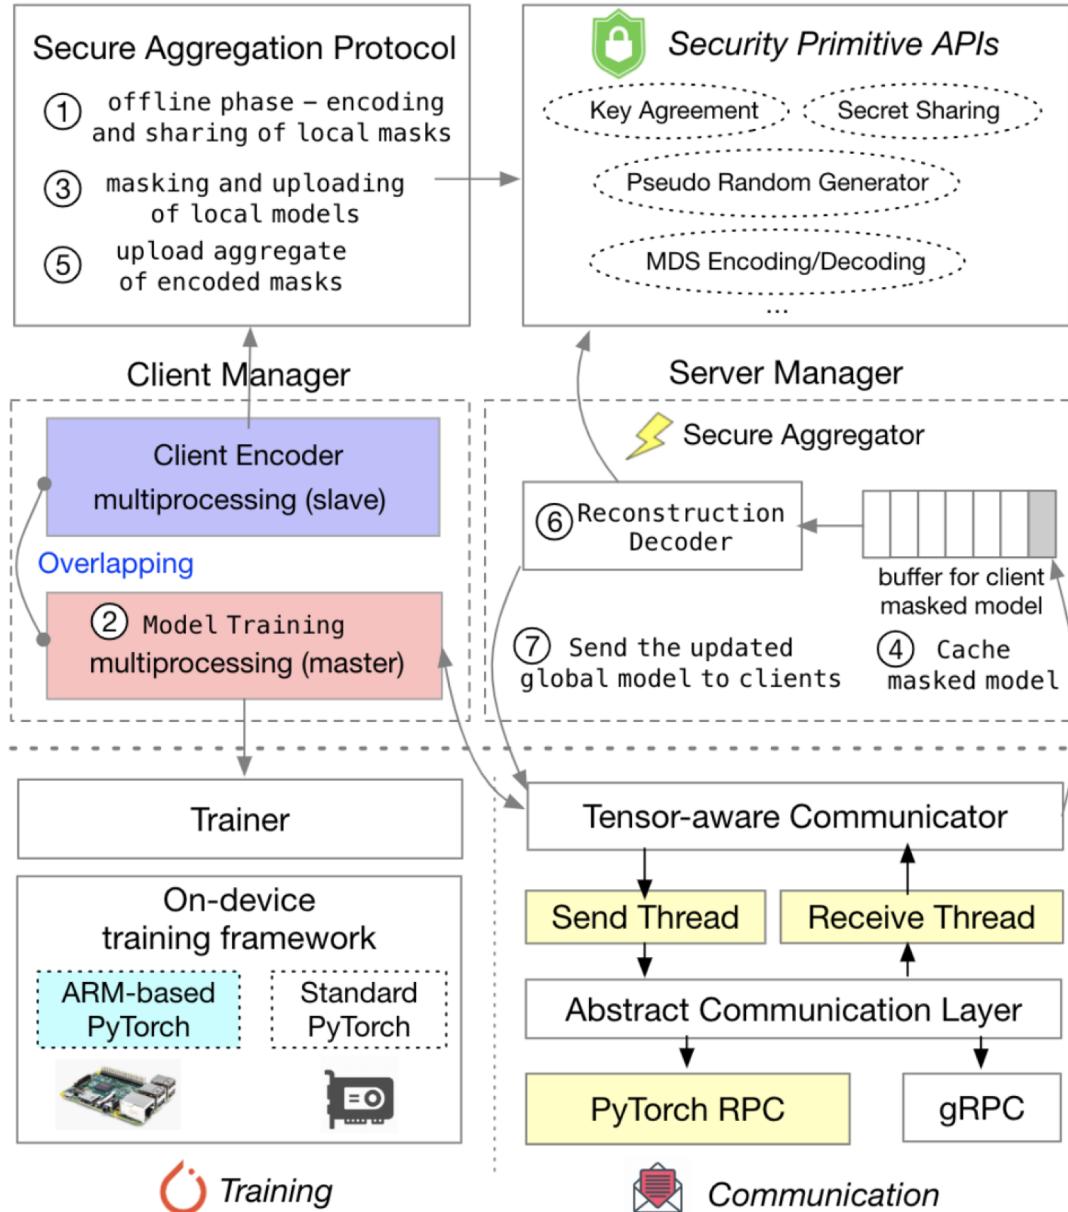


(a) Non-overlapped



(b) Overlapped

2. Tensor-aware RPC (Remote Procedure Call)

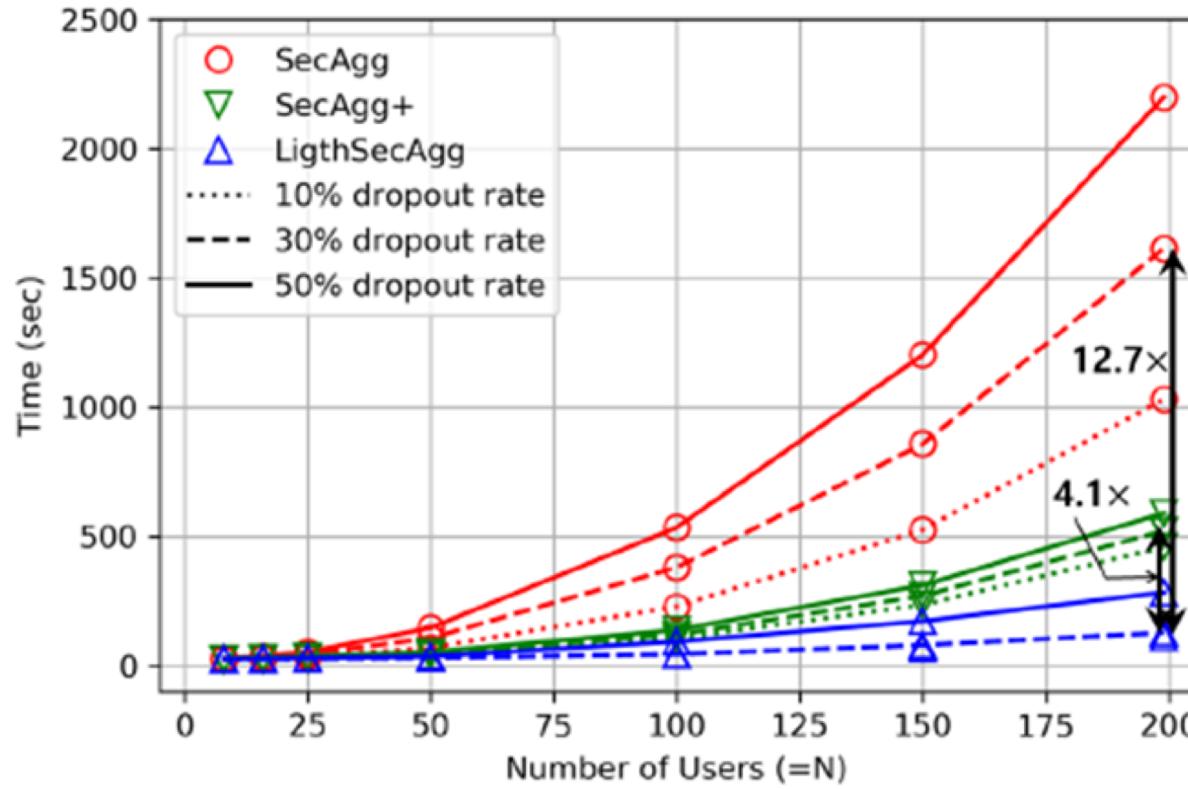


Comparison with Third-Party RPC

Experiments

- Experiment setup:
 - Amazon EC2 cloud using m3.medium machine instances
 - Four different machine learning tasks
 - Communication using the MPI4Py message passing interface on Python
 - Each user drops with a fixed dropout rate $p = 0.1$, $p = 0.3$, and $p = 0.5$

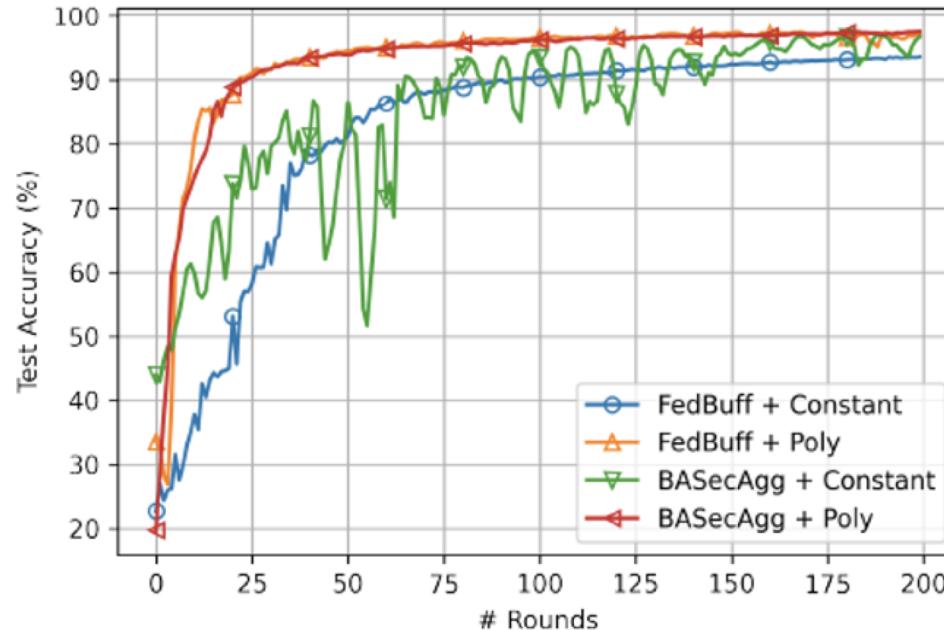
Experiments



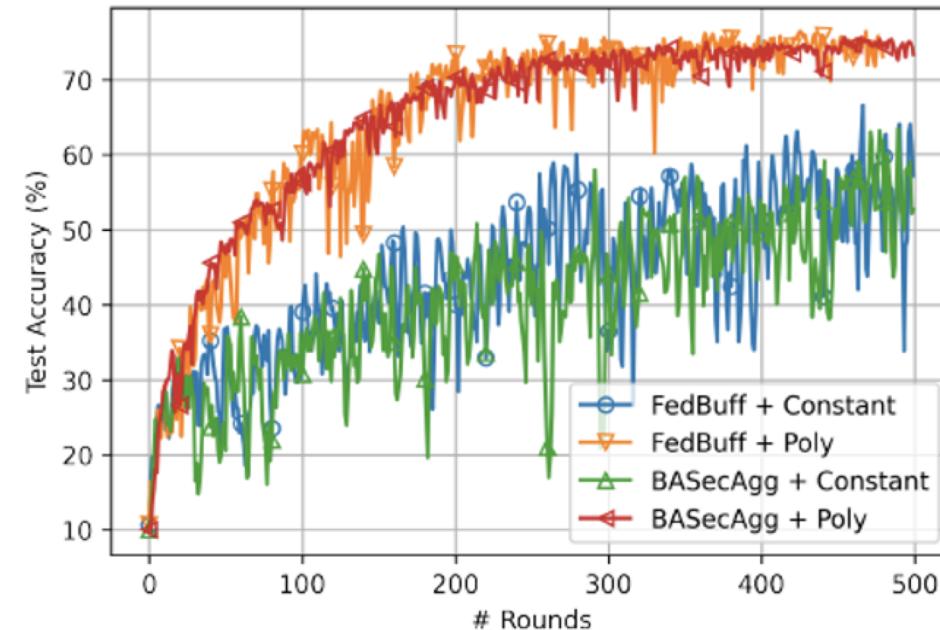
LightSecAgg achieves a performance gain of up to **12.7x !!**

Experiments

- LightSecAgg has almost the same performance as FedBuff while LightSecAgg includes quantization noise to for privacy.



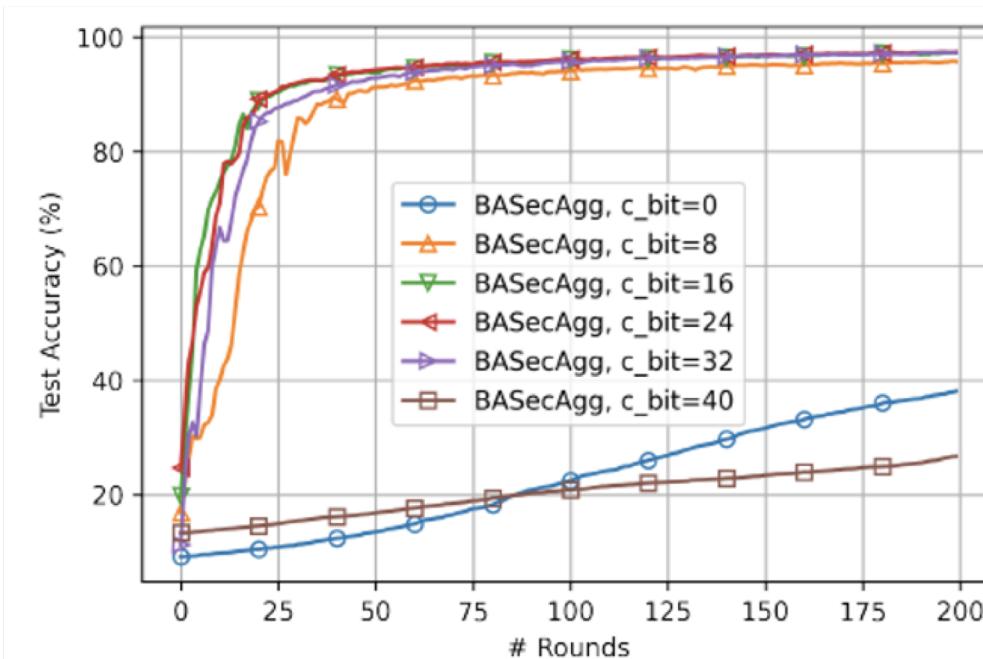
(a) MNIST dataset.



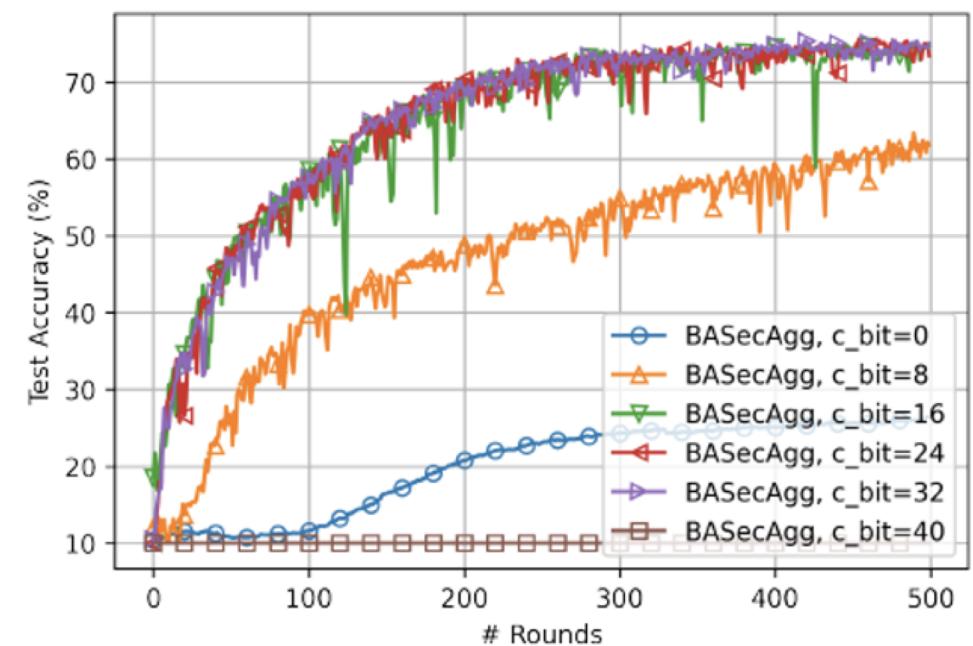
(b) CIFAR-10 dataset.

Experiments

- Effect of the quantization parameter: Trade-off between rounding and wrap-around errors



(a) MNIST dataset.



$$c_l = 2^{c_{bit}}$$

(b) CIFAR-10 dataset.

Large c_l results in large wrap-around error, while small c_l results in large rounding error.

Experiments

Table 2: Summary of four implemented machine learning tasks and performance gain of LightSecAgg with respect to SecAgg [4] and SecAgg+ [2]. All learning tasks are for image classification. MNIST, FEMNIST and CIFAR-100 are low-resolution datasets, while images in GLD-23K are high resolution, which cost much longer training time for one mini-batch; LR and CNN are shallow models, but MobileNetV3 and EfficientNet-B0 are much larger models, but they are tailored for efficient edge training and inference.

No.	Dataset	Model	Model Size (d)	Gain	
				Non-overlapped	Overlapped
1	MNIST [14]	Logistic Regression	7,850	6.7 \times , 2.5 \times	8.0 \times , 2.9 \times
2	FEMNIST [5]	CNN [17]	1,206,590	11.3 \times , 3.7 \times	12.7 \times , 4.1 \times
3	CIFAR-100 [13]	MobileNetV3 [11]	3,111,462	7.6 \times , 2.8 \times	9.5 \times , 3.3 \times
4	GLD-23K [27]	EfficientNet-B0 [24]	5,288,548	3.3 \times , 1.6 \times	3.4 \times , 1.7 \times

LightSecAgg can survive and speedup the training of **large** deep neural network models on high resolution image datasets.

Concluding Remarks

- We propose a new perspective for secure model aggregation in FL, by turning the focus from “pairwise random-seed reconstruction of the dropped users” to “one-shot aggregate-mask reconstruction of the surviving users”.
- We propose LightSecAgg that provides the same level of privacy and dropout-resiliency guarantees as the state-of-the-art while substantially reducing the aggregation complexity.
- In a realistic FL framework, extensive empirical results show that LightSecAgg can provide substantial speedup over baseline protocols for training diverse machine learning models with a performance gain up to 12.7x.