

# Erasure-Coded Key-Value Stores with Side Information

Ramy E. Ali

Algorithms, Analytics & Augmented Intelligence group, Math  
of Communications department, August 2018

# Outline

- Key-value Stores Overview
- Background: Replication & Erasure Coding
- Coding with Side Information: Problem Formulation
- Impossibility Results
- Code Constructions
- Case Study: Latency-Storage Trade-off in AWS
- Discussion

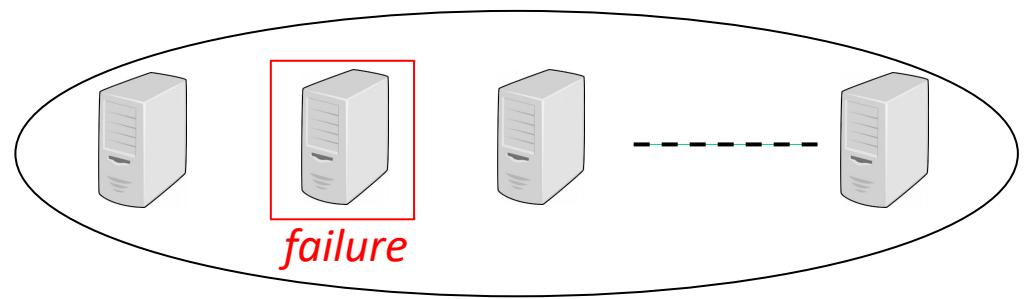
# Key-value Stores

- Applications: reservation systems, financial transactions, distributed computing, ...
- Numerous key-value stores: Amazon Dynamo, Apache Cassandra, and CouchDB



# Distributed Key-value Stores

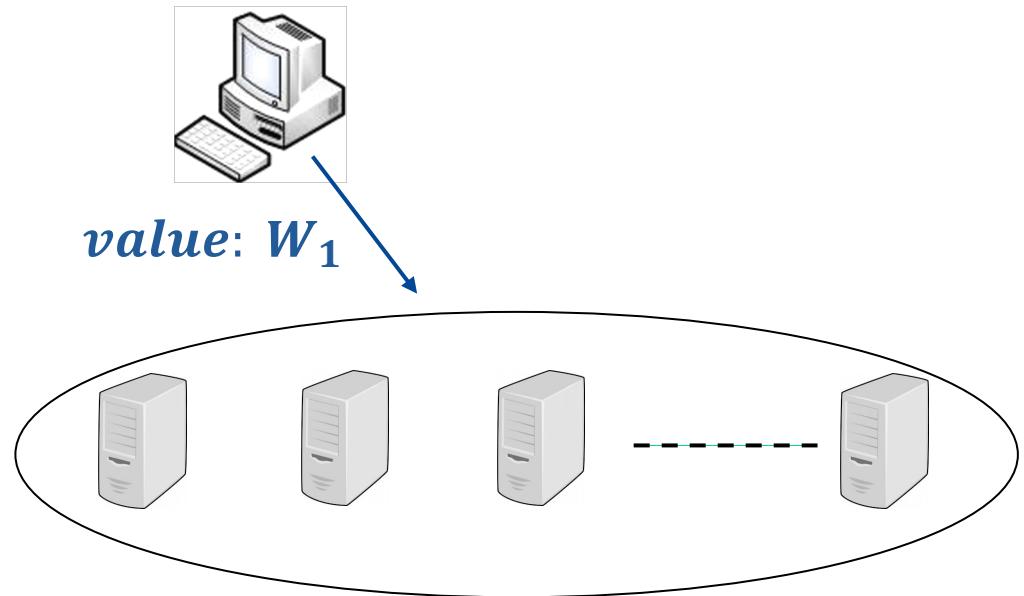
- Data is stored over multiple nodes.



# Distributed Key-value Stores


- Data is stored over multiple nodes.
- Data is asynchronously updated.

Write client




# Distributed Key-value Stores

- Data is stored over multiple nodes.
- Data is asynchronously updated.



# Distributed Key-value Stores

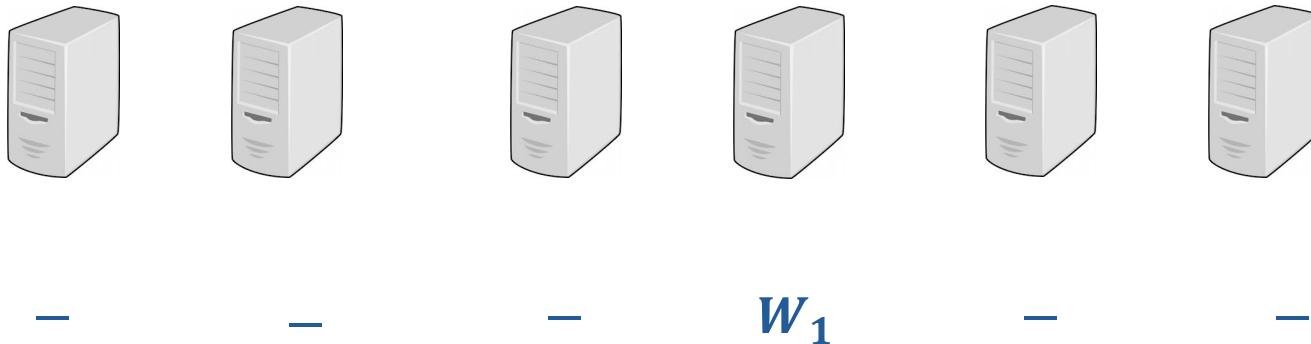
- Data is stored over multiple nodes.
- Data is asynchronously updated.
- Client must get the *latest possible version* of the data [Lamport 1979, ABD 1995].



# Distributed Key-value Stores

## 1. Asynchrony

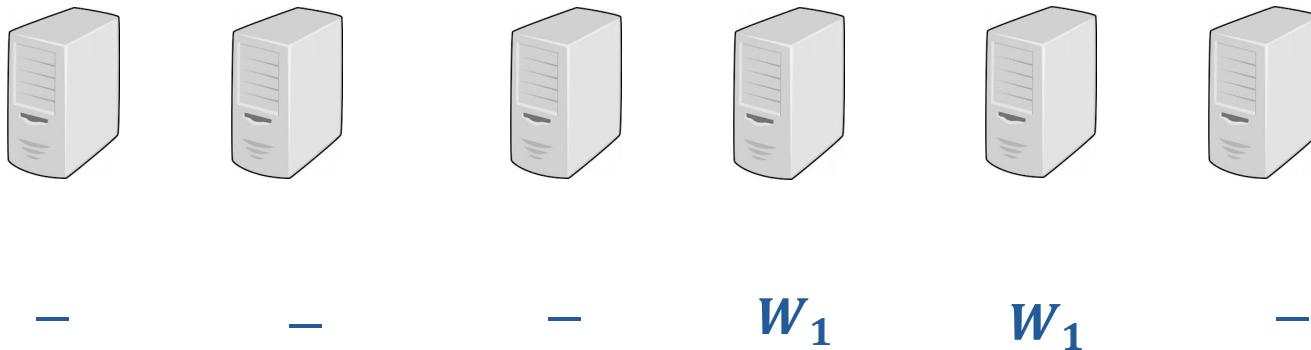
*Data updates may not arrive at all servers simultaneously.*



# Distributed Key-value Stores

## 1. Asynchrony

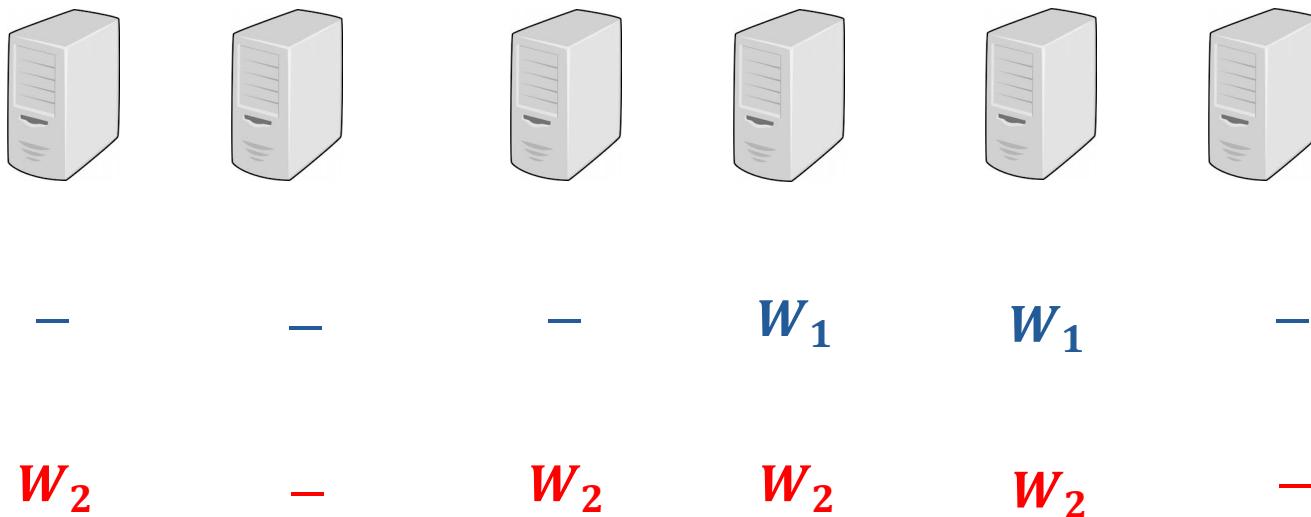
*Data updates may not arrive at all servers simultaneously.*



# Distributed Key-value Stores

## 1. Asynchrony

*Data updates may not arrive at all servers simultaneously.*



# Distributed Key-value Stores

## 1. Asynchrony

*Data updates may not arrive at all servers simultaneously.*

## 2. Decentralized Nature

*A server may not be aware of which updates received by others.*



—

$W_2$

$$S(i) = \{2\}$$



$W_1$

$W_2$

$$S(j) = \{1, 2\}$$


# Distributed Key-value Stores

## 1. Asynchrony

*Data updates may not arrive at all servers simultaneously.*

## 2. Decentralized Nature

*A server may not be aware of which updates received by others.*



# Distributed Key-value Stores

## 1. Asynchrony

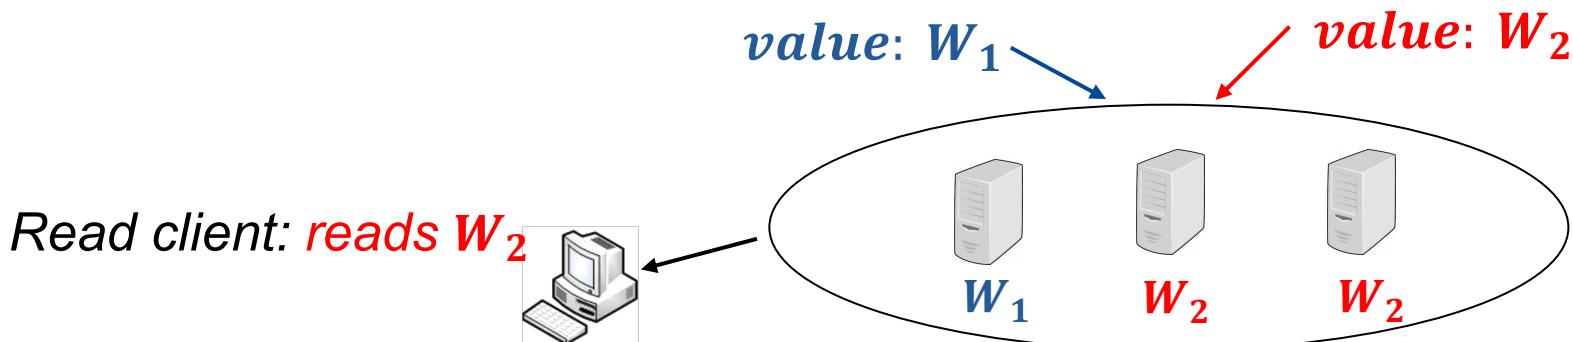
*Data updates may not arrive at all servers simultaneously.*

## 2. Decentralized Nature

*A server may not be aware of which updates received by others.*


## 3. Consistency

*A client must retrieve the **latest possible update**.*



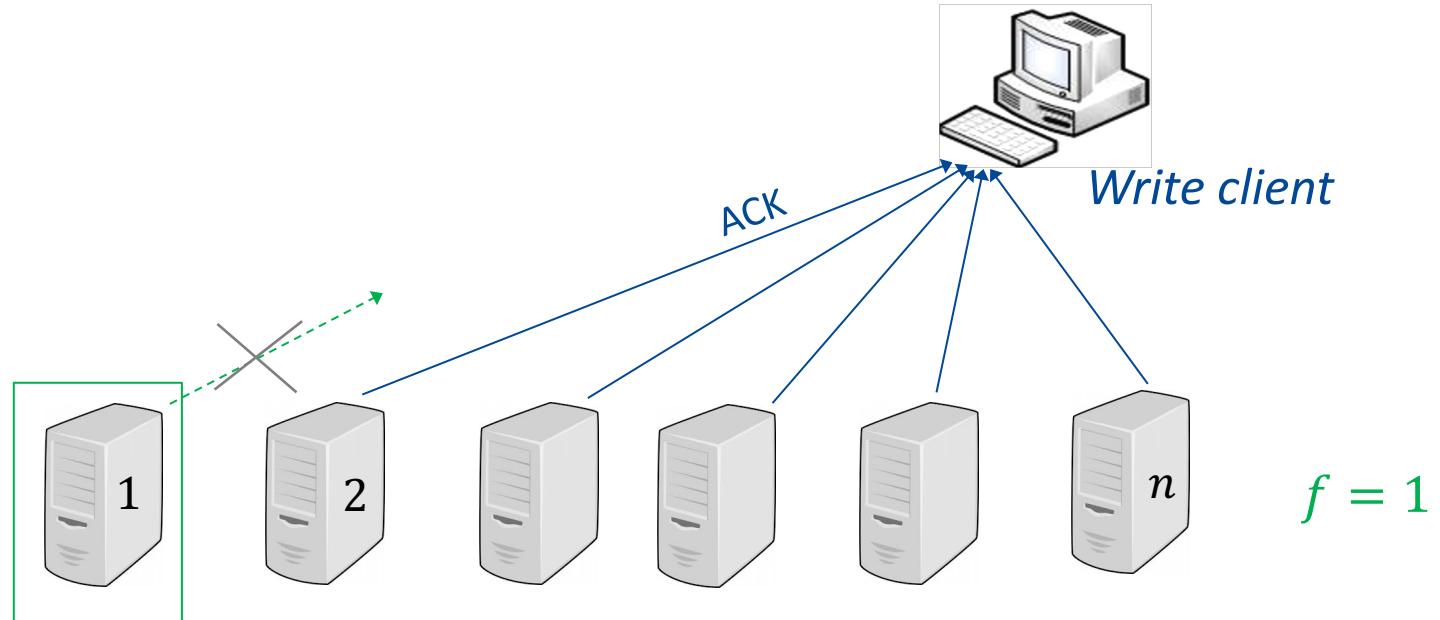
# How to handle asynchrony & failures?

- Fault tolerance:  $f$  failures
- A complete write: write to  $c_W \leq n - f$  servers



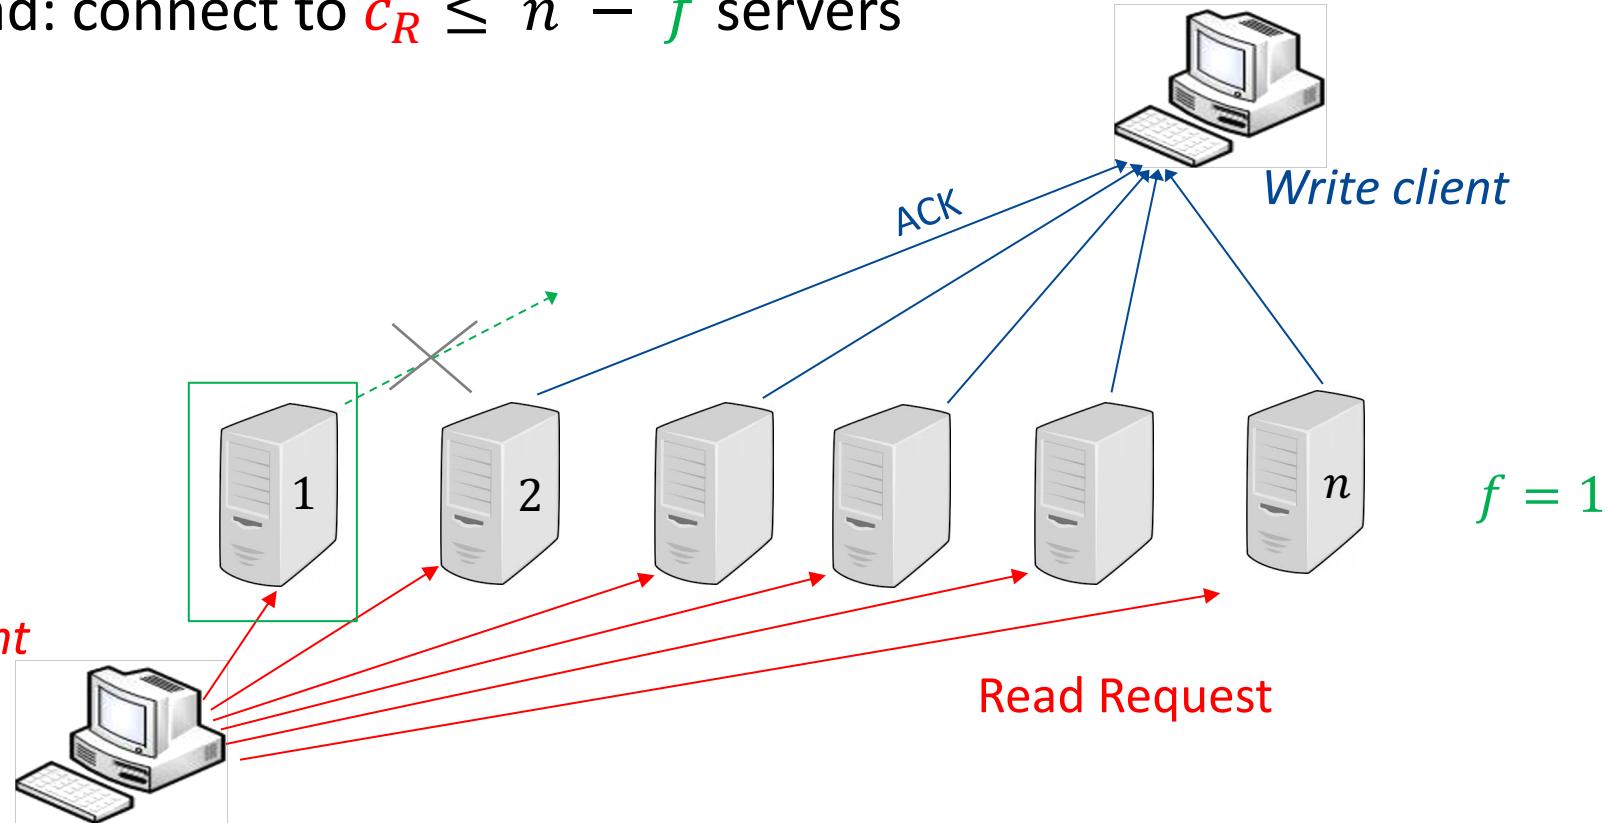
# How to handle asynchrony & failures?

- Fault tolerance:  $f$  failures
- A complete write: write to  $c_W \leq n - f$  servers



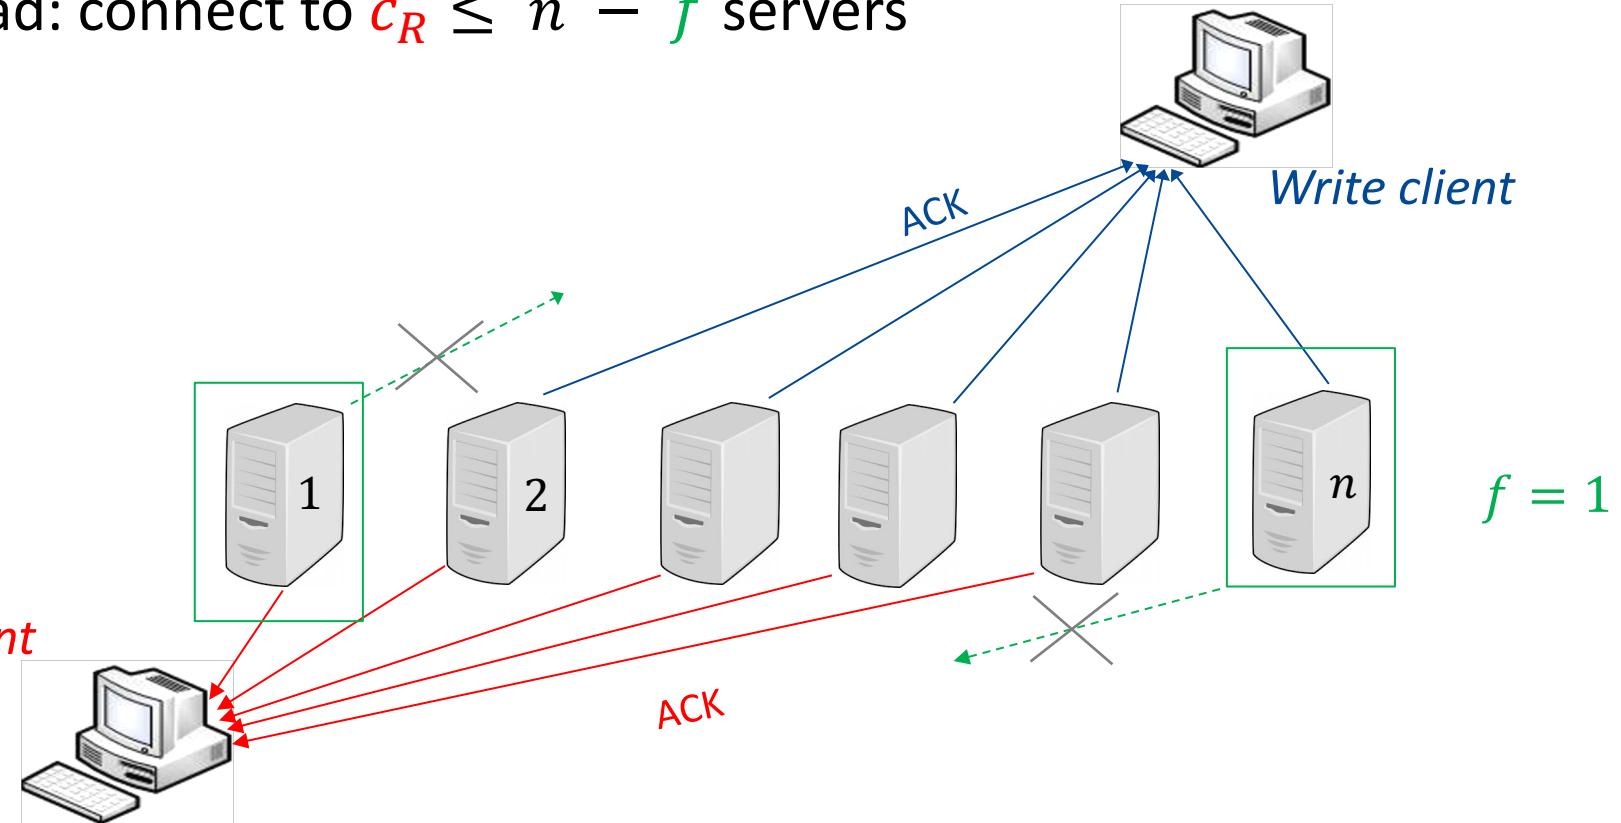
# How to handle asynchrony & failures?

- Fault tolerance:  $f$  failures
- A complete write: write to  $c_W \leq n - f$  servers
- Read: connect to  $c_R \leq n - f$  servers



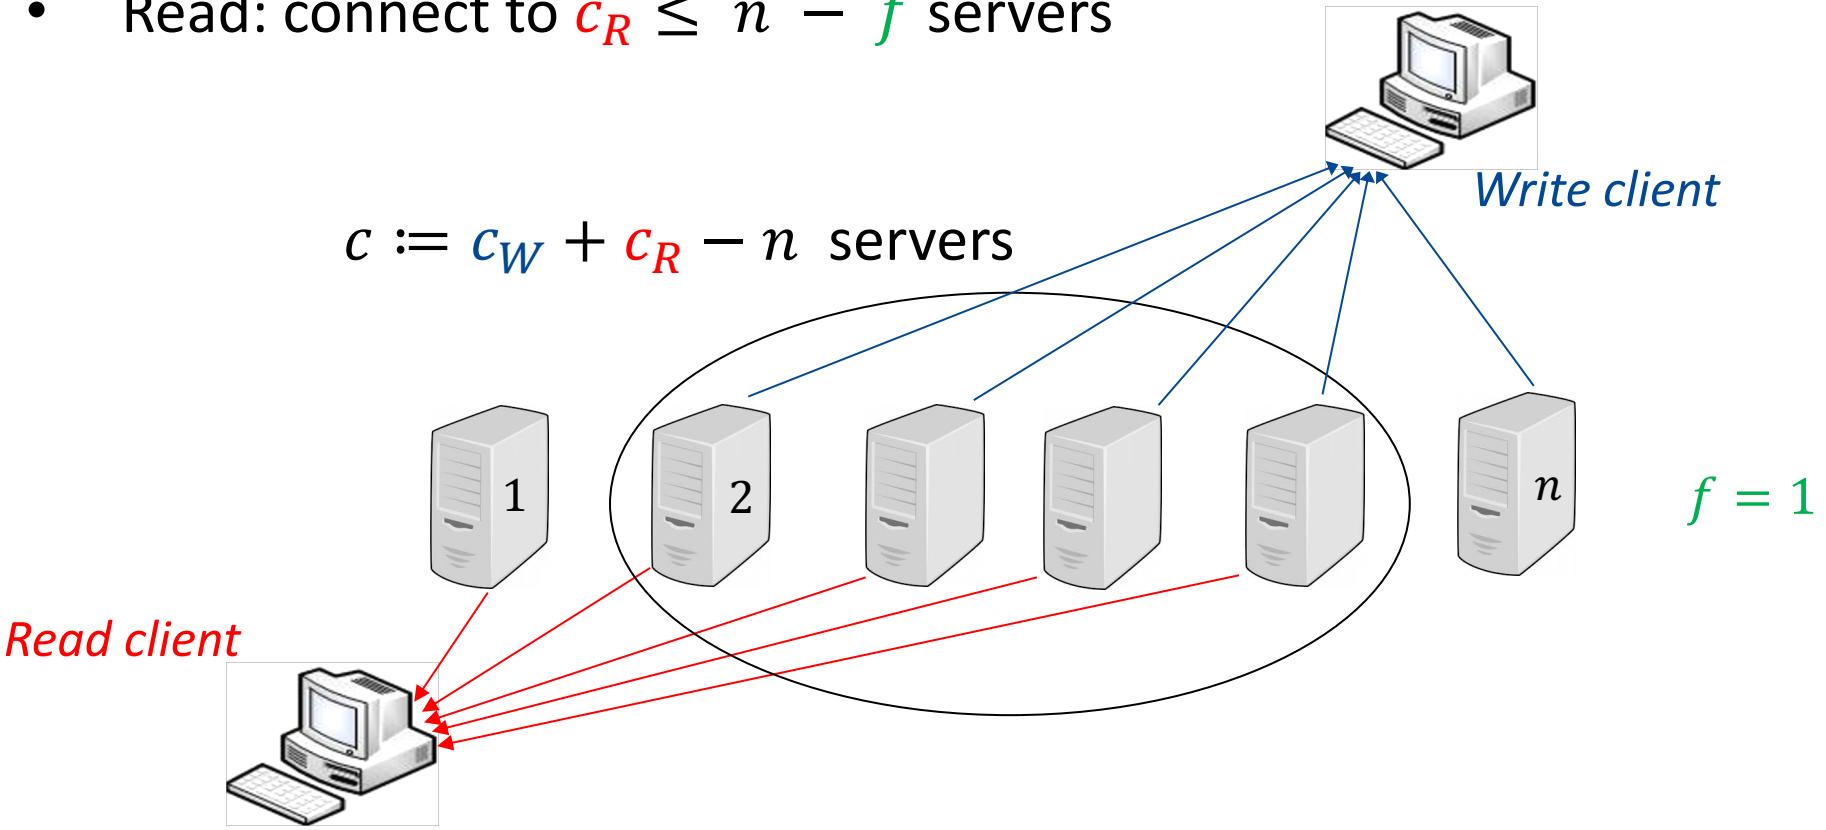
# How to handle asynchrony & failures?

- Fault tolerance:  $f$  failures
- A complete write: write to  $c_W \leq n - f$  servers
- Read: connect to  $c_R \leq n - f$  servers



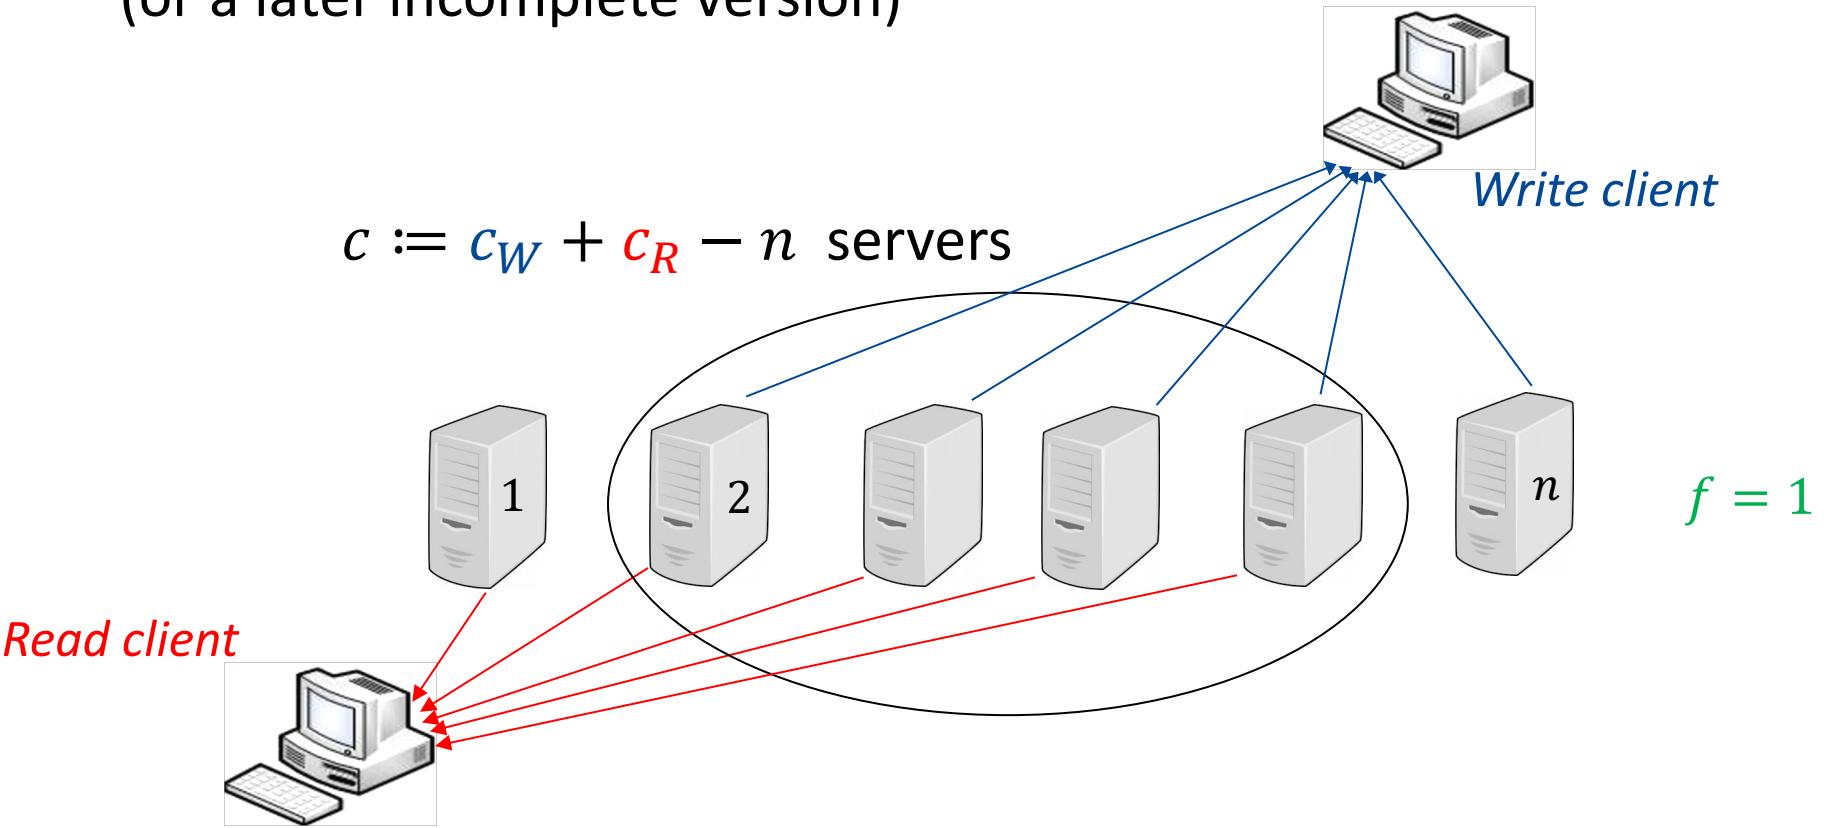
# How to handle asynchrony & failures?

- Fault tolerance:  $f$  failures
- A complete write: write to  $c_W \leq n - f$  servers
- Read: connect to  $c_R \leq n - f$  servers



# How to handle asynchrony & failures?


- **Strong Consistency:** decode the latest complete version (or a later incomplete version)



# How to handle asynchrony & failures?

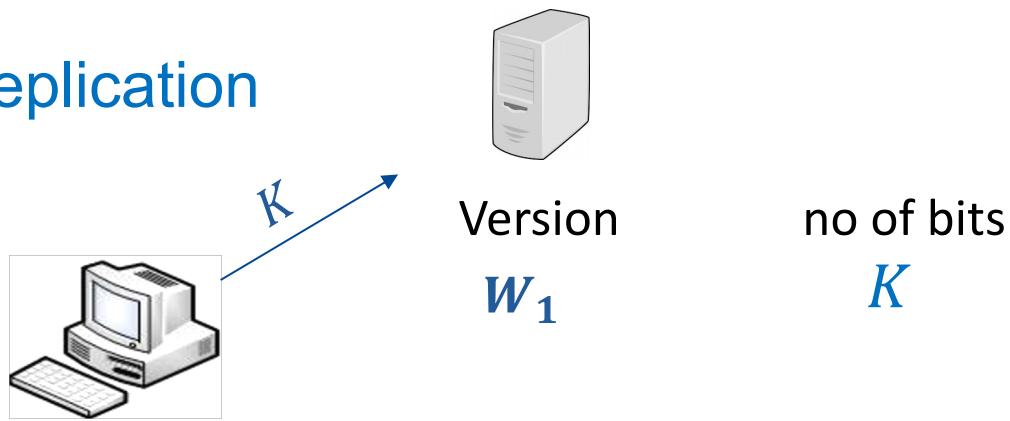
- **Strong Consistency:** decode the latest complete version (or a later incomplete version)

$$c := c_W + c_R - n \text{ servers}$$



# Background: Replication


Replication



*Write client*

# Background: Replication

## Replication




*Write client*

Storage Cost=  $K$

node stores only latest version

# Background: Replication

## Replication



*Write client*

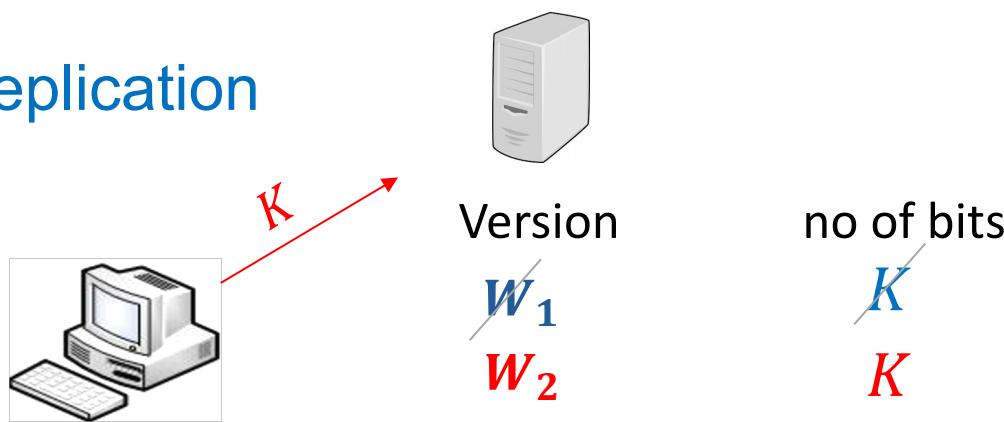
Storage Cost=  $K$

node stores only latest version

Significant Communication and Storage Costs

# Background: Replication

## Replication




*Write client*

Storage Cost =  $K$

node stores only latest version

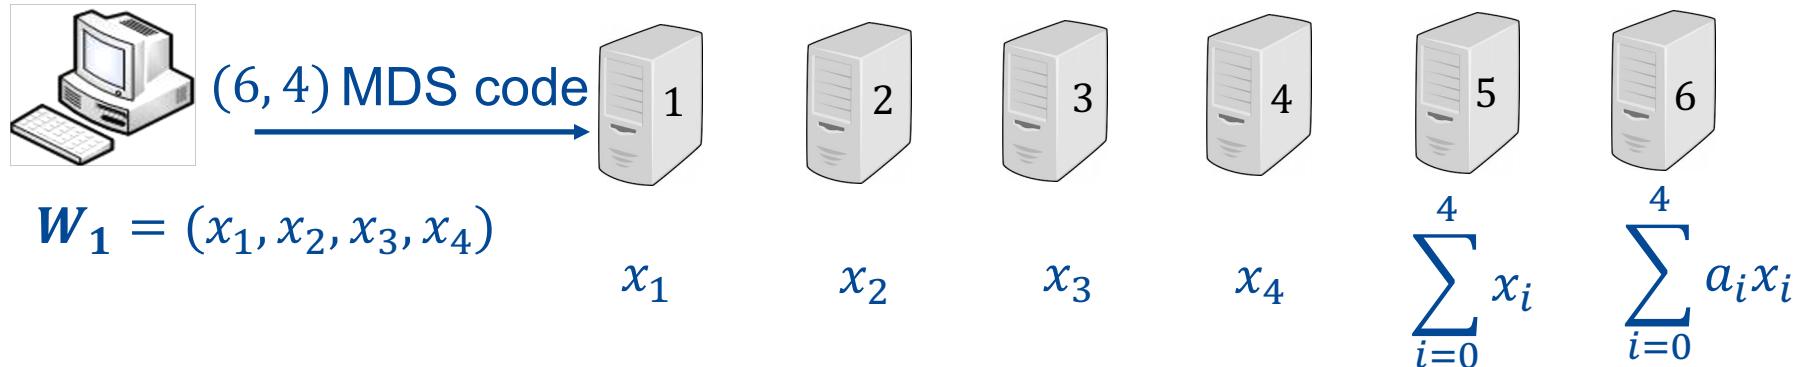
## Significant Communication and Storage Costs

→ Use  $(n, c)$  MDS code, where each node stores  $\frac{1}{c}$  of the data



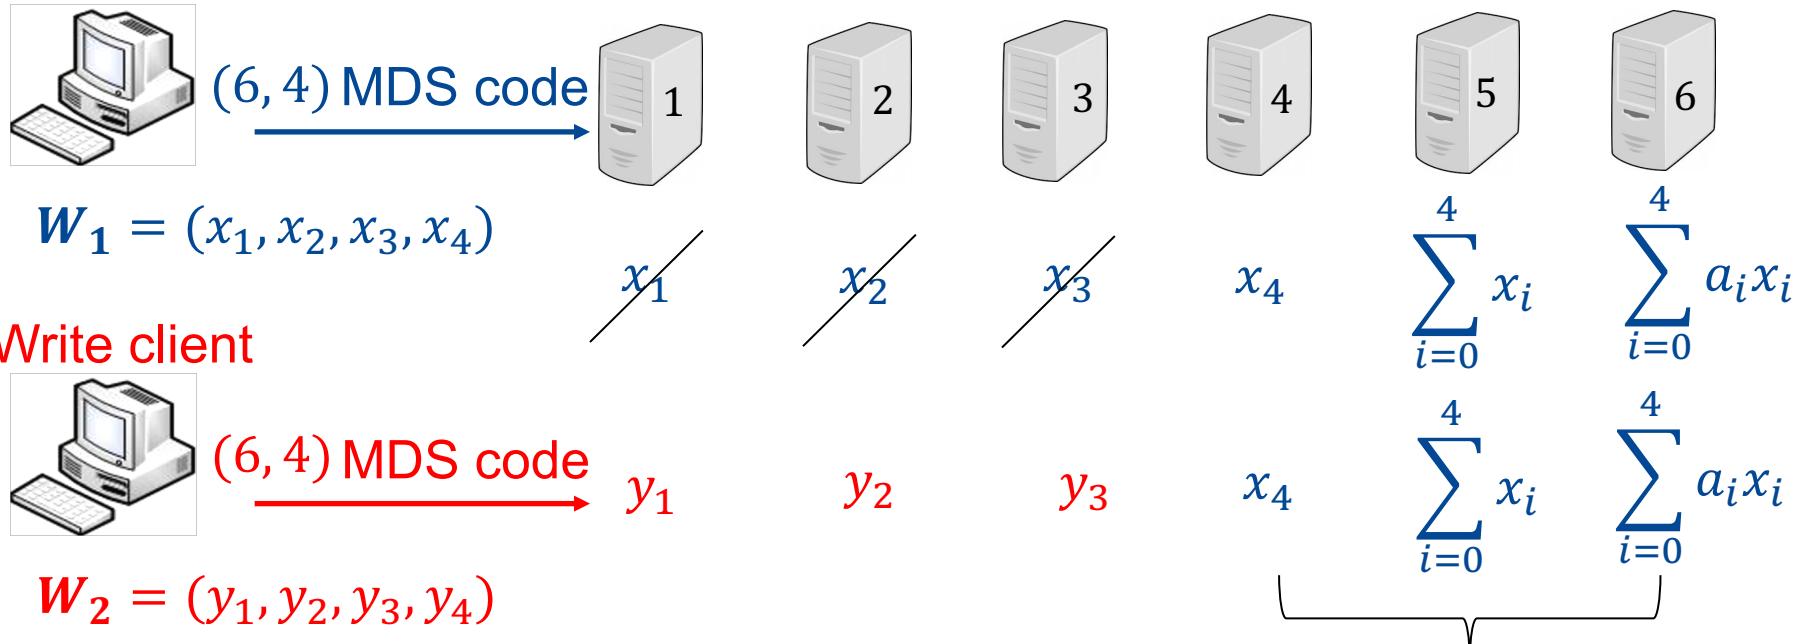
# Background: Erasure Coding Challenges

Write client



# Background: Erasure Coding Challenges

Write client




did not get the new version

# Background: Erasure Coding Challenges

Write client




(6, 4) MDS code



Write client




(6, 4) MDS code



cannot decode

$W_1$  nor  $W_2$



Read client

needs 4 symbols of the same version


did not get the new version

# Background: Erasure Coding Challenges

Write client



(6, 4) MDS code



$$W_1 = (x_1, x_2, x_3, x_4)$$

Write client

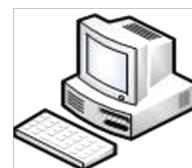


(6, 4) MDS code



$$W_2 = (y_1, y_2, y_3, y_4)$$

can decode  $W_1$




Read client

nodes have to store  
multiple versions

# Background: Erasure Coding Challenges


Replication



Storage Cost=  $K$

*Write client*


Simple Erasure  
Coding



*Write client*


# Background: Erasure Coding Challenges

## Replication



*Write client*

## Simple Erasure Coding



*Write client*

# Background: Erasure Coding Challenges

## Replication



Version

~~$W_1$~~

$W_2$

no of bits

~~$K$~~

$K$

Storage Cost =  $K$

node stores only latest version

## Simple Erasure Coding



Version

$W_1$

$W_2$

no of bits

$K/c$

$K/c$

Storage Cost =  $\nu \frac{1}{c} K$

node stores multiple versions

# Background: Erasure Coding Challenges

Replication



Storage Cost =  $K$

Version

~~$W_1$~~   
 $W_2$

no of bits

~~$K$~~   
 $K$

Simple Erasure  
Coding



Version

$W_1$   
 $W_2$

no of bits

$K/c$   
 $K/c$

Erasure coding gain



Storage Cost =  $\nu \frac{1}{c} K$

# Background: Erasure Coding Challenges

Replication



Storage Cost =  $K$

Version

$\cancel{W_1}$

$W_2$

no of bits

$K$

$K$

Simple Erasure  
Coding



Version

$W_1$

$W_2$

no of bits

$K/c$

$K/c$

Erasure coding gain

Storage Cost =  $\nu \frac{1}{c} K$

Offsets the gain

# Background: Erasure Coding Challenges

Replication



Storage Cost =  $K$

Version

~~$W_1$~~   
 $W_2$

no of bits

$K$   
 $K$

Simple Erasure  
Coding



Version

$W_1$   
 $W_2$

no of bits

$K/c$   
 $K/c$

Erasure coding gain

Storage Cost =  $\nu \frac{1}{c} K$

Offsets the gain

Can we do better?

# Background: Erasure Coding Challenges

Replication



Storage Cost =  $K$

Version

$\cancel{W_1}$

$W_2$

no of bits

$K$

Erasure coding gain

Simple Erasure  
Coding



Version

$W_1$

$W_2$

no of bits

$K/c$

$K/c$

Offsets the gain

[Wang et al. 2014]

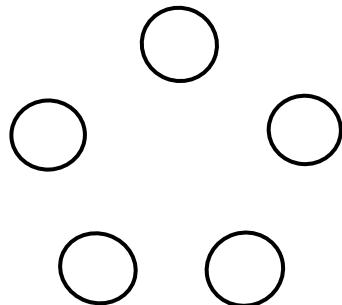
Storage Cost  $\geq \left\lceil \frac{\nu}{2} \frac{1}{c} K - \Theta(1) \right\rceil, \quad \nu < c$

Can we do better?

# Erasure-coded Key-value Stores with Side Information

Decentralized [Wang et al. 2014]

$$\text{Storage Cost} \geq \left( \frac{\nu}{c} - \frac{\nu(\nu - 1)}{c^2} + o\left(\frac{1}{c^2}\right) \right) K$$

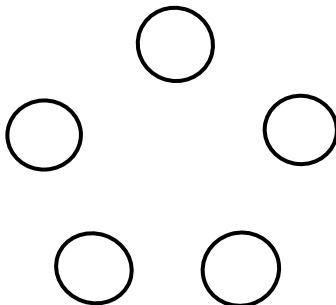


# Erasure-coded Key-value Stores with Side Information

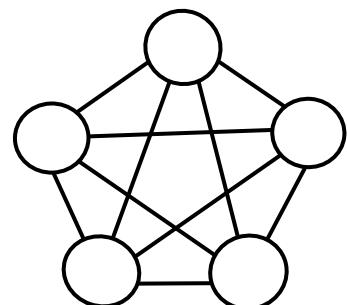
Decentralized [Wang et al. 2014]

Centralized

$$\text{Storage Cost} \geq \left( \frac{\nu}{c} - \frac{\nu(\nu - 1)}{c^2} + o\left(\frac{1}{c^2}\right) \right) K$$



$$\text{Storage Cost} = \frac{1}{c} K$$




# Erasure-coded Key-value Stores with Side Information

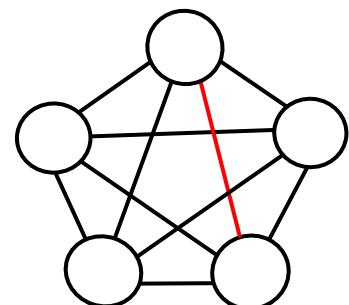
Decentralized [Wang et al. 2014]

Centralized

$$\text{Storage Cost} \geq \left( \frac{\nu}{c} - \frac{\nu(\nu - 1)}{c^2} + o\left(\frac{1}{c^2}\right) \right) K$$



$$\text{Storage Cost} = \frac{1}{c} K$$



High Latency



Geo-distributed  
key-value store

# Erasure-coded Key-value Stores with Side Information


Decentralized [Wang et al. 2014]

Centralized

$$\text{Storage Cost} \geq \left( \frac{\nu}{c} - \frac{\nu(\nu - 1)}{c^2} + o\left(\frac{1}{c^2}\right) \right) K$$



$$\text{Storage Cost} = \frac{1}{c} K$$



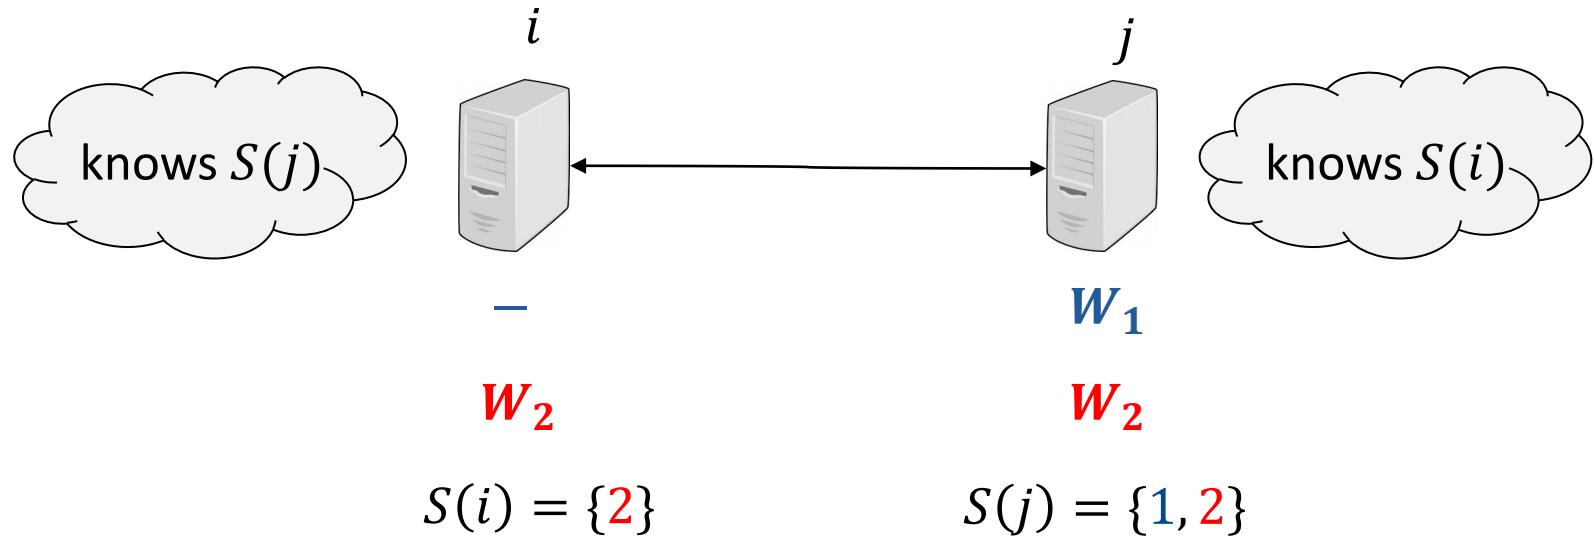
This Work: Coding with Partial Side Information  
Latency-Storage Trade-off

High Latency  
Geo-distributed  
key-value store

# Coding with Side Information

- Topology is given by a directed graph with degree  $H$

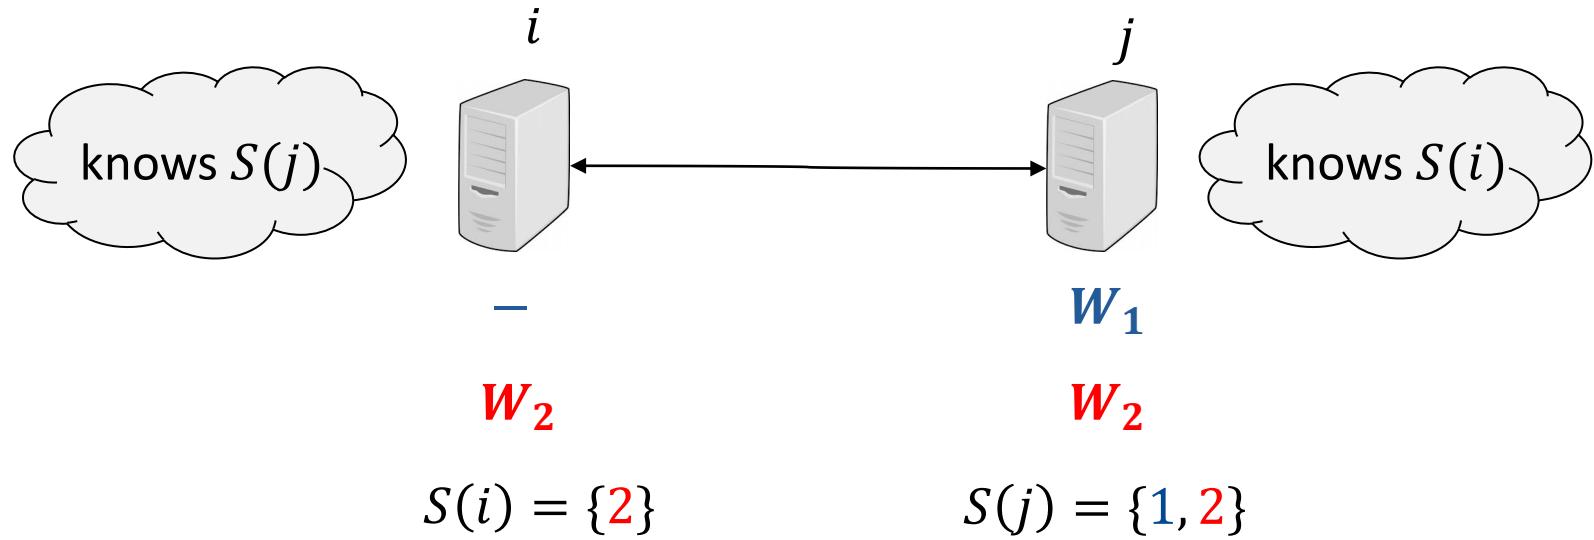
$$\mathcal{G} = (\mathcal{N}, \mathcal{E})$$



# Coding with Side Information

- Topology is given by a directed graph with degree  $H$

$$\mathcal{G} = (\mathcal{N}, \mathcal{E})$$

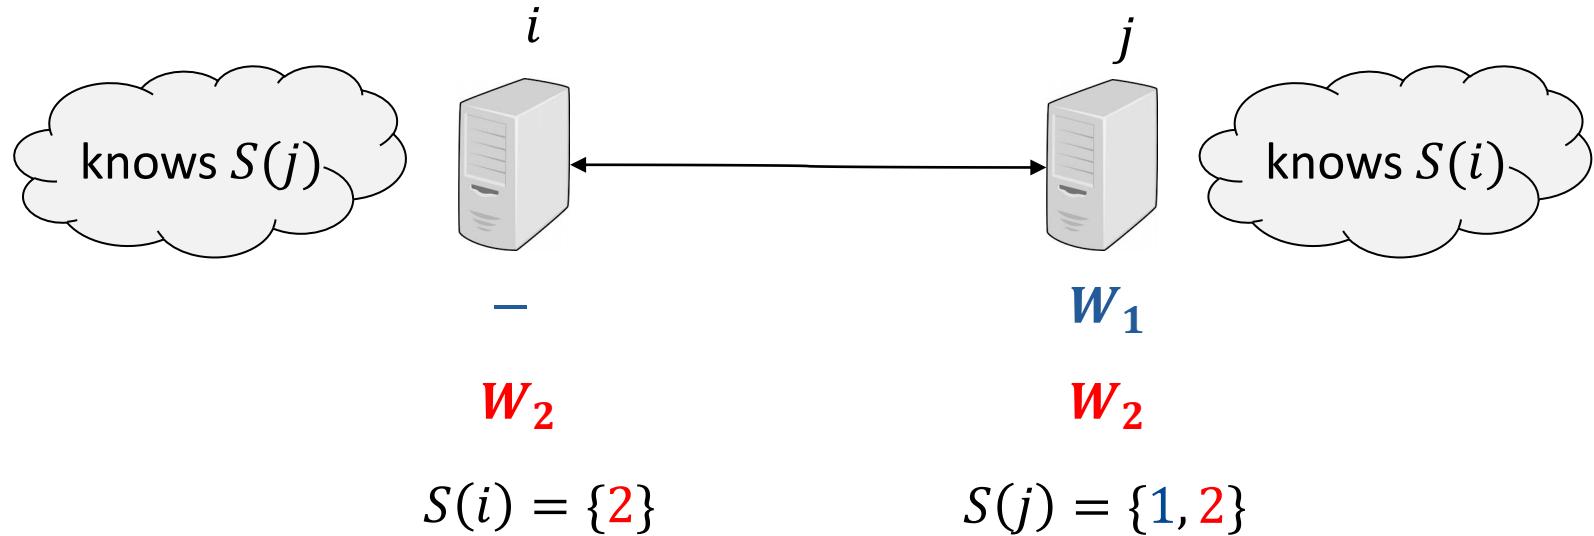


Decoding Requirement: latest complete version (or a later version)

# Coding with Side Information

- Topology is given by a directed graph with degree  $H$

$$\mathcal{G} = (\mathcal{N}, \mathcal{E})$$

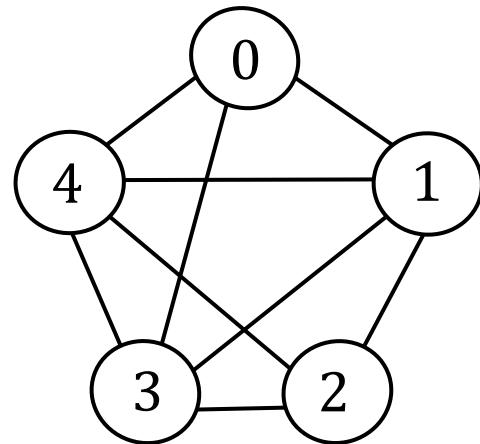


Decoding Requirement: latest complete version (or a later version)

Idea: Can the servers guess which version is the latest complete?

# Coding with Side Information: Challenges

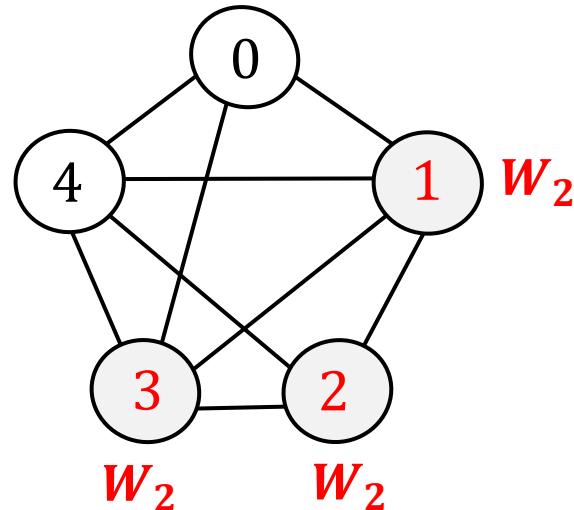
Can the servers guess which version is the latest complete?



$$c_W = 4$$

# Coding with Side Information: Challenges

Can the servers guess which version is the latest complete?




$$c_W = 4$$

$W_2$  is incomplete

# Coding with Side Information: Challenges

Can the servers guess which version is the latest complete?

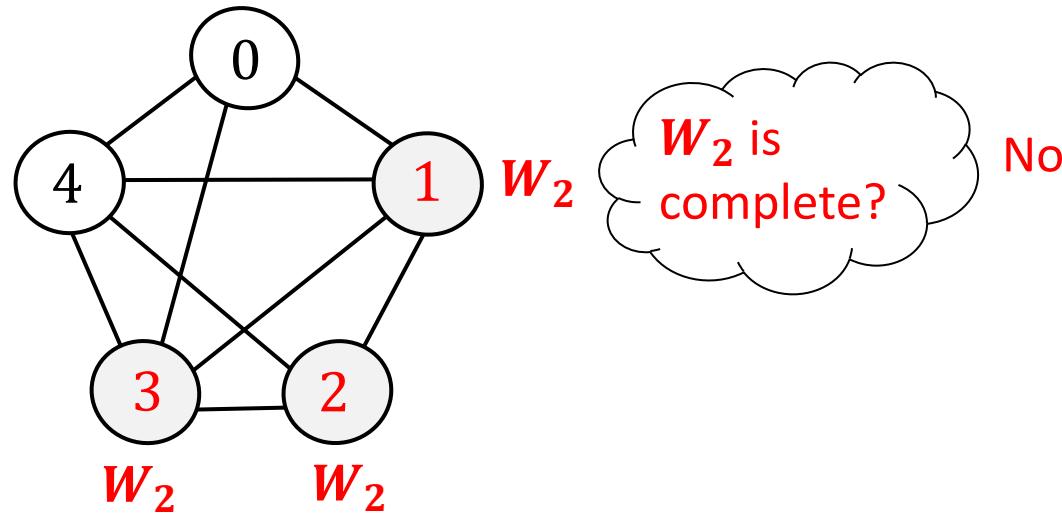


$$c_W = 4$$

$W_2$  is incomplete

# Coding with Side Information: Challenges

Can the servers guess which version is the latest complete?




$$c_W = 4$$

$W_2$  is incomplete

# Coding with Side Information: Challenges

Can the servers guess which version is the latest complete?

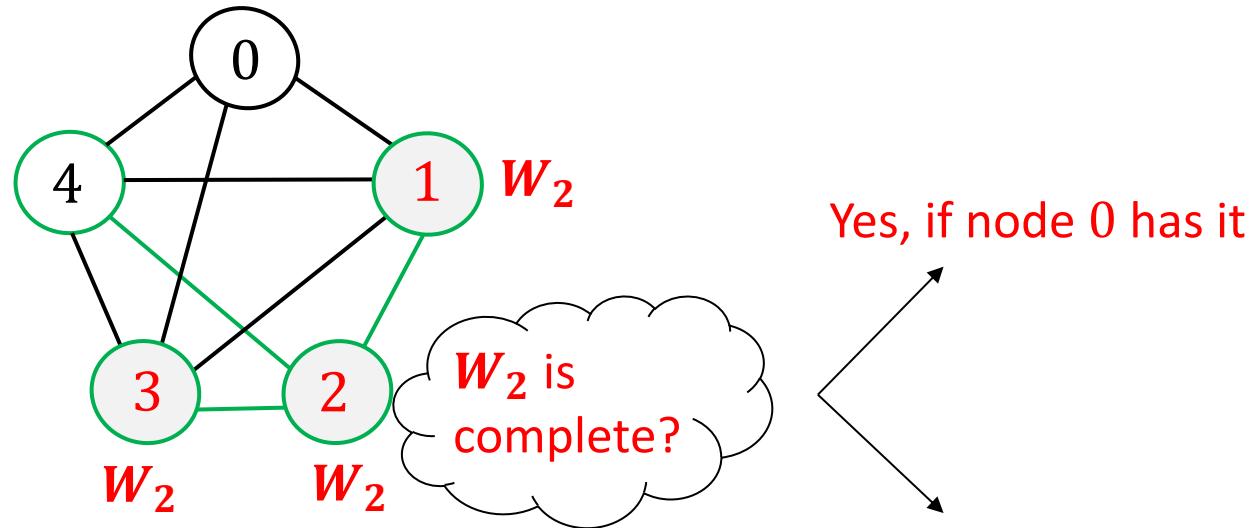


$$c_W = 4$$

$W_2$  is incomplete

# Coding with Side Information: Challenges

Can the servers guess which version is the latest complete?



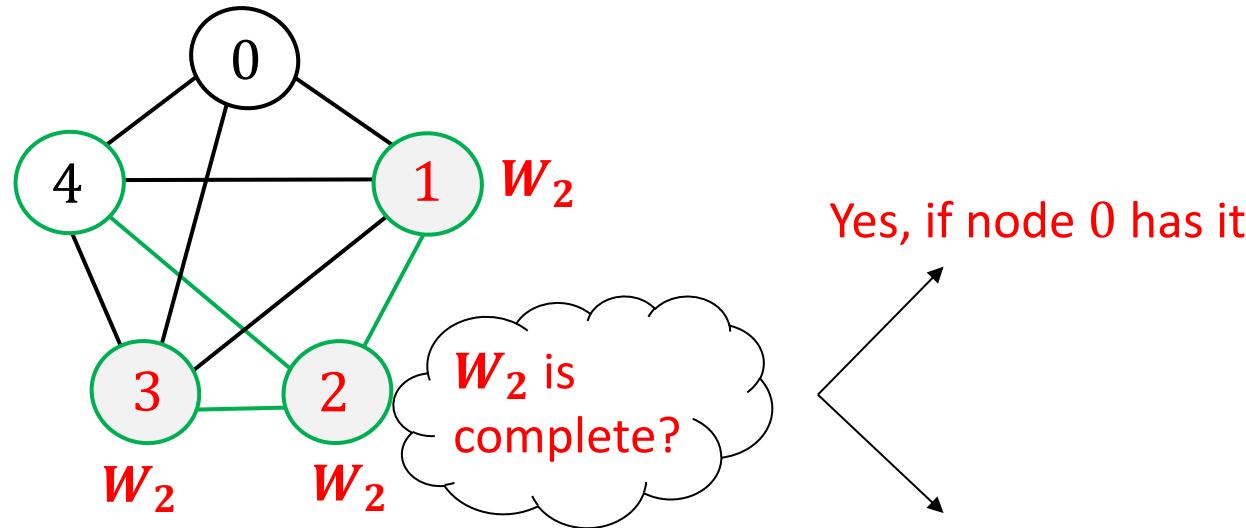
$c_W = 4$   
W<sub>2</sub> is incomplete

Yes, if node 0 has it

No, if node 0 does  
not have it

# Coding with Side Information: Challenges


Can the servers guess which version is the latest complete?



node 2 does not know that  $W_2$  is incomplete!

# Coding with Side Information: Challenges

Can the servers guess which version is the latest complete?




$$c_W = 4$$

Yes, if node 0 has it  
No, if node 0 does not have it

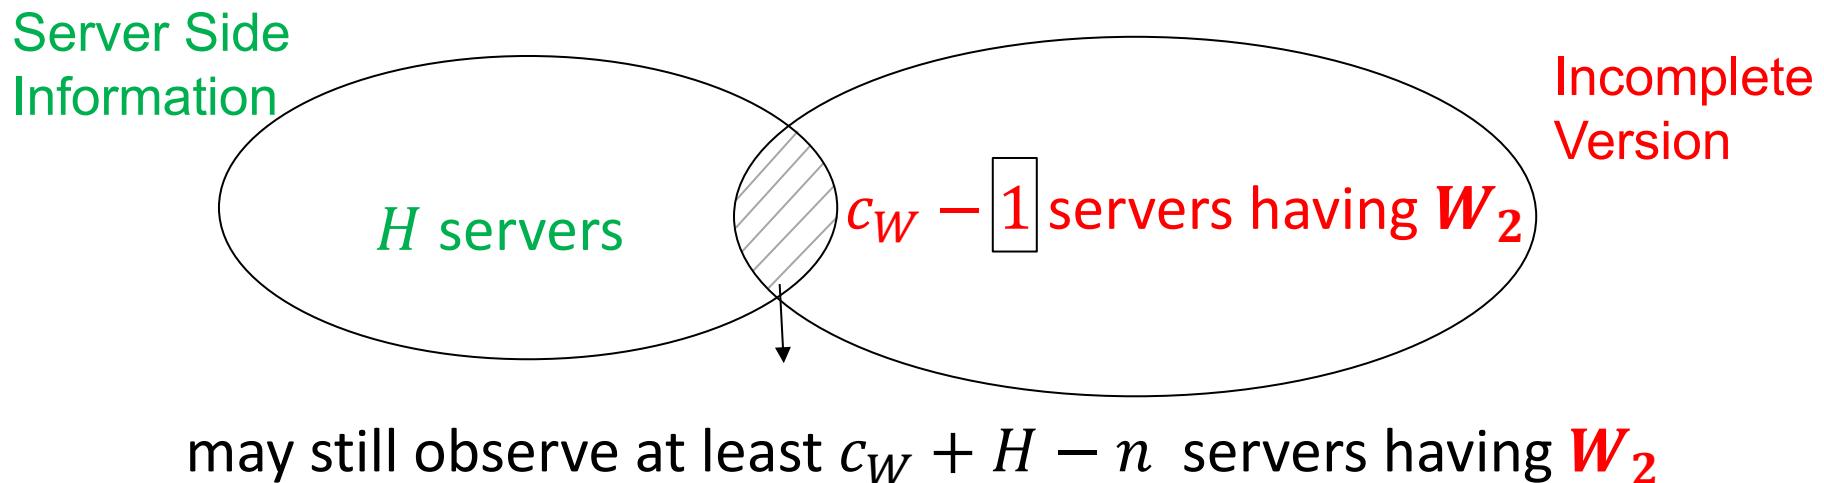

node 2 does not know that  $W_2$  is incomplete!

Given  $\mathcal{G}$ , how many servers cannot guess correctly?


# Coding with Side Information: Challenges



# Coding with Side Information: Challenges




# Coding with Side Information: Challenges



Given an incomplete version, how many servers may assume it is complete?


# Coding with Side Information: Challenges



Given an incomplete version, how many servers may assume it is complete?

$$\bar{m}(\mathcal{G}) = \max_{\mathcal{G}'=(\mathcal{N}', \mathcal{E}') \subset \mathcal{G}: |\mathcal{N}'|=c_W-1} \left| \{i' \in \mathcal{N}': \deg_{\mathcal{G}'}^+(i') \geq c_W + H - n\} \right|$$

# Coding with Side Information: Challenges



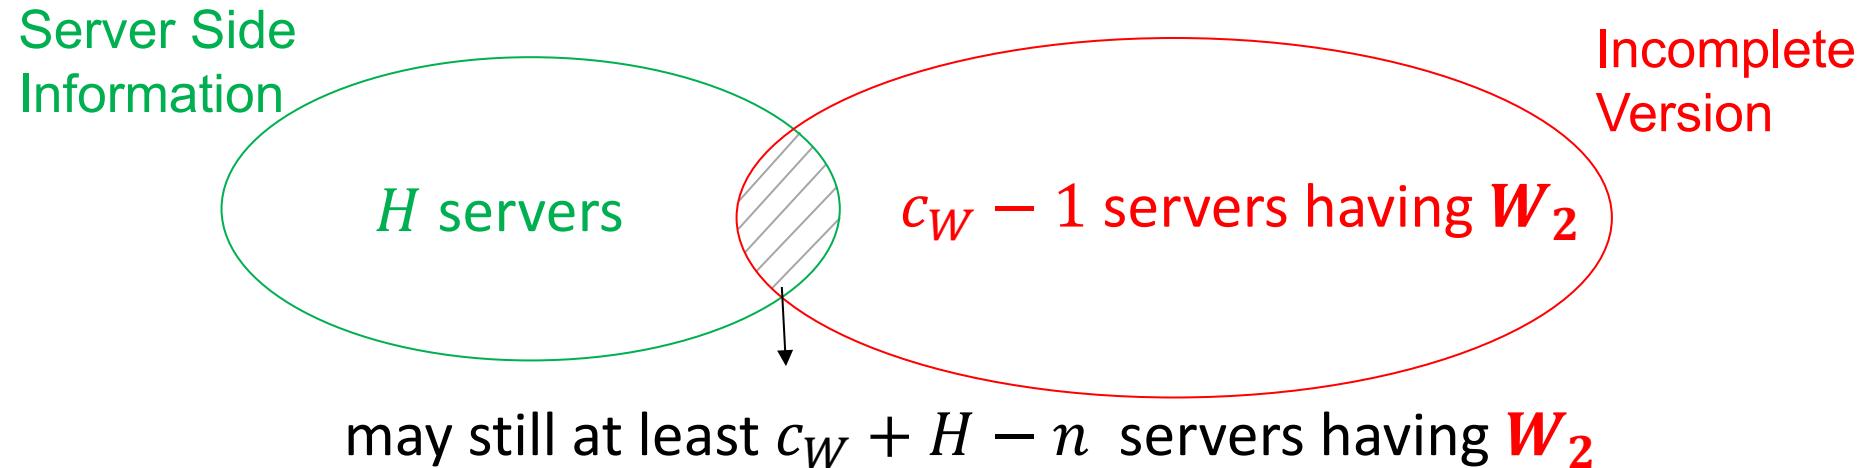
Given an incomplete version, how many servers may assume it is complete?

$$\bar{m}(G) = \max_{G'=(\mathcal{N}', \mathcal{E}') \subset G: |\mathcal{N}'|=c_W-1} \left| \{i' \in \mathcal{N}': \deg_{G'}^+(i') \geq c_W + H - n\} \right|$$

We need to consider  $\binom{n}{c_W - 1}$  graphs

Computationally challenging  
for large graphs

# Coding with Side Information: Challenges



Given an incomplete version, how many servers may assume it is complete?

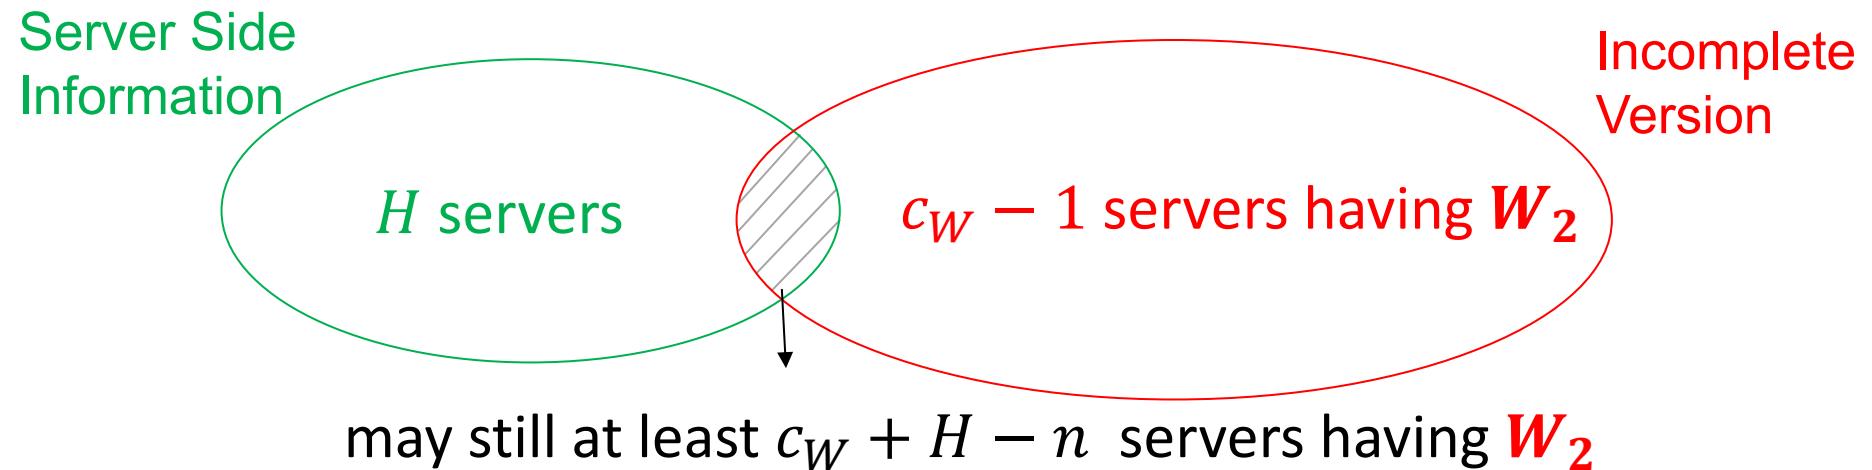
$$\bar{m}(G) = \max_{G'=(\mathcal{N}', \mathcal{E}') \subset G: |\mathcal{N}'|=c_W-1} \left| \{i' \in \mathcal{N}': \deg_{G'}^+(i') \geq c_W + H - n\} \right|$$

We need to consider  $\binom{n}{c_W - 1}$  graphs

Computationally challenging  
for large graphs

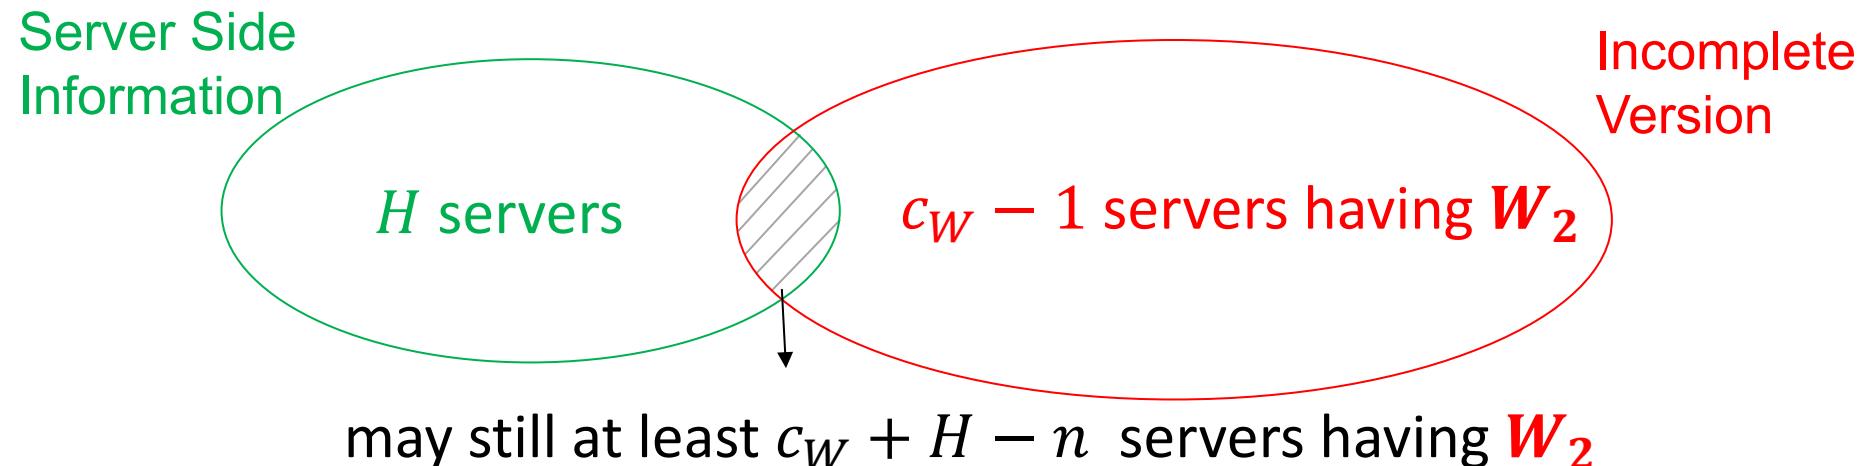
$$\bar{m}(G) \leq (n - c_W + 1) (n - H)$$

# Coding with Side Information: Construction



Coding Strategy: a server stores part of  $\mathbf{W}_2$  if it observes at least  $c_W + H - n$  servers having it.

# Coding with Side Information: Construction



Coding Strategy: a server stores part of **W**<sub>2</sub> if it observes at least  $c_W + H - n$  servers having it.

At most  $\bar{m}(\mathcal{G})$  servers store **W**<sub>2</sub> when it is incomplete.

$$\text{Storage Cost} = \left( \frac{\mathbf{1}}{\mathbf{c}} + \frac{(\nu - 1)\bar{m}(\mathcal{G})}{c^2} + o\left(\frac{\bar{m}(\mathcal{G})}{c^2}\right) \right) K$$




# Coding with Side Information

Decentralized


Partial Information

Centralized

$$\text{Storage Cost} \geq \left( \frac{\nu}{c} - \frac{\nu(\nu - 1)}{c^2} + o\left(\frac{1}{c^2}\right) \right) K$$



$$\text{Storage Cost} = \frac{1}{c} K$$



$$\text{Storage Cost} = \left( \frac{1}{c} + \frac{(\nu - 1)\bar{m}(\mathcal{G})}{c^2} + o\left(\frac{\bar{m}(\mathcal{G})}{c^2}\right) \right) K$$

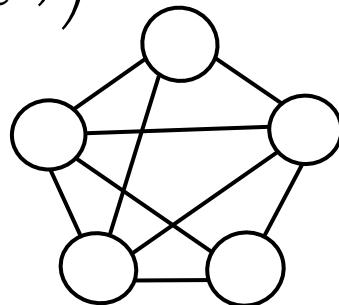
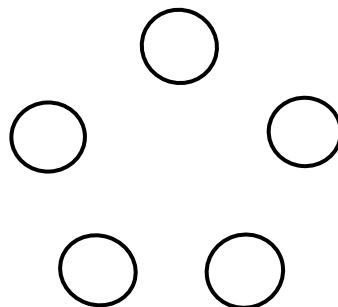
# Coding with Side Information

Decentralized


Partial Information

Centralized

$$\text{Storage Cost} \geq \left( \frac{\nu}{c} - \frac{\nu(\nu - 1)}{c^2} + o\left(\frac{1}{c^2}\right) \right) K$$




$$\text{Storage Cost} = \frac{1}{c} K$$



Storage Reduction = 11%


$$(n = 5, c_W = c_R = 4, \nu = 2)$$

# Impossibility Results



$W_1$  is the latest complete version

# Impossibility Results

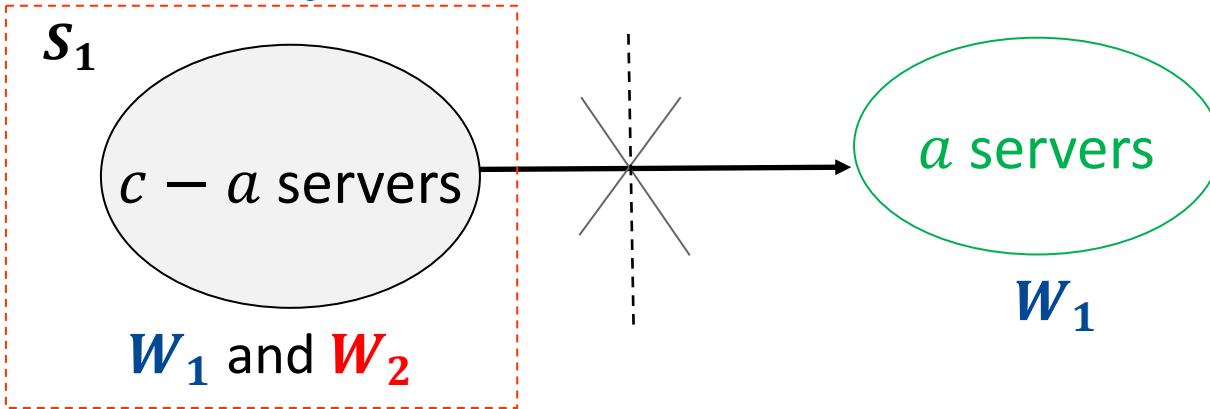


$W_1$  is the latest complete version

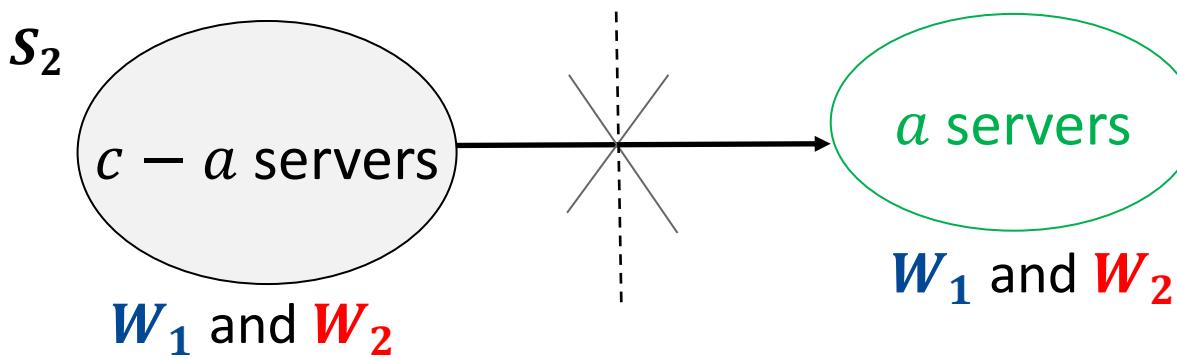
$W_2$  is the latest complete version

cannot differentiate between  $S_1$  and  $S_2$

# Impossibility Results

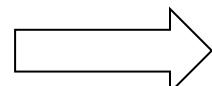


$W_1$  is the latest complete version



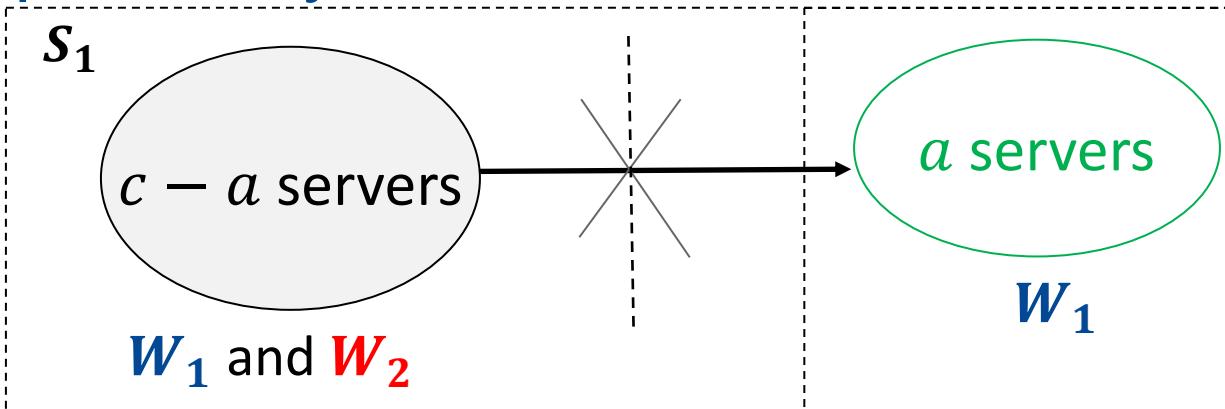
$W_2$  is the latest complete version

Decoding  $W_2$  in  $S_1$

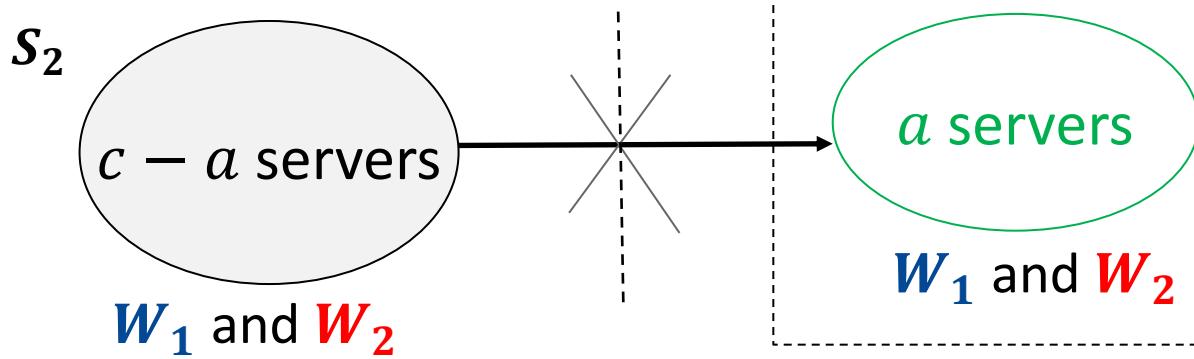


$$\text{Storage Cost} \geq \frac{1}{c - a} K$$

# Impossibility Results



$W_1$  is the latest complete version




$W_2$  is the latest complete version

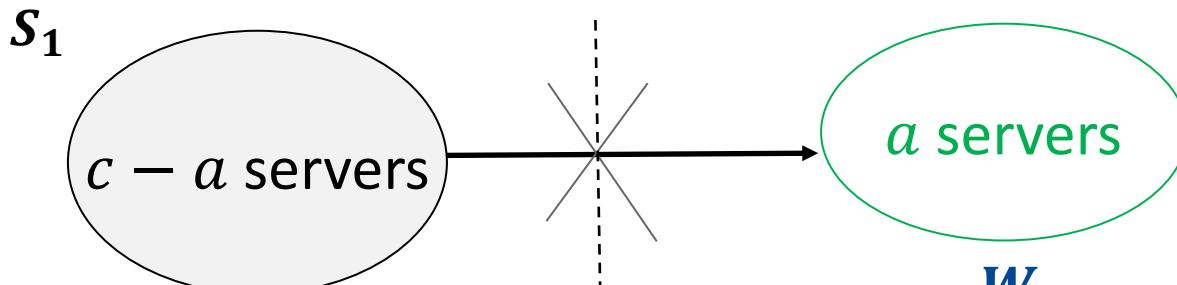
Decoding  $W_1$  in  $S_1$

$$\text{Storage Cost} \geq \frac{2}{c+a} K$$

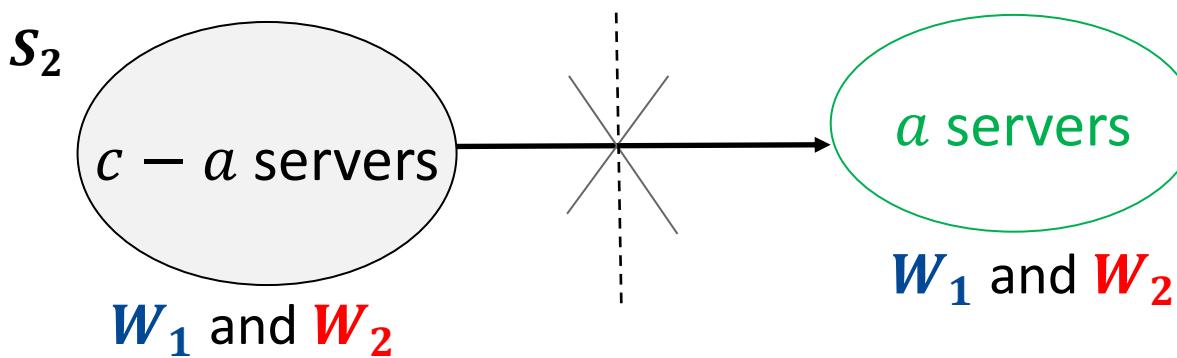
$c$  servers of  $S_1$  and the  $a$  servers of  $S_2$



# Impossibility Results



W<sub>1</sub> is the latest complete version




W<sub>2</sub> is the latest complete version

$$\text{Storage Cost} \geq \min \left\{ \frac{1}{c-a}, \frac{2}{c+a} \right\} K$$

# Impossibility Results

$$\text{Storage Cost} \geq \min \left\{ \frac{1}{c-a}, \frac{2}{c+a} \right\} K$$

Implication:



Side Information is not useful

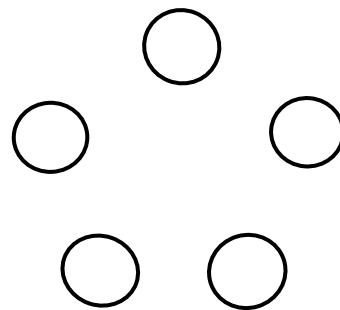
$$(n = 5, c_W = c_R = 4, \nu = 2) \\ (c = 3, a = 1)$$

$$\text{Storage Cost} \geq K/2 \rightarrow$$

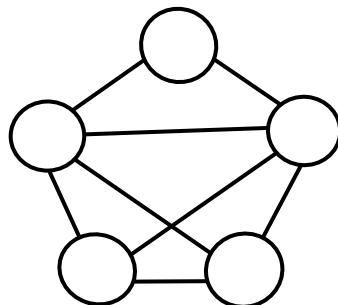
Can be achieved without side information [Wang et al. 2014]

# Coding with Side Information

Decentralized



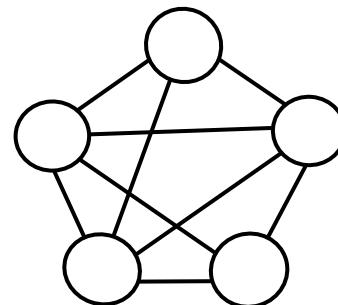
Partial Information




Side Information  
is not useful

$$(n = 5, c_W = c_R = 4, \nu = 2)$$

Centralized

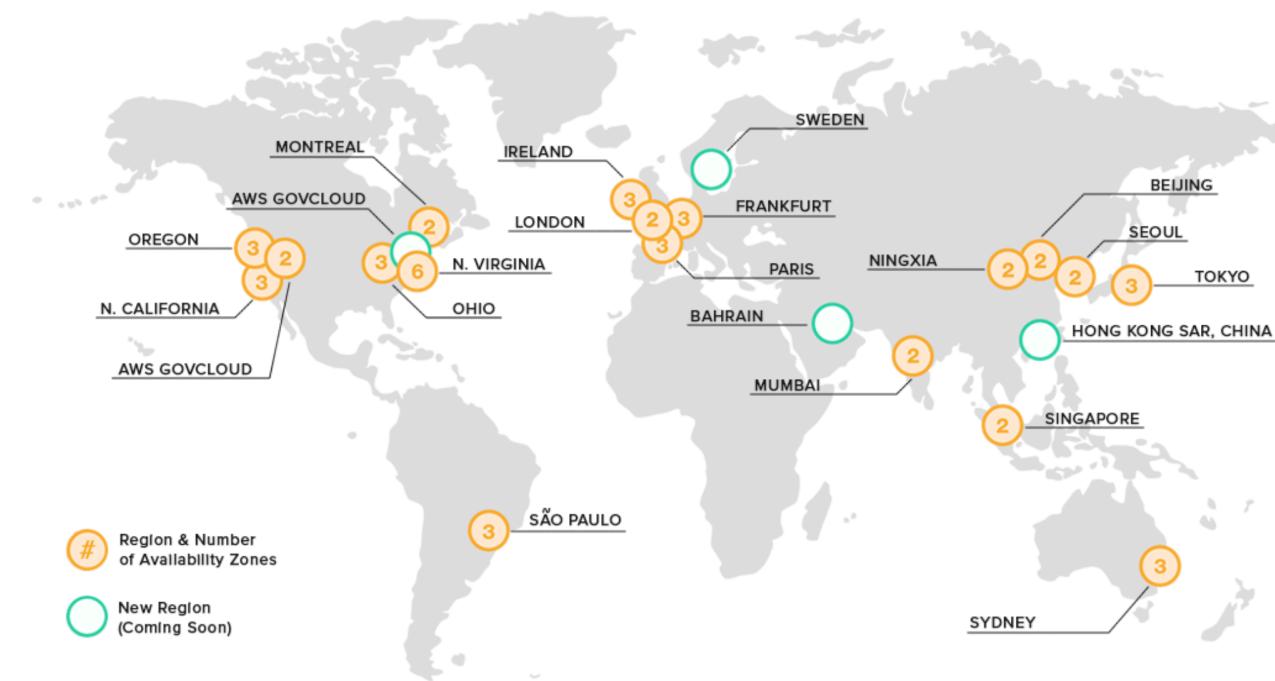


Side Information  
is useful



A careful study of the network topology is necessary

# Case Study: Amazon Web Services (AWS)



| Data center | Location  | Data center | Location    | Data center | Location      |
|-------------|-----------|-------------|-------------|-------------|---------------|
| 1           | Tokyo     | 6           | Frankfurt   | 11          | Ohio          |
| 2           | Seoul     | 7           | Ireland     | 12          | N. California |
| 3           | Mumbai    | 8           | London      | 13          | Oregon        |
| 4           | Singapore | 9           | Paris       |             |               |
| 5           | Canada    | 10          | N. Virginia |             |               |


# Case Study: AWS Inter-Region Latency

| Data center | 1     | 2     | 3     | 4     | 5     | 6     | 7     | 8     | 9     | 10    | 11    | 12    | 13    |
|-------------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|
| 1           | 0     | 37.8  | 157.2 | 90.8  | 177.2 | 249.7 | 234.4 | 259.4 | 259.4 | 167.5 | 166.2 | 119.6 | 106.5 |
| 2           | 37.9  | 0     | 160.1 | 105.7 | 199.7 | 269.9 | 255.7 | 269.3 | 268.2 | 190.7 | 189.3 | 153   | 128.2 |
| 3           | 136.9 | 181.5 | 0     | 68.8  | 212.8 | 129.9 | 134.4 | 128   | 118.3 | 187.7 | 202.2 | 240.8 | 225   |
| 4           | 90    | 112.4 | 82.3  | 0     | 240.9 | 189.7 | 186.4 | 181.3 | 178.5 | 267.8 | 232.6 | 184.7 | 194.7 |
| 5           | 159.2 | 189.5 | 202   | 222.3 | 0     | 103.1 | 81.7  | 92    | 95.4  | 17.8  | 27.2  | 82    | 81.7  |
| 6           | 241.3 | 267.3 | 115.3 | 174.8 | 107   | 0     | 24.2  | 19.1  | 12.8  | 90.4  | 98.9  | 147.8 | 165.4 |
| 7           | 230   | 258.4 | 128.4 | 180   | 85.2  | 23.8  | 0     | 14.6  | 21.6  | 72.7  | 84.6  | 152.8 | 137.4 |
| 8           | 236.9 | 265.3 | 116.9 | 168   | 93.9  | 15.7  | 13.2  | 0     | 10.7  | 78    | 88.7  | 141.7 | 148.5 |
| 9           | 233.5 | 301.6 | 111.6 | 173   | 97.6  | 14.4  | 20.4  | 11    | 0     | 81.7  | 99.4  | 140.7 | 157.8 |
| 10          | 164.3 | 188.8 | 195.8 | 239.9 | 18.8  | 92    | 73.1  | 79.8  | 110.5 | 0     | 13.66 | 67.2  | 79.3  |
| 11          | 162.4 | 189.9 | 199.7 | 226   | 27.6  | 121.5 | 87.7  | 91.3  | 94.6  | 16.4  | 0     | 55.9  | 74.53 |
| 12          | 111.4 | 157.9 | 253.4 | 178.3 | 81.7  | 148.7 | 150.7 | 140   | 146.7 | 67.8  | 53.9  | 0     | 23.4  |
| 13          | 109.8 | 139.7 | 226   | 166.5 | 73.4  | 167.8 | 137.8 | 150.8 | 160.4 | 84    | 73    | 25.8  | 0     |

Source: <https://www.cloudping.co/>


An edge exists between node  $i$  and node  $j$  if the latency between them  $\leq$  maximum allowable latency

# Case Study: Latency-Storage Trade-off in AWS




# Discussion

Decentralized




Partial Information




Side Information  
is not useful

Centralized



Side Information  
is useful



$$(n = 5, c_W = c_R = 4, \nu = 2)$$

*Questions?  
Thank You*